

Building user-level storage data planes with PAIO

Ricardo Macedo
INESC TEC & University of Minho

part 1

background and motivation

Data-centric systems


- Data-centric systems have become an integral part of modern I/O stacks
- Good performance for these systems often requires storage optimizations
 - Scheduling, caching, tiering, replication, ...
- Optimizations are implemented in sub-optimal manner



Challenge #1

✖ Tightly coupled optimizations

- I/O optimizations are single purposed
- Require deep understanding of the system's internal operation model
- Require profound system refactoring
- Limited portability across systems



[1] “SILK: Preventing Latency Spikes in Log-Structured Merge Key-Value Stores”. Balmau et al. USENIX ATC 2019.

[2] “AC-Key: Adaptive Caching for LSM-based Key-Value Stores”. Wu et al. USENIX ATC 2020.


[3] “SpanDB: A Fast, Cost-Effective LSM-tree Based KV Store on Hybrid Storage”. Chen et al. USENIX FAST 2021.

[4] “Evolution of Development Priorities in Key-Value Stores Serving Large-scale Applications: The RocksDB Experience”. Dong et al. USENIX FAST 2021.

Challenge #1

✖ Tightly coupled optimizations

- I/O optimizations are single purposed
- Require deep understanding of the system's internal operation model
- Require profound system refactoring
- Limited portability across systems




SILK's I/O Scheduler

- Reduce tail latency spikes in RocksDB
- Controls the interference between foreground and background tasks
- Required changing several modules, such as *background operation handlers*, *internal queuing logic*, and *thread pools*

Challenge #1

✓ Decoupled optimizations


- I/O optimizations should be disaggregated from the internal logic
- Moved to a dedicated I/O layer
- Generally applicable
- Portable across different scenarios



Challenge #2

✖ Rigid interfaces


- Decoupled optimizations lose granularity and internal application knowledge
- I/O layers communicate through rigid interfaces
- Discard information that could be used to classify and differentiate requests



Challenge #2

✖ Rigid interfaces


- Decoupled optimizations lose granularity and internal application knowledge
- I/O layers communicate through rigid interfaces
- Discard information that could be used to classify and differentiate requests



Challenge #2

✓ Information propagation


- Application-level information must be propagated throughout layers
- Decoupled optimizations can provide the same level of control and performance



Challenge #3

✖ Partial visibility

- Optimizations are oblivious of other systems
- Lack of coordination
- Conflicting optimizations, I/O contention, and performance variation




! **Note:** the storage backend can either be local (e.g., ext4, xfs) or distributed (e.g., Lustre, GPFS)

Challenge #3

✓ Global I/O control

- Optimizations should be aware of the surrounding system stack
- Operate in coordination
- Holistic control of I/O workflows and shared resources



part 2

designing a storage data plane framework


PAIO

- **User-level** framework for building **portable** and **generally applicable** optimizations
- Adopts ideas from **Software-Defined Storage** [6]
 - I/O optimizations are implemented **outside** applications as **data plane stages**
 - **Stages** are controlled through a **control plane** for coordinated access to resources
- Enables the propagation of application-level information through **context propagation**
- Porting I/O layers to use PAIO requires **none to minor** code changes

[5] “PAIO: General, Portable I/O Optimizations with Minor Application Modifications”. Macedo et al. USENIX FAST 2022.


[6] “A Survey and Classification of Software-Defined Storage Systems”. Macedo et al. ACM CSUR 2020.

PAIO design



PAIO design

- I/O differentiation
- I/O enforcement
- Control plane interaction




Policy: *limit the rate of RocksDB's flush operations to X MiB/s*

I/O differentiation

Context propagation:


instrumentation + propagation phases




Identify the origin of POSIX operations (i.e., **foreground**, **compaction**, or **flush** operations)

I/O differentiation


Context propagation:
propagation + classification phases



I/O differentiation




I/O enforcement




PAIO currently supports **Noop** (passthrough) and **DRL** (token-bucket) enforcement objects

I/O enforcement



Requests return to their
original I/O path

Control plane interaction



Implements the **control algorithms** for orchestrating stages (e.g., **tail latency control**, **per-application bandwidth guarantees**)

More about PAIO



PAIO paper

- Context propagation
- PAIO interfaces
- Control algorithms
- Micro and macro experiments

part 3

building storage data planes

Per-application bandwidth control

ABCI supercomputer

- Jobs can be co-located in the same compute node
- Each job runs with dedicated CPU cores, memory, GPU, and storage quota
- Local **disk bandwidth** is **shared**, leading to **I/O interference** and **performance variation**

BLKIO

- cgroup's block I/O controller allows static rate limiting read and write operations
- Adjusting the rate requires stopping and restarting jobs
- Cannot leverage from leftover bandwidth

PAIO

- Stage provides the I/O mechanisms to dynamically rate limit workflows at each instance
 - Integrating PAIO in TensorFlow did not required any code changes (`LD_PRELOAD`)
- Control plane provides a proportional sharing algorithm to ensure per-application bandwidth QoS guarantees

Per-application bandwidth control

ABCI supercomputer

- Jobs can be co-located in the same compute node
- Each job runs with dedicated CPU cores, memory, GPU, and storage quota
- Local disk bandwidth is shared, leading to I/O interference and performance variation

BLKIO

- cgroup's block I/O controller allows **static rate limiting** read and write operations
- Adjusting the rate requires **stopping** and **restarting jobs**
- **Cannot leverage** from **leftover bandwidth**

PAIO

- Stage provides the I/O mechanisms to dynamically rate limit workflows at each instance
 - Integrating PAIO in TensorFlow did not required any code changes (LD_PRELOAD)
- Control plane provides a proportional sharing algorithm to ensure per-application bandwidth QoS guarantees

Per-application bandwidth control

ABCI supercomputer

- Jobs can be co-located in the same compute node
- Each job runs with dedicated CPU cores, memory, GPU, and storage quota
- Local disk bandwidth is shared, leading to I/O interference and performance variation


BLKIO

- cgroup's block I/O controller allows static rate limiting read and write operations
- Adjusting the rate requires stopping and restarting jobs
- Cannot leverage from leftover bandwidth

PAIO


- Stage provides the I/O mechanisms to **dynamically rate limit** workflows at each instance
 - Integrating PAIO in TensorFlow did **not required any code changes** (`LD_PRELOAD`)
- Control plane provides a **proportional sharing algorithm** to ensure per-application bandwidth QoS guarantees

Per-application bandwidth control




I₃ and I₄ cannot meet their bandwidth targets during 31 and 34 minutes

Per-application bandwidth control



Instances cannot be dynamically provisioned with available disk bandwidth

Per-application bandwidth control



PAIO ensures that policies are met at all times, and whenever leftover bandwidth is available, PAIO shares it across active instances

Tail latency control in LSM-based KVS

RocksDB

- Interference between foreground and background tasks generates high latency spikes
- Latency spikes occur due to L₀-L₁ compactions and flushes being slow or on hold

SILK

- I/O scheduler
 - Allocates bandwidth for internal operations when client load is low
 - Prioritizes flushes and low level compactions
 - Preempts high level compactions with low level ones
- Required changing several core modules made of thousands of LoC

PAIO

- Stage provides the I/O mechanisms for prioritizing and rate limiting background flows
 - Integrating PAIO in RocksDB only required adding 85 LoC
- Control plane provides a SILK-based I/O scheduling algorithm

Tail latency control in LSM-based KVS

RocksDB

- Interference between foreground and background t
- Latency spikes occur due to L₀-L₁ compactions and

SILK

- I/O scheduler
 - **Allocates bandwidth for internal operations when client load is low**
 - **Prioritizes flushes and low level compactions**
 - ~~Preempts high level compactions with low level ones~~
- Required changing several core modules made of thousands of LoC


PAIO

- Stage provides the I/O mechanisms for prioritizing and rate limiting background flows
 - Integrating PAIO in RocksDB only required adding 85 LoC
- Control plane provides a SILK-based I/O scheduling algorithm

! Note: By propagating application-level information to the stage, PAIO can enable similar control and performance as system-specific optimizations



Mixture workload

50% read 50% write



System configuration and workload

- 8 client threads and 8 background threads
- Memory limited to 1GB and I/O BW to 200MB/s
- Bursty workload with peaks and valleys




Throughput: high variability due to constant flushes and compactions


99th latency: high tail latency with peaks with an average range between 3 and 15 ms

Mixture workload

50% read 50% write



Throughput: suffers periodic throughput drops due to accumulated backlog




99th latency: low sustained tail latency

PAIO and SILK observe a 4x decrease in absolute tail latency

Data planes for Deep Learning


Storage tiering (Monarch)

- Framework-agnostic storage middleware
- Leverages existing storage tiers of supercomputers
- Accelerates DL training time by up to 28% and 37% in TensorFlow and PyTorch
- Decreases the operations submitted to the PFS



Parallel data prefetching (Prisma)

- Data plane for prefetching training data samples
- Significantly outperforms baseline PyTorch and TensorFlow configurations
- Achieves similar performance as carefully engineered I/O optimizations in TensorFlow



[7] "Accelerating Deep Learning Training Through Transparent Storage Tiering". Dantas et al. ACM/IEEE CCGrid 2022.

[8] "Monarch: Hierarchical Storage Management for Deep Learning Frameworks". Dantas et al. IEEE Cluster@Rex-IO 2021.

[9] "The Case for Storage Optimization Decoupling in Deep Learning Frameworks". Macedo et al. IEEE Cluster@Rex-IO 2021.

Summary and takeaways

- PAIO, a **user-level** framework to build **custom-made** storage **data plane stages**
- Combines ideas from **Software-Defined Storage** and **context propagation**
- **Decouples** system-specific optimizations to **dedicated I/O layers**
- **User-level data planes** enable similar **control** and **I/O performance** as system-specific optimizations
 - Can be applied over (a lot of) different storage scenarios ...

Accelerated Data Analytics and Computing Institute Seminar


Building user-level storage data planes with PAIO

Ricardo Macedo
INESC TEC & University of Minho

 ricardo.g.macedo@inesctec.pt

 github.com/dsrhaslab

 dsr-haslab.github.io



The screenshot shows the GitHub organization profile for `dsrhaslab`. The profile picture is a photograph of a building with "DSR HASLab" written on it. The main navigation bar includes "Overview", "Repositories 10", "Projects", "Packages", "Teams", "People 8", and "Settings". Below the navigation, there is a "Pinned" section with the following repositories:

- paio** (Public, C++, 8 stars, 2 forks): PAIO: General, Portable I/O Optimizations With Minor Application Modifications.
- monarch** (Public, Python, 5 stars): Hierarchical Storage Management for Deep Learning Frameworks.
- cat** (Public, Python, 3 stars): CaT: Content-aware Tracing and Analysis for Distributed Systems.
- bdus** (Public, c): Forked from `albertofaria/bdus`. A framework for implementing Block Devices in User Space.
- prisma** (Public, C++, 5 stars): A data prefetching storage data plane for accelerating DL training performance.