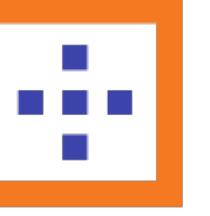


# User-level Software- Defined Storage Data Planes

Ricardo Macedo  
INESC TEC & U. Minho





# Part 1

# background and motivation

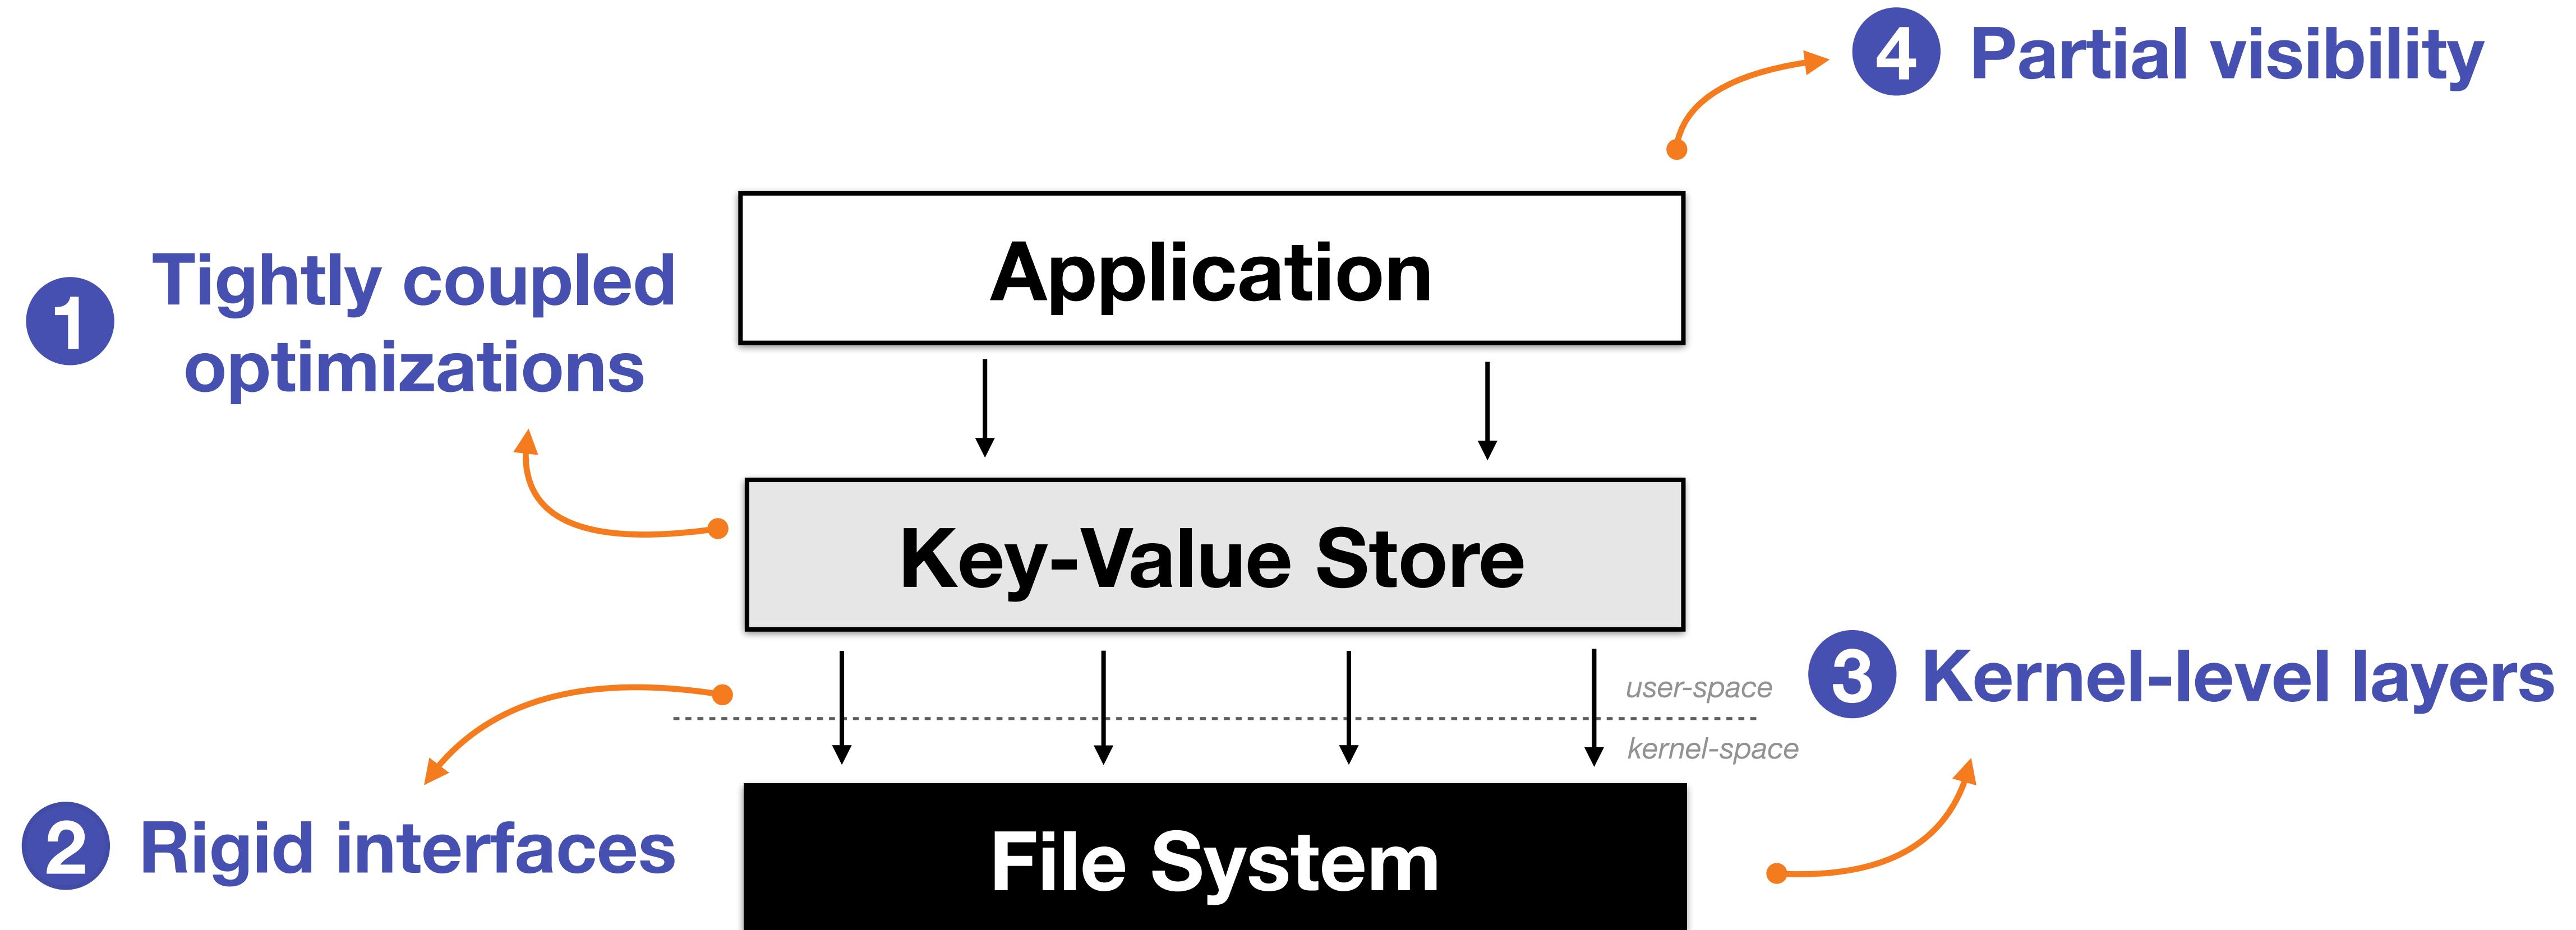
# Data-centric systems

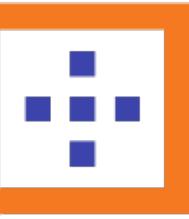
- Data-centric systems have become an integral part of modern I/O stacks
- Good performance for these systems often requires storage optimizations
  - Scheduling, caching, tiering, replication, ...
- Optimizations are implemented in sub-optimal manner





# Challenges

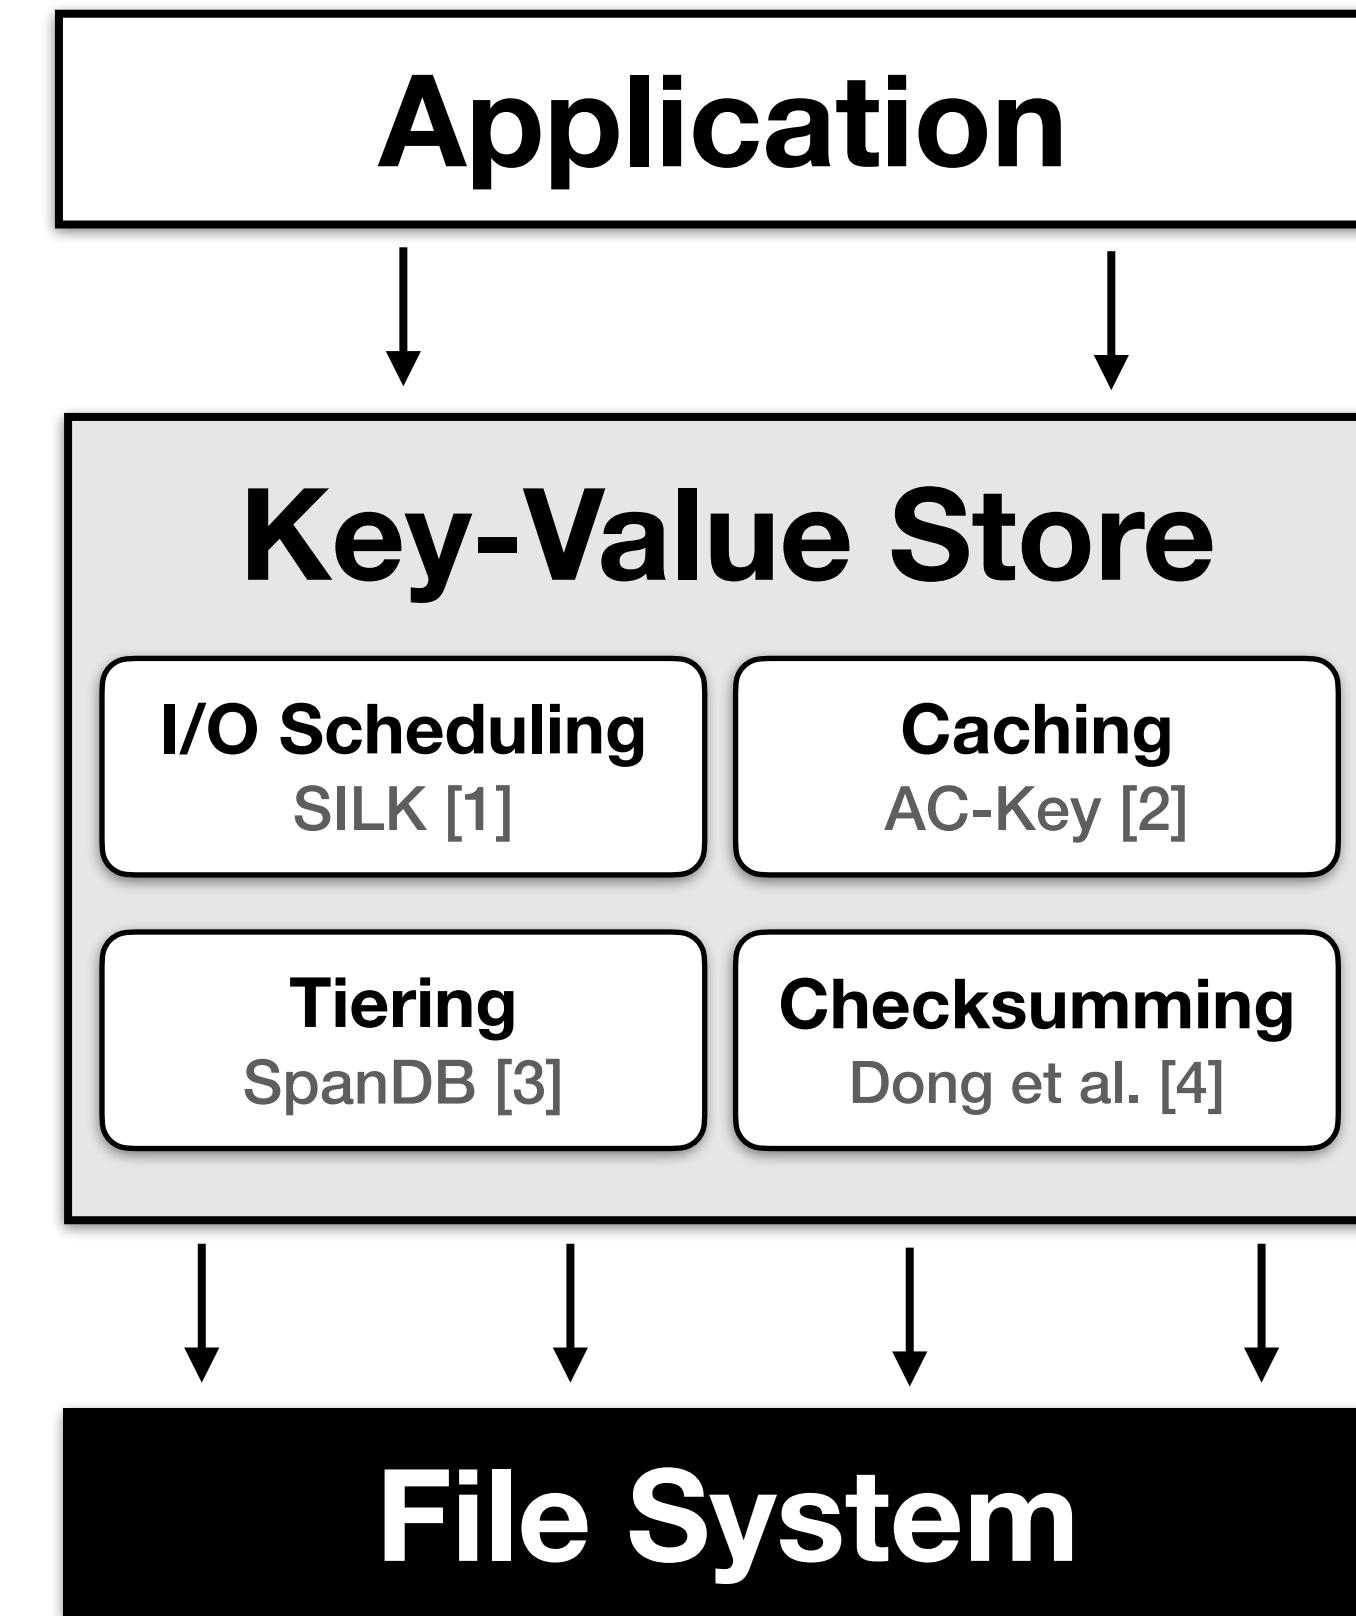




# Challenge #1

## ✖ Tightly coupled optimizations

- I/O optimizations are single purposed
- Require deep understanding of the system's internal operation model
- Require profound system refactoring
- Limited portability across systems

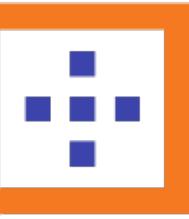


[1] "SILK: Preventing Latency Spikes in Log-Structured Merge Key-Value Stores". Balmau et al. USENIX ATC 2019.

[2] "AC-Key: Adaptive Caching for LSM-based Key-Value Stores". Wu et al. USENIX ATC 2020.

[3] "SpanDB: A Fast, Cost-Effective LSM-tree Based KV Store on Hybrid Storage". Chen et al. USENIX FAST 2021.

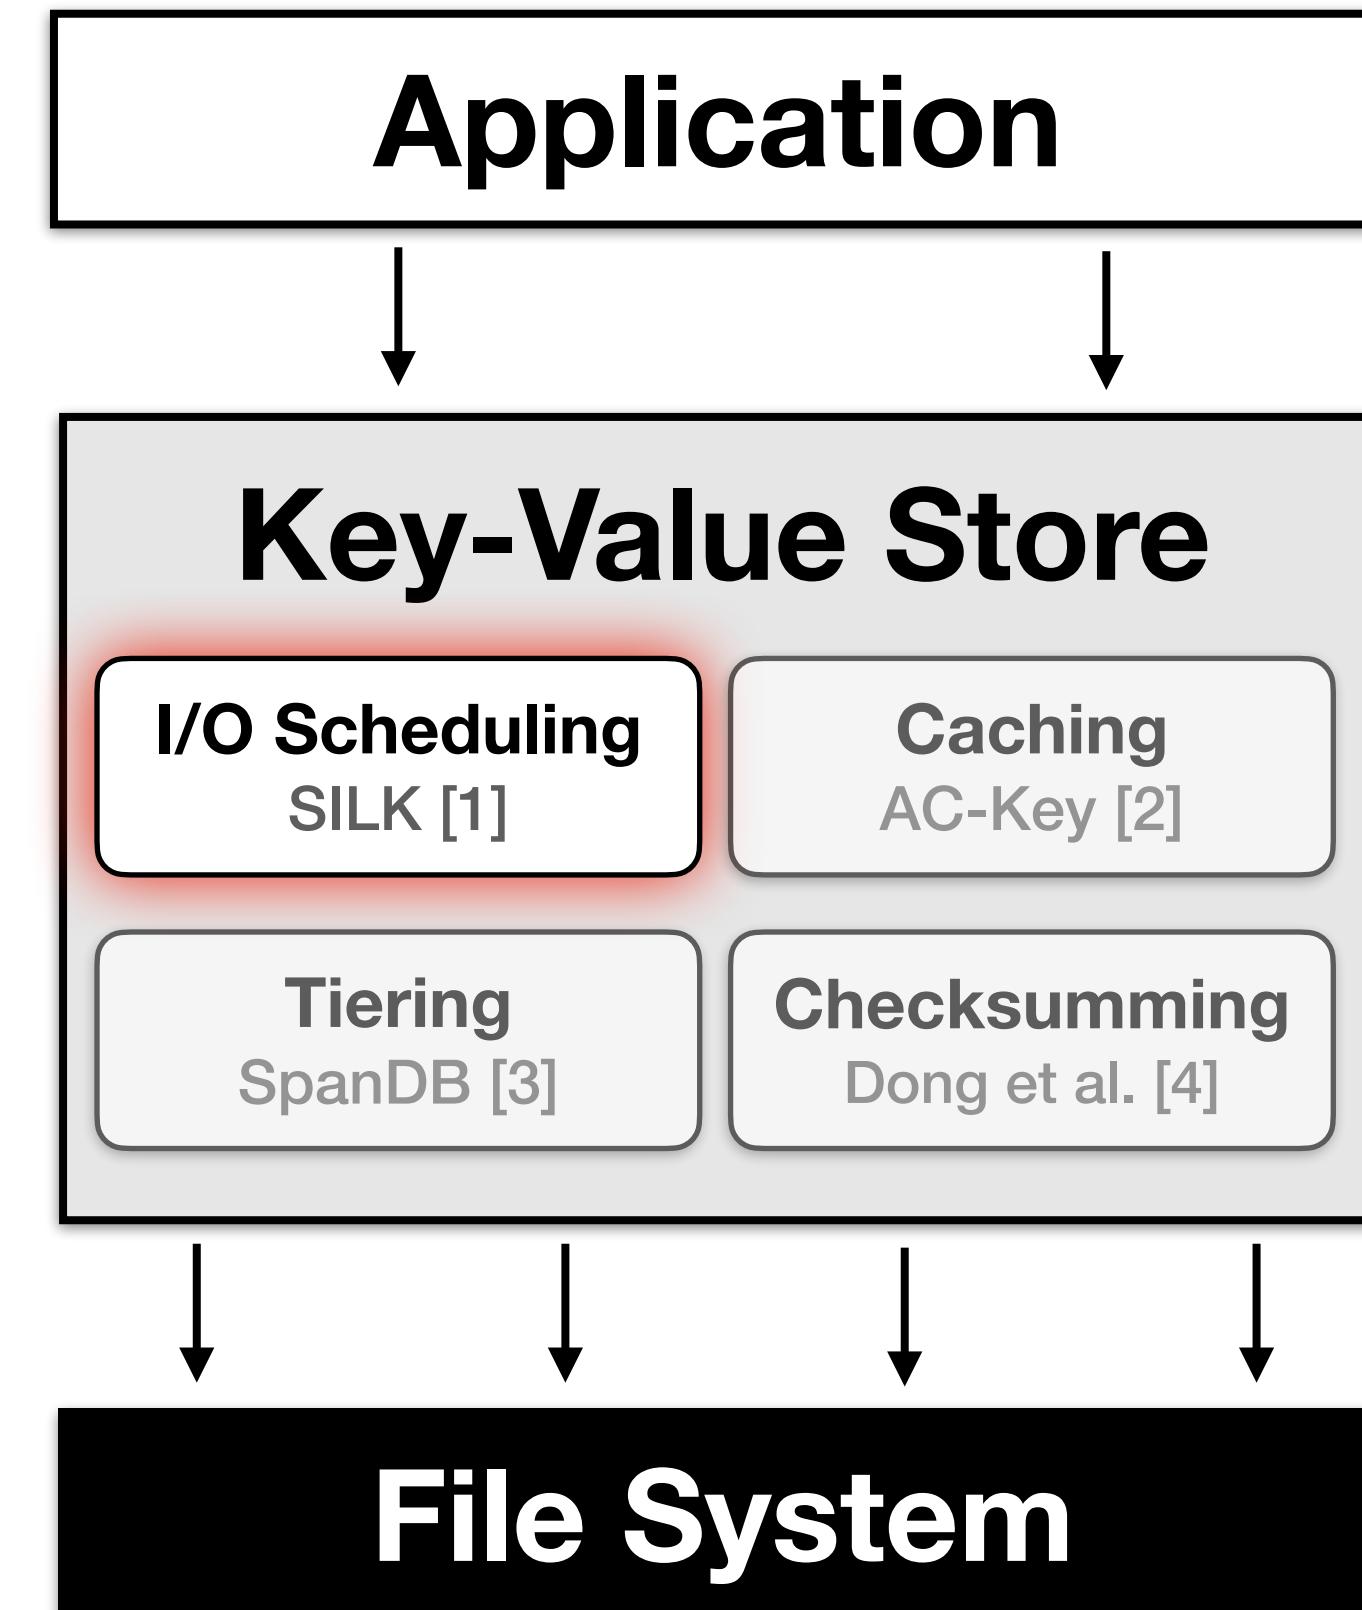
[4] "Evolution of Development Priorities in Key-Value Stores Serving Large-scale Applications: The RocksDB Experience". Dong et al. USENIX FAST 2021.



# Challenge #1

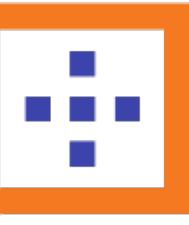
## ✖ Tightly coupled optimizations

- I/O optimizations are single purposed
- Require deep understanding of the system's internal operation model
- Require profound system refactoring
- Limited portability across systems



## SILK's I/O Scheduler

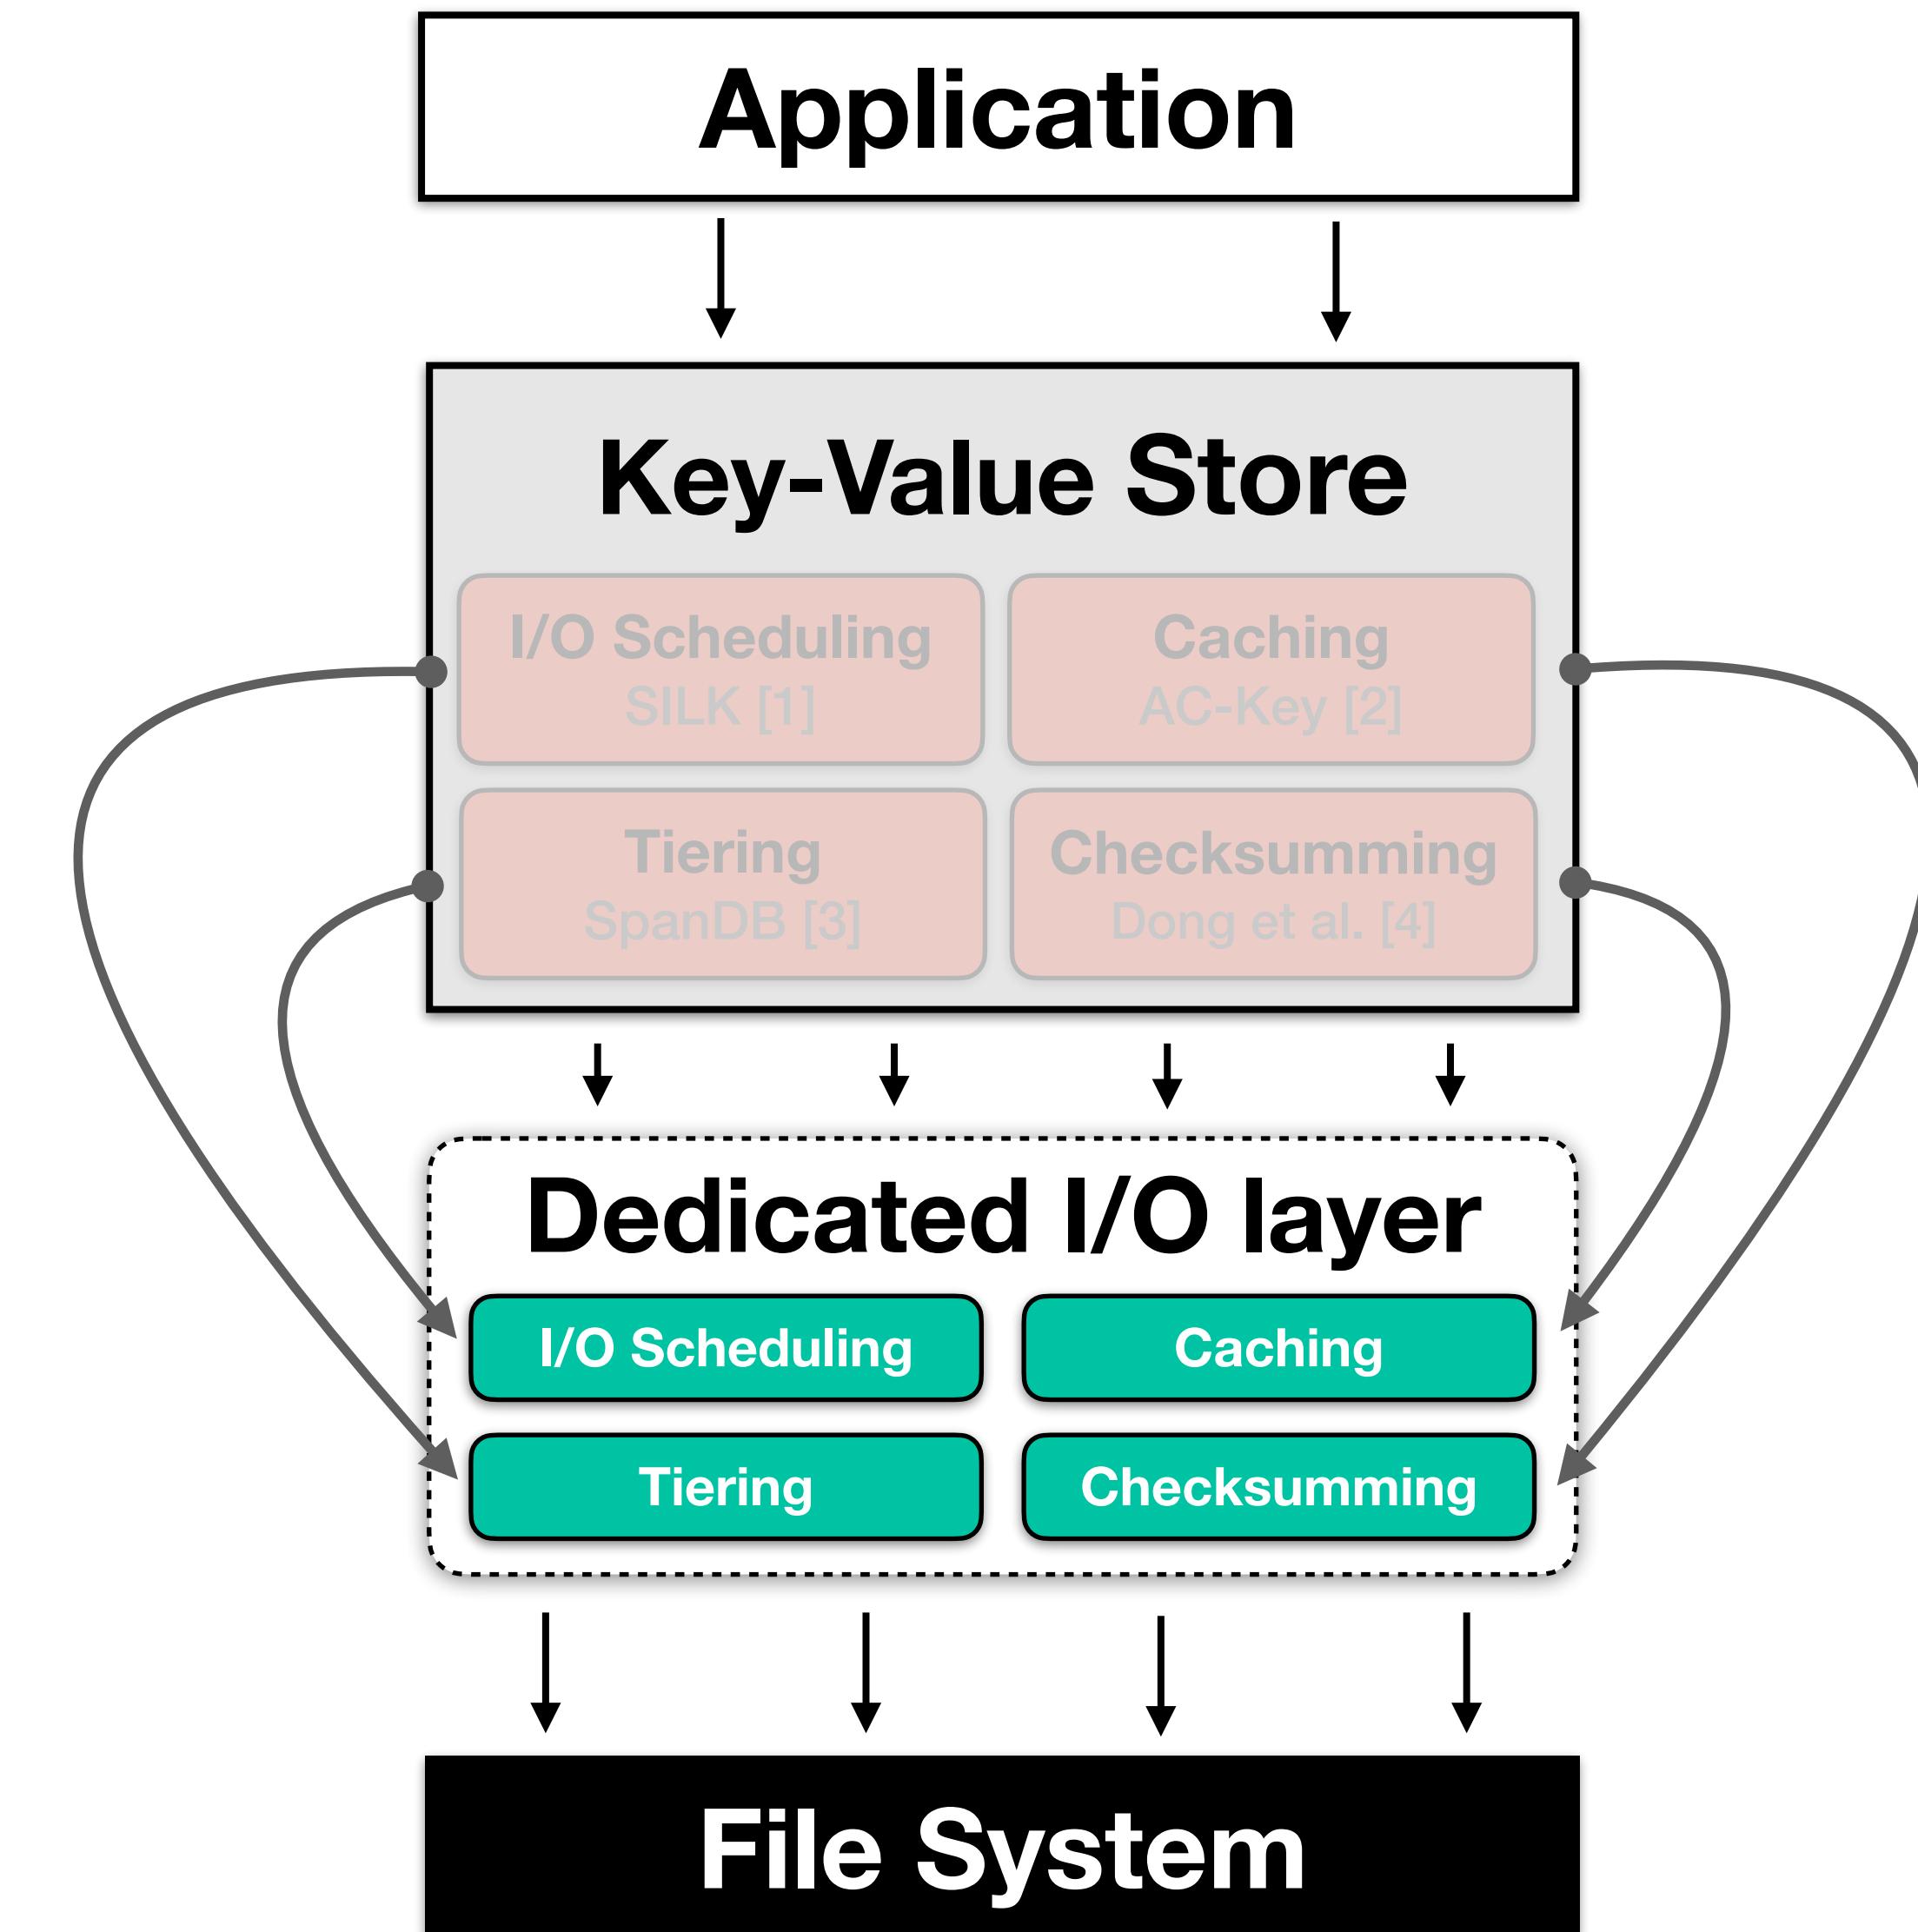
- Reduce tail latency spikes in RocksDB
- Controls the interference between foreground and background tasks
- Required changing several modules, such as *background operation handlers*, *internal queuing logic*, and *thread pools*

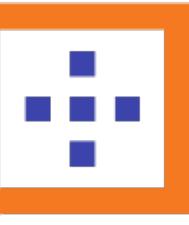


# Challenge #1

## ✓ Decoupled optimizations

- I/O optimizations should be disaggregated from the internal logic
- Moved to a dedicated I/O layer
- Generally applicable
- Portable across different scenarios

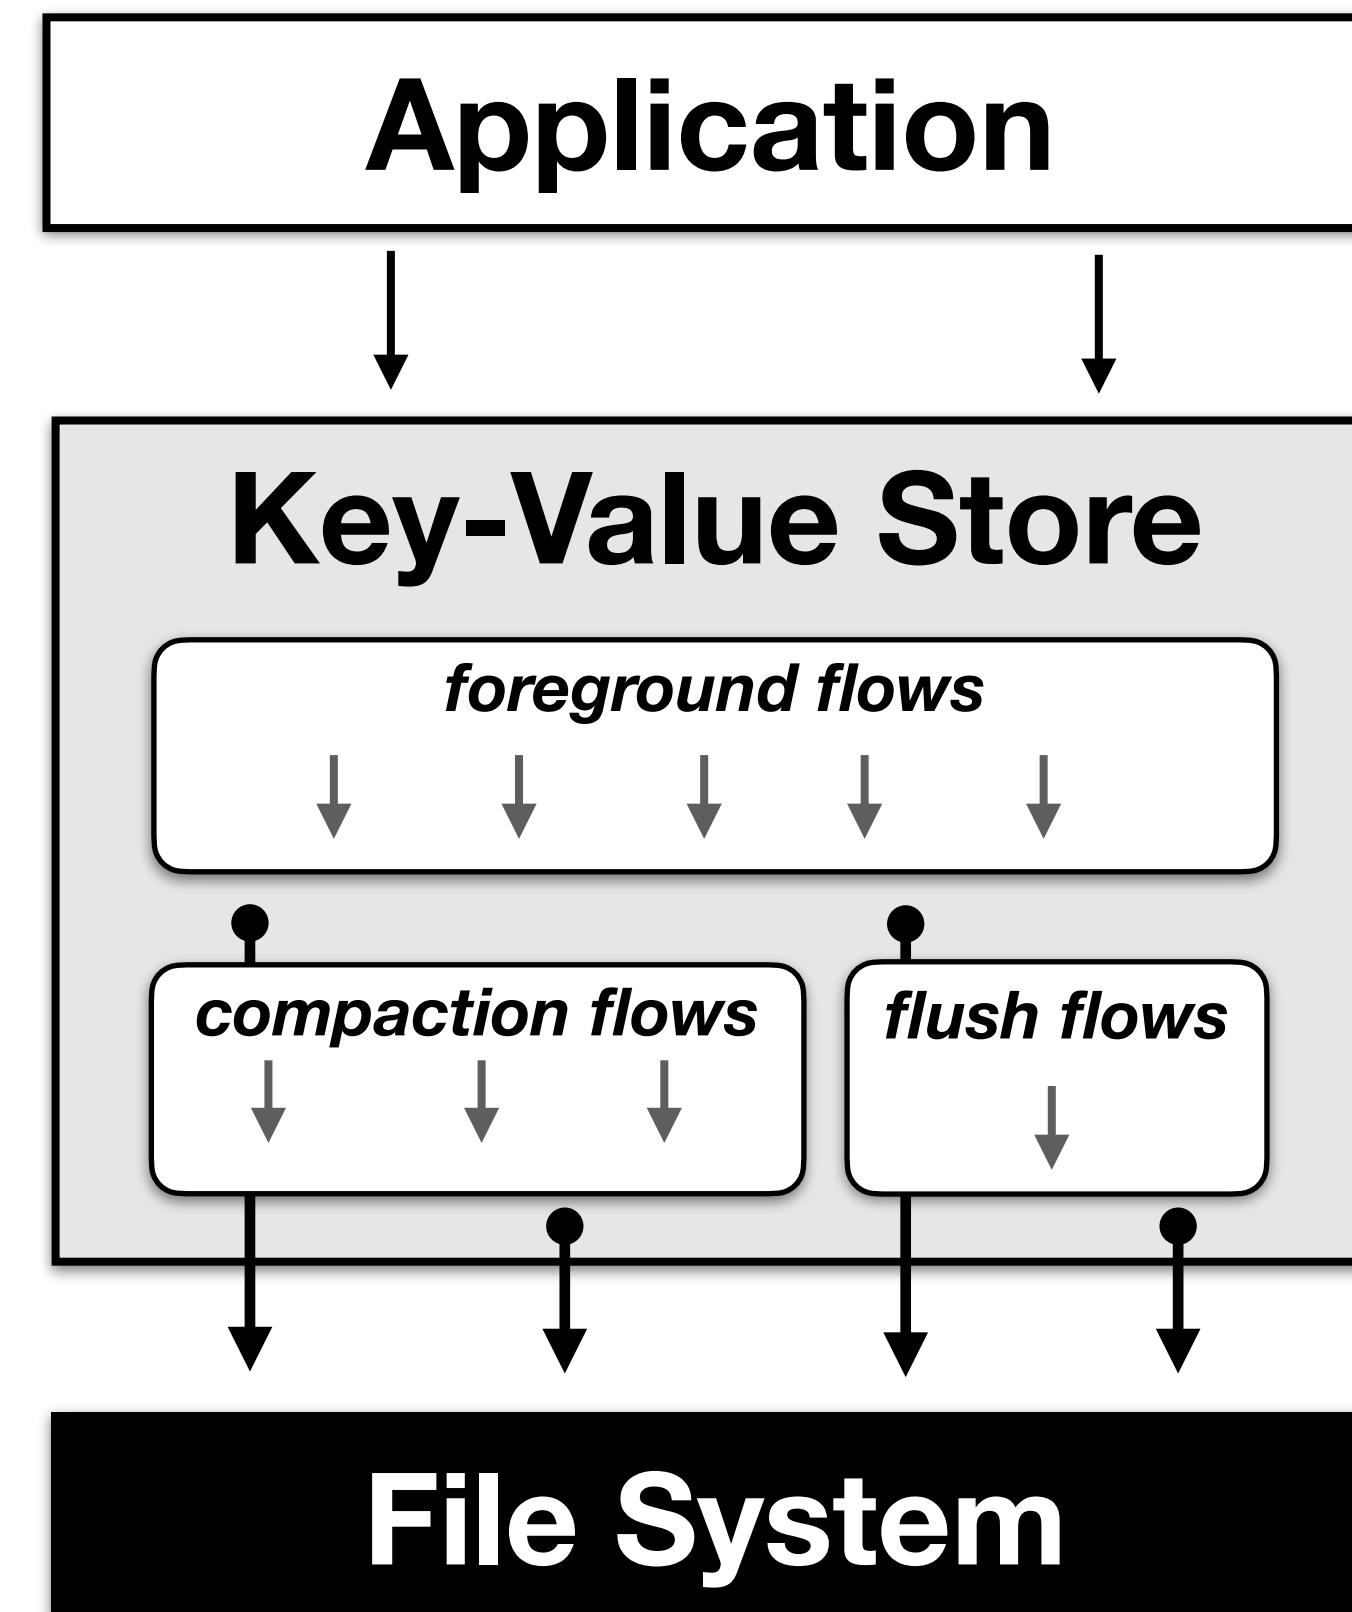


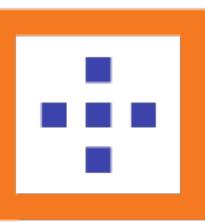


# Challenge #2

## ✖ Rigid interfaces

- Decoupled optimizations lose granularity and internal application knowledge
- I/O layers communicate through rigid interfaces
- Discard information that could be used to classify and differentiate requests

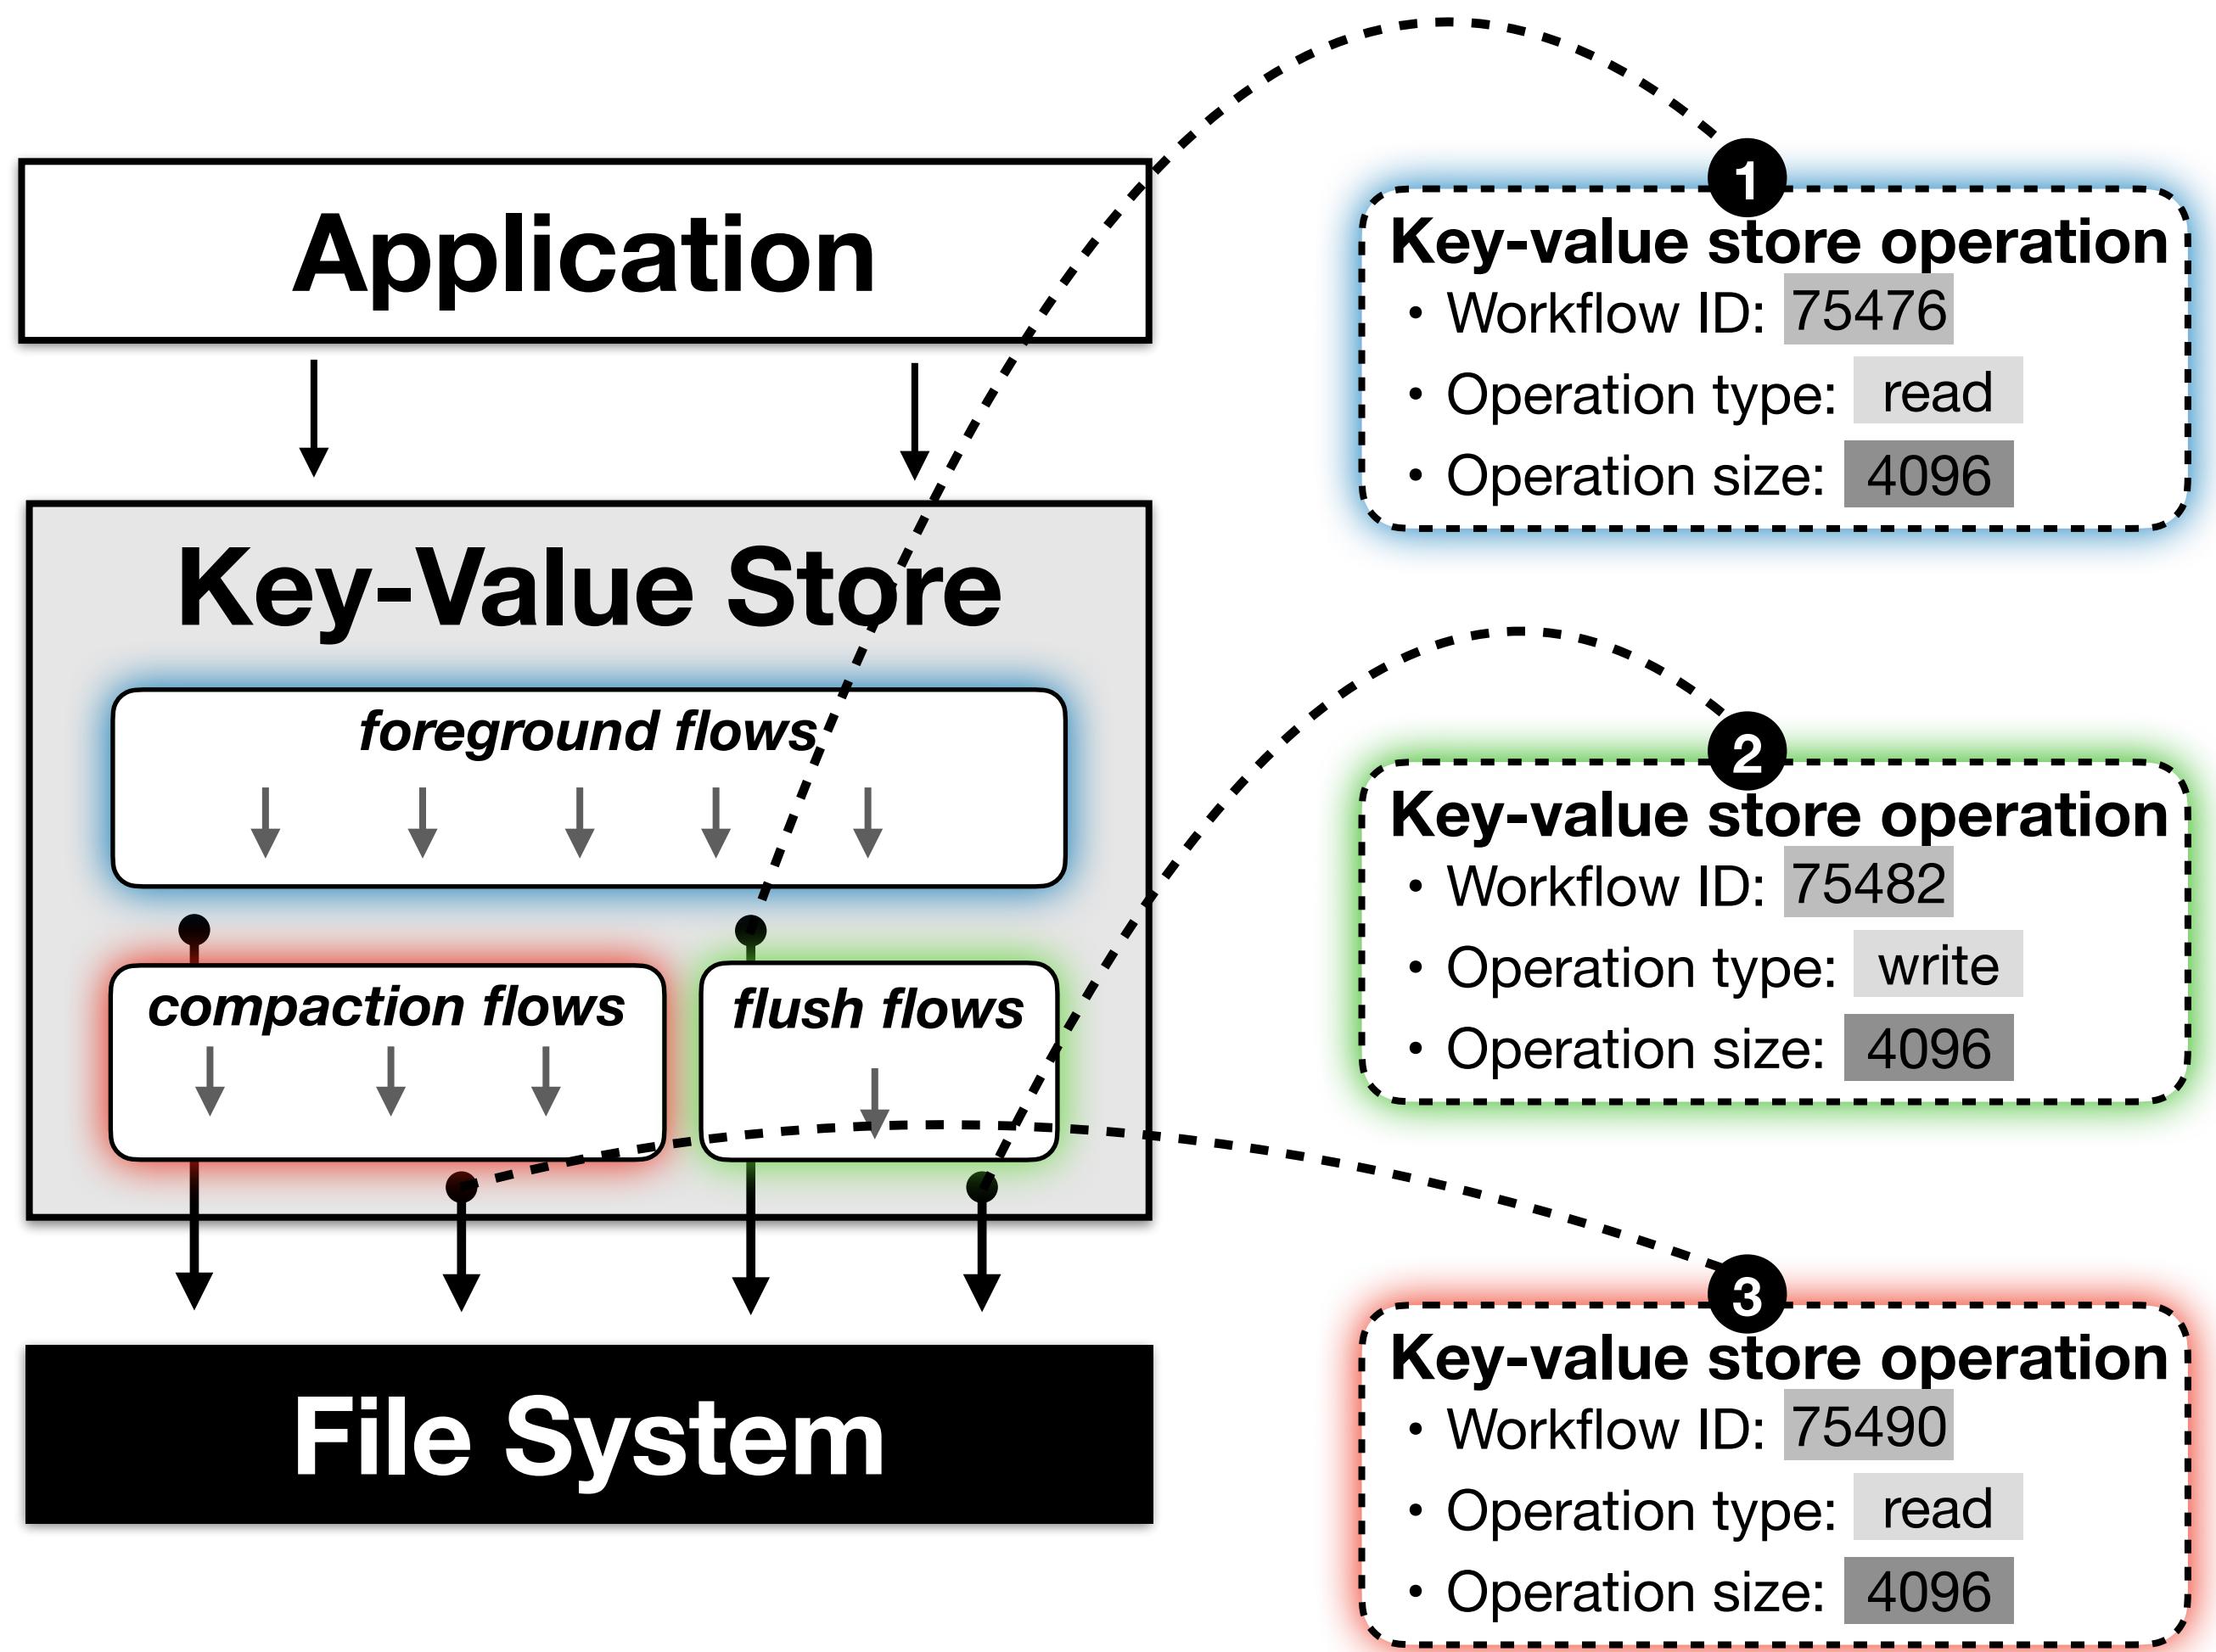




# Challenge #2

## ✖ Rigid interfaces

- Decoupled optimizations lose granularity and internal application knowledge
- I/O layers communicate through rigid interfaces
- Discard information that could be used to classify and differentiate requests



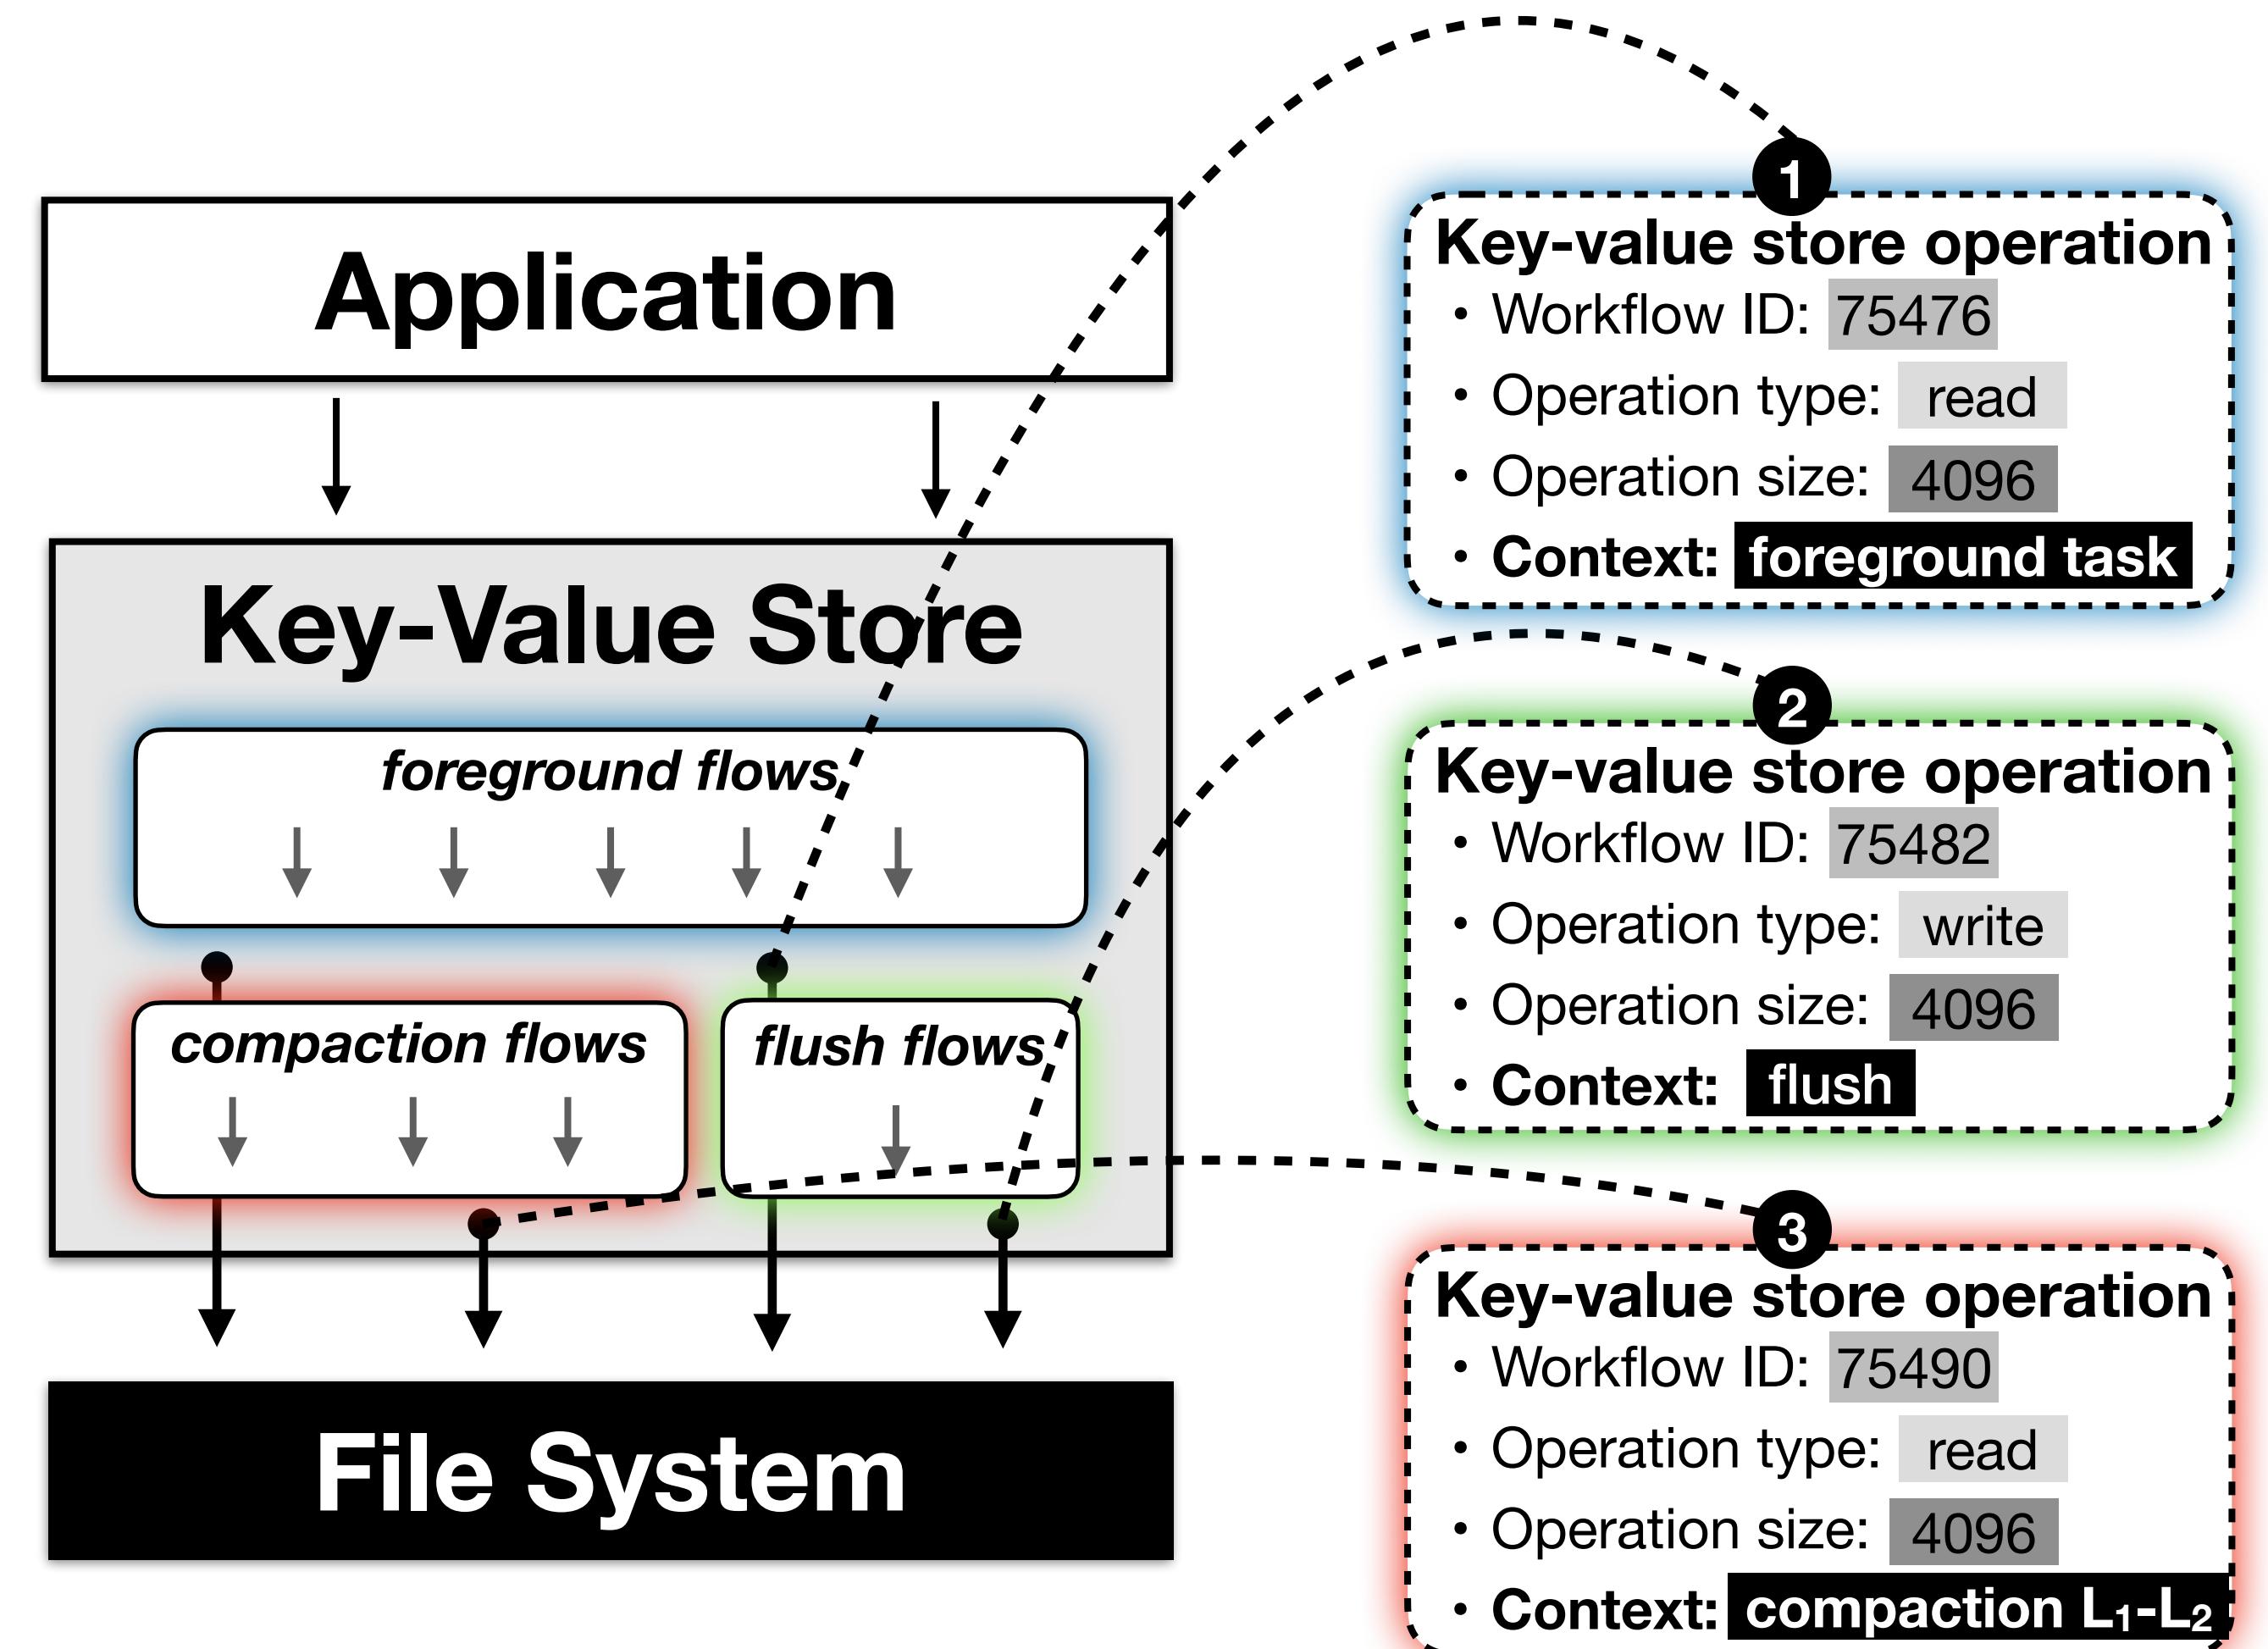


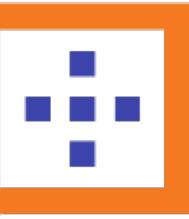
# Challenge #2



## Information propagation

- Application-level information must be propagated throughout layers
- Decoupled optimizations can provide the same level of control and performance

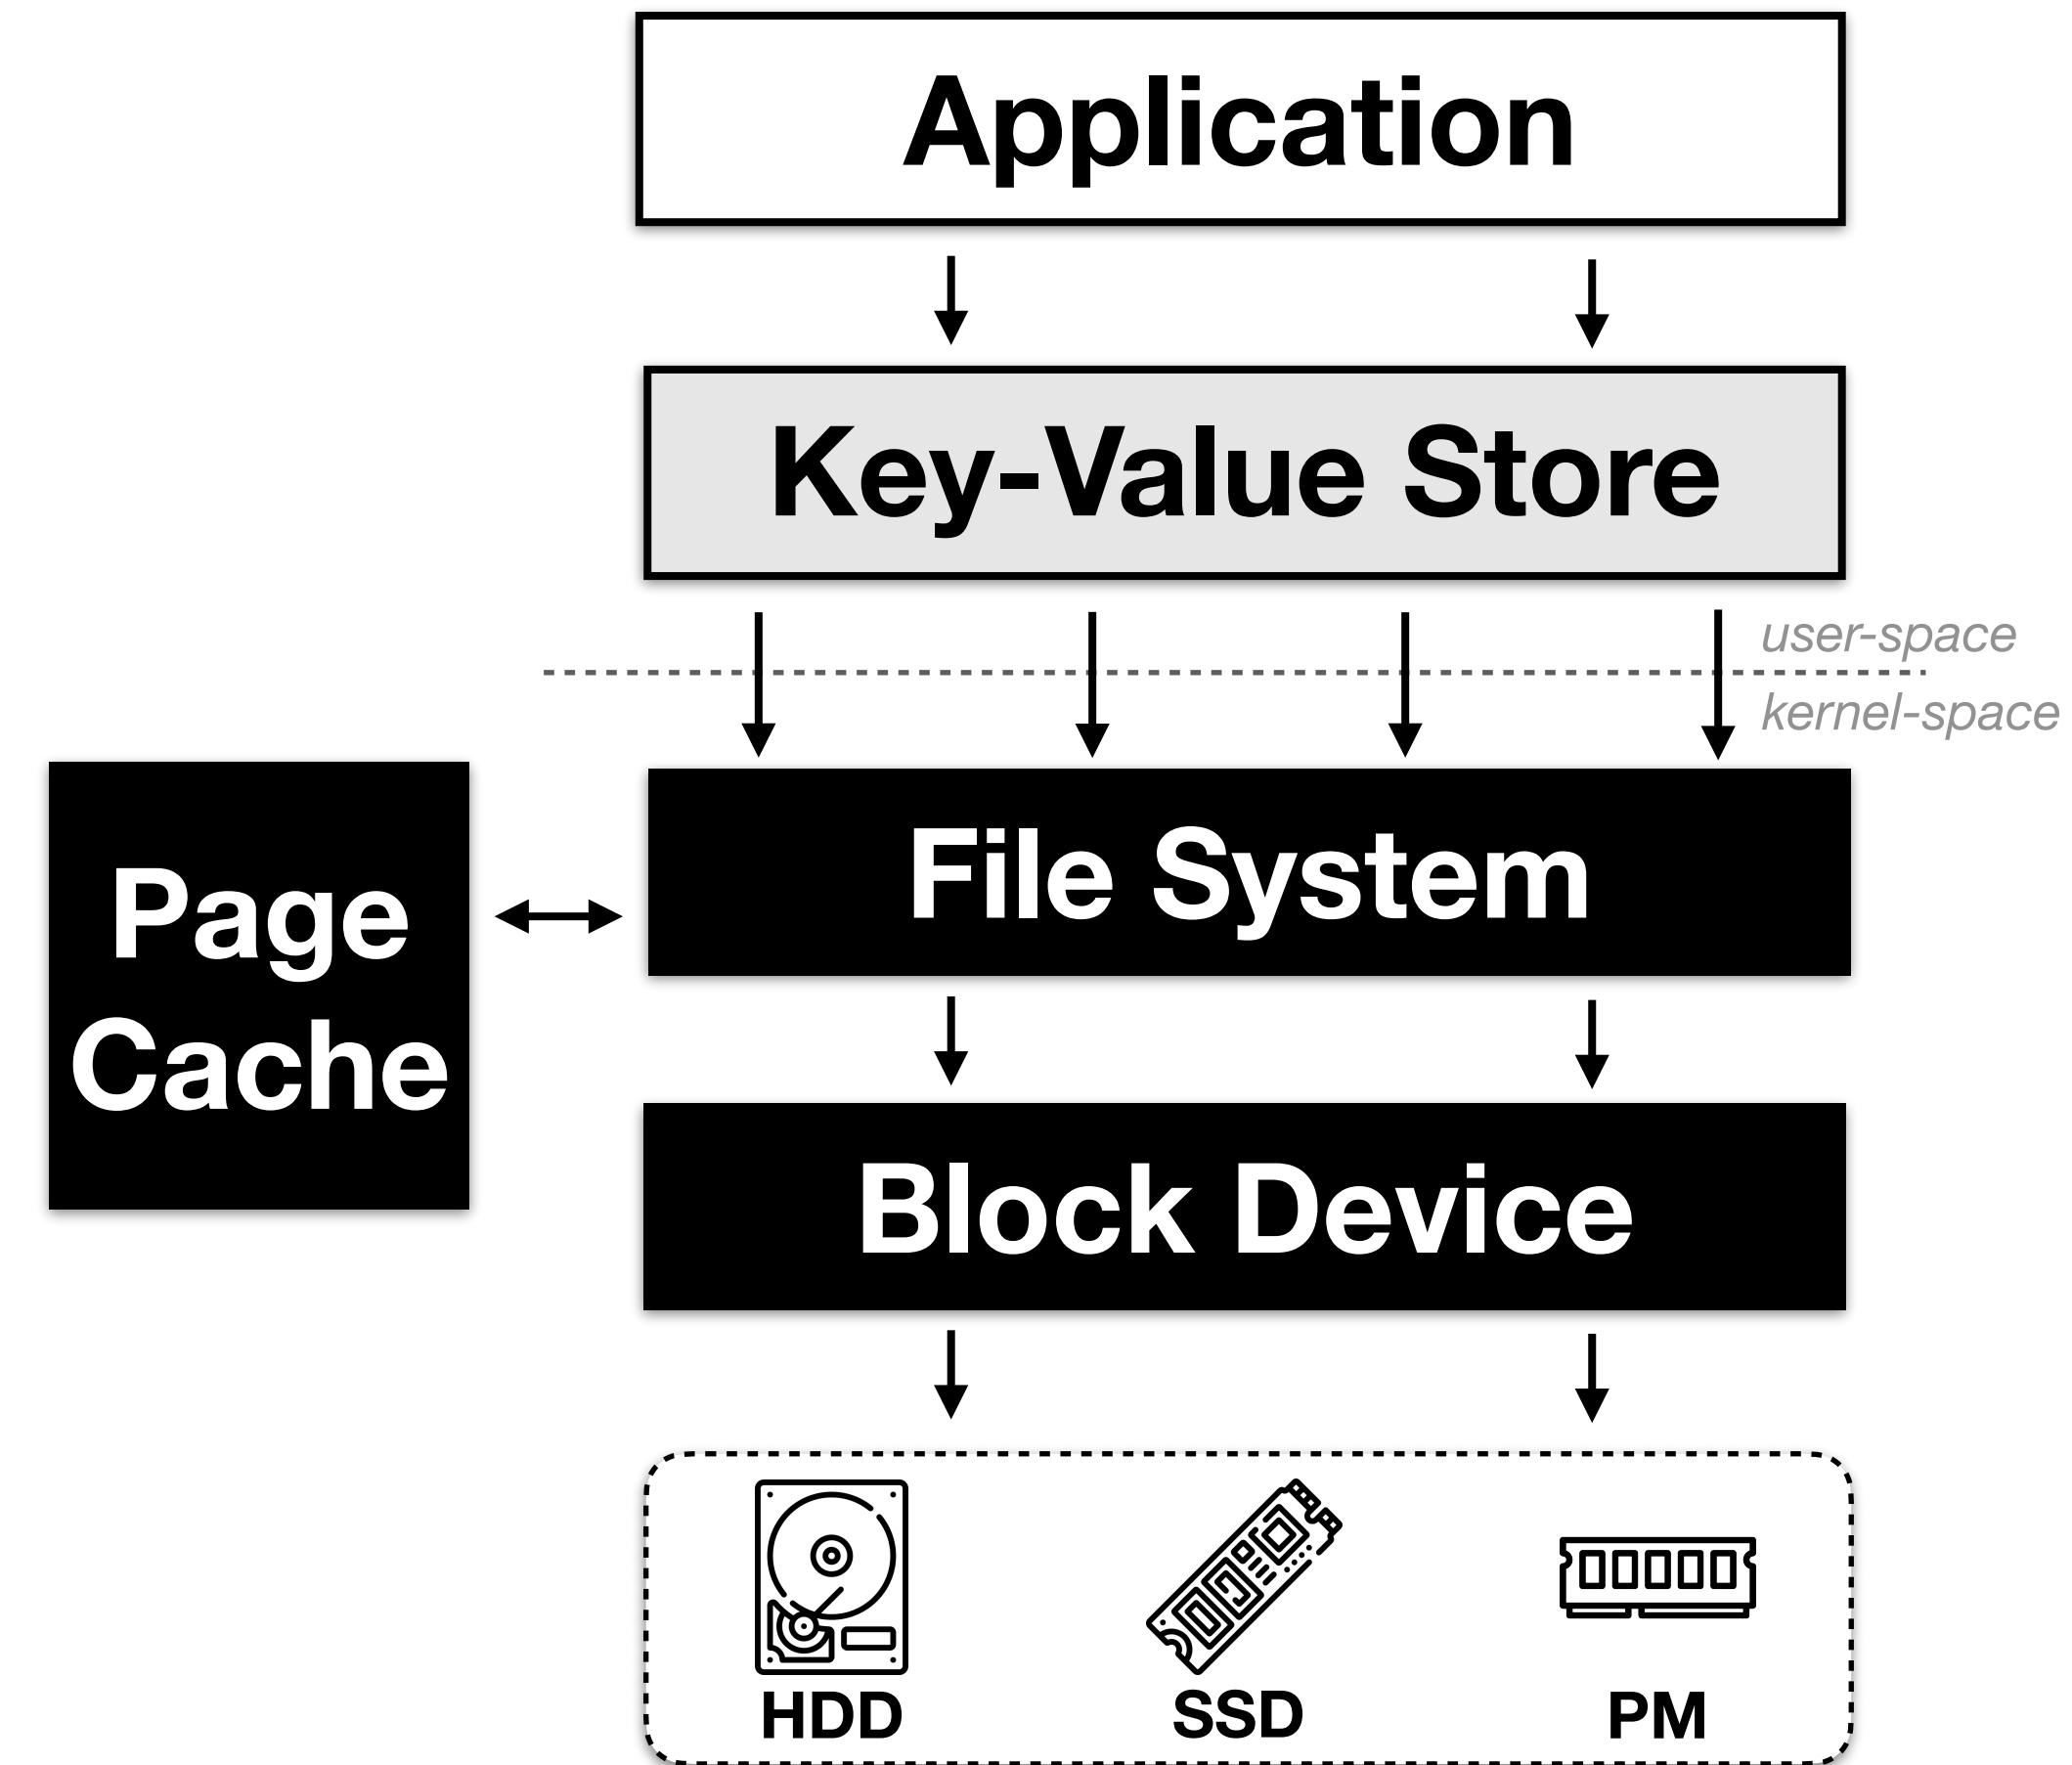


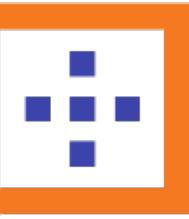


# Challenge #3

## ✗ Kernel-level layers

- Propagating context to kernel requires breaking user-to-kernel and kernel-internal APIs
- Kernel-level development is more restricted and error-prone
- Optimizations would be ineffective under kernel-bypass storage stacks (e.g., SPDK, PMDK)

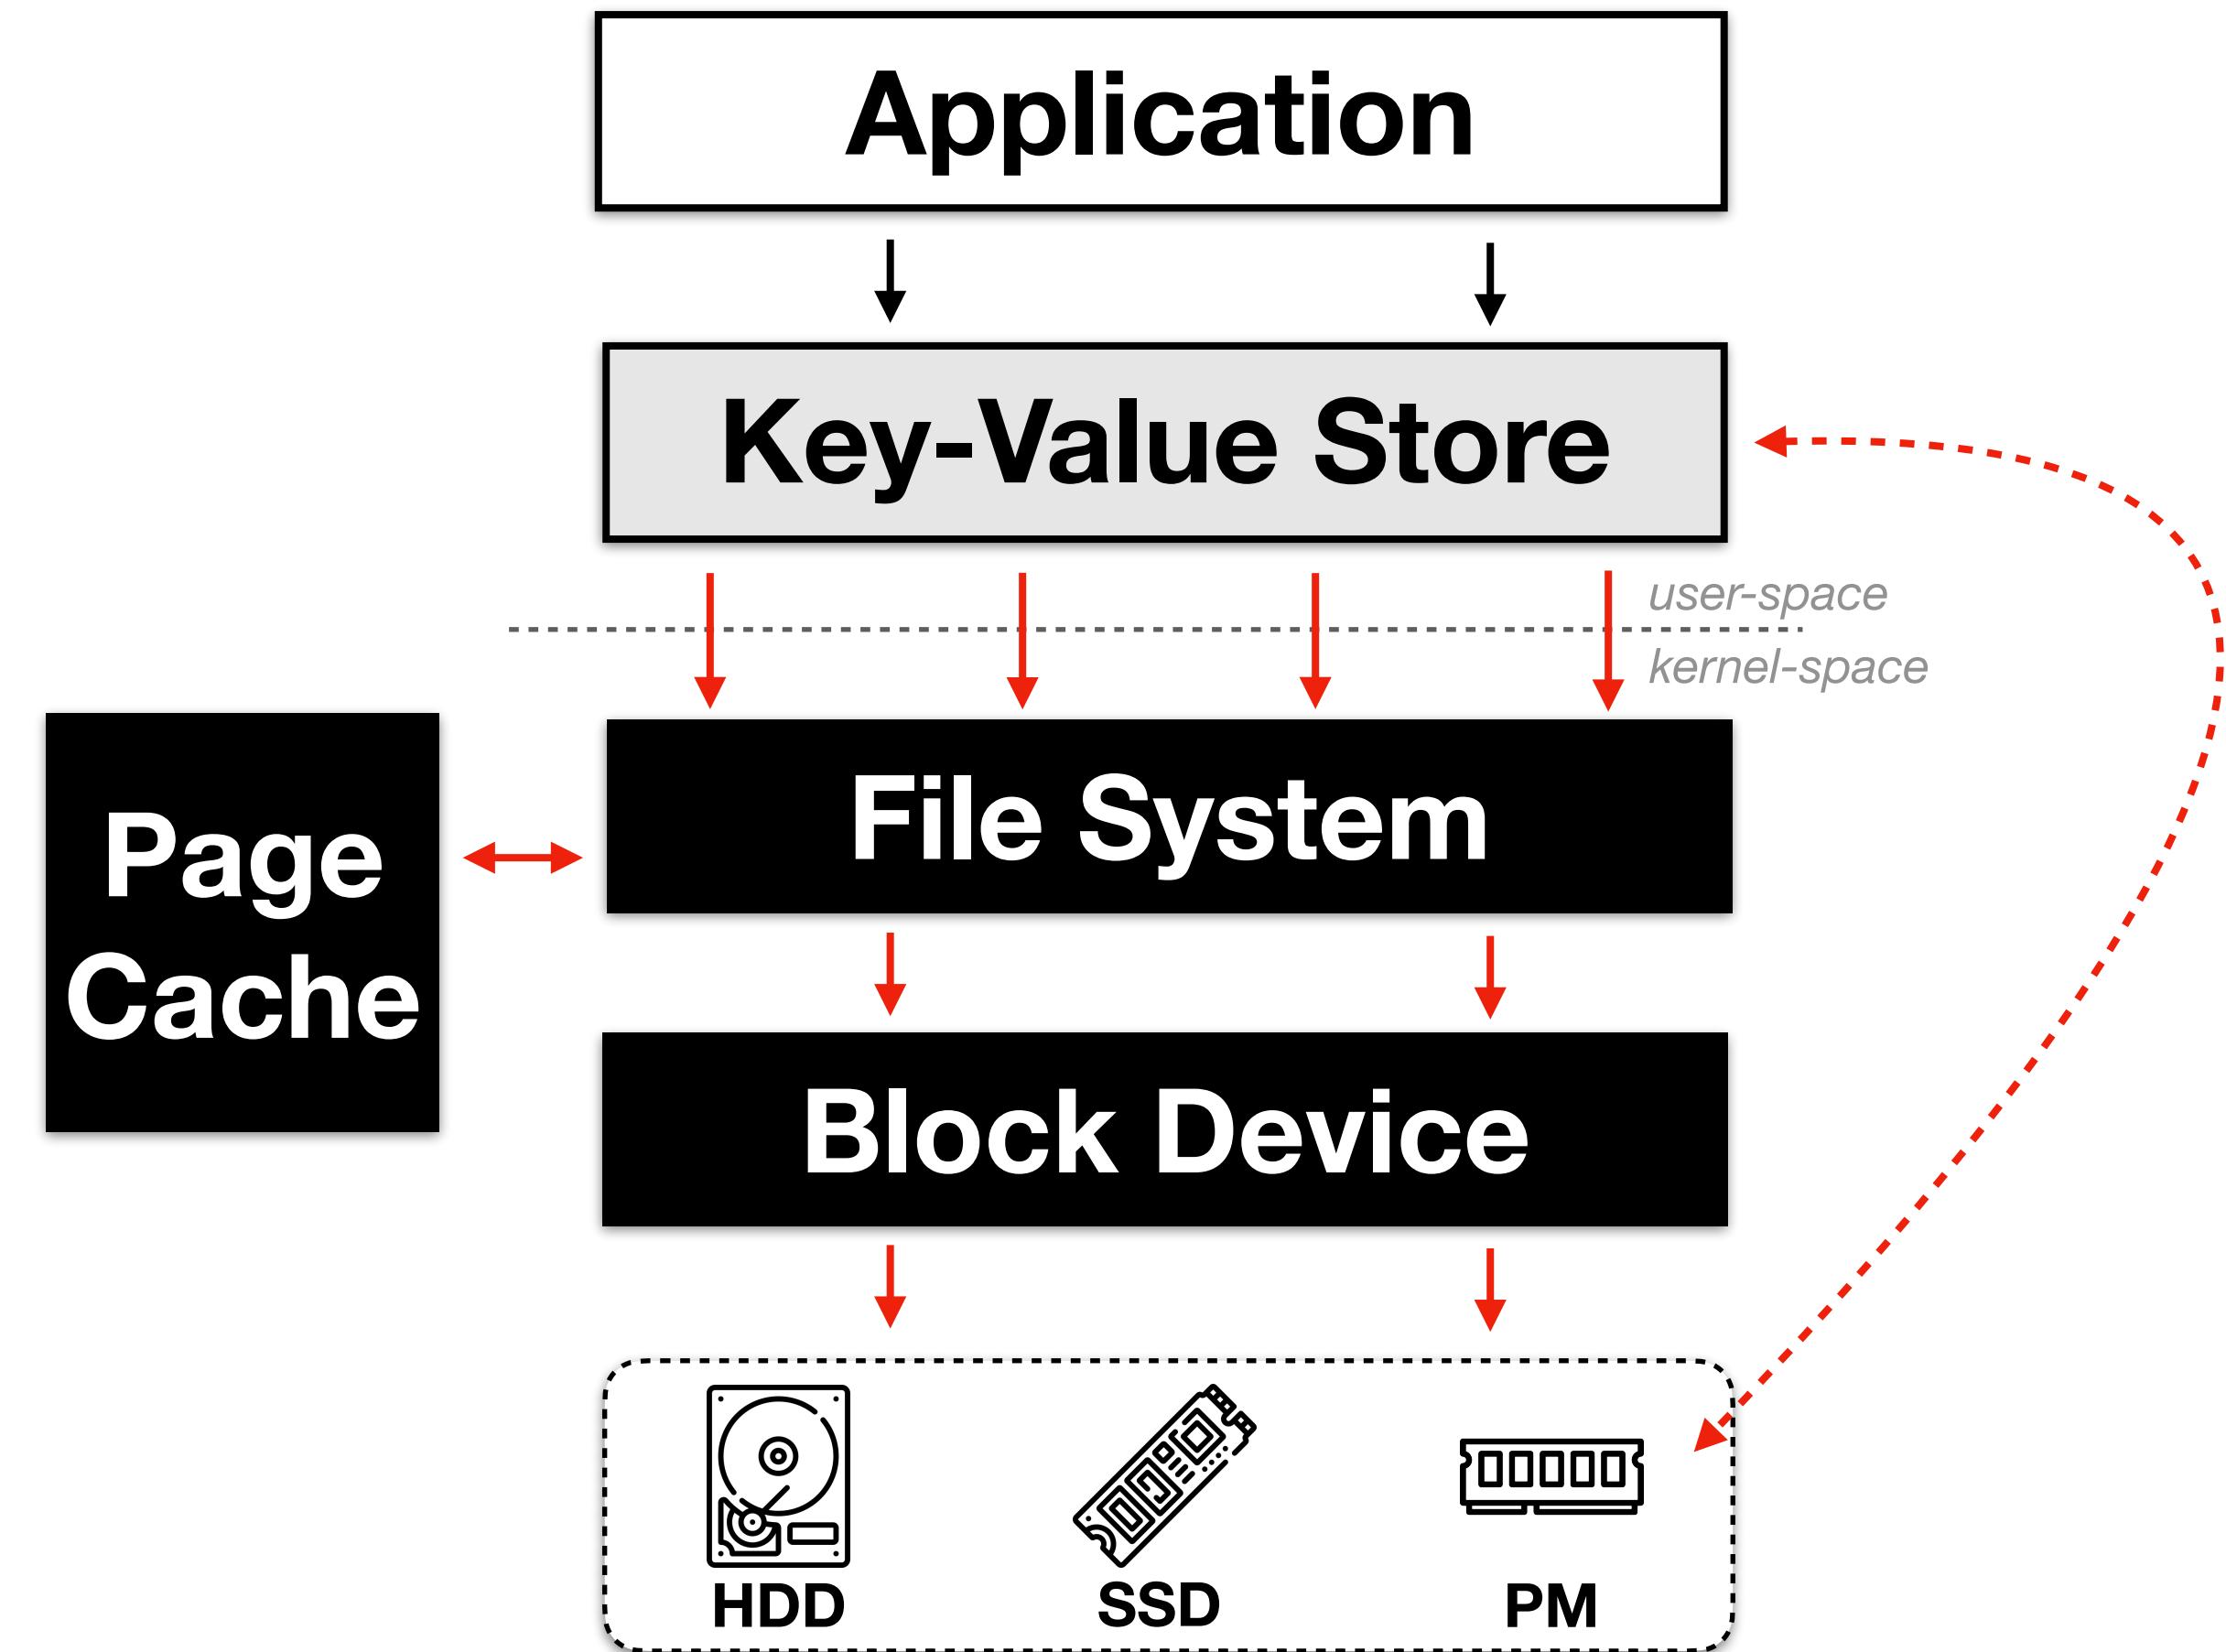


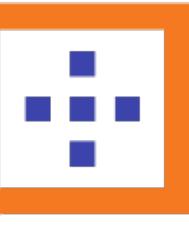


# Challenge #3

## ✗ Kernel-level layers

- Propagating context to kernel requires breaking user-to-kernel and kernel-internal APIs
- Kernel-level development is more restricted and error-prone
- Optimizations would be ineffective under kernel-bypass storage stacks (e.g., SPDK, PMDK)

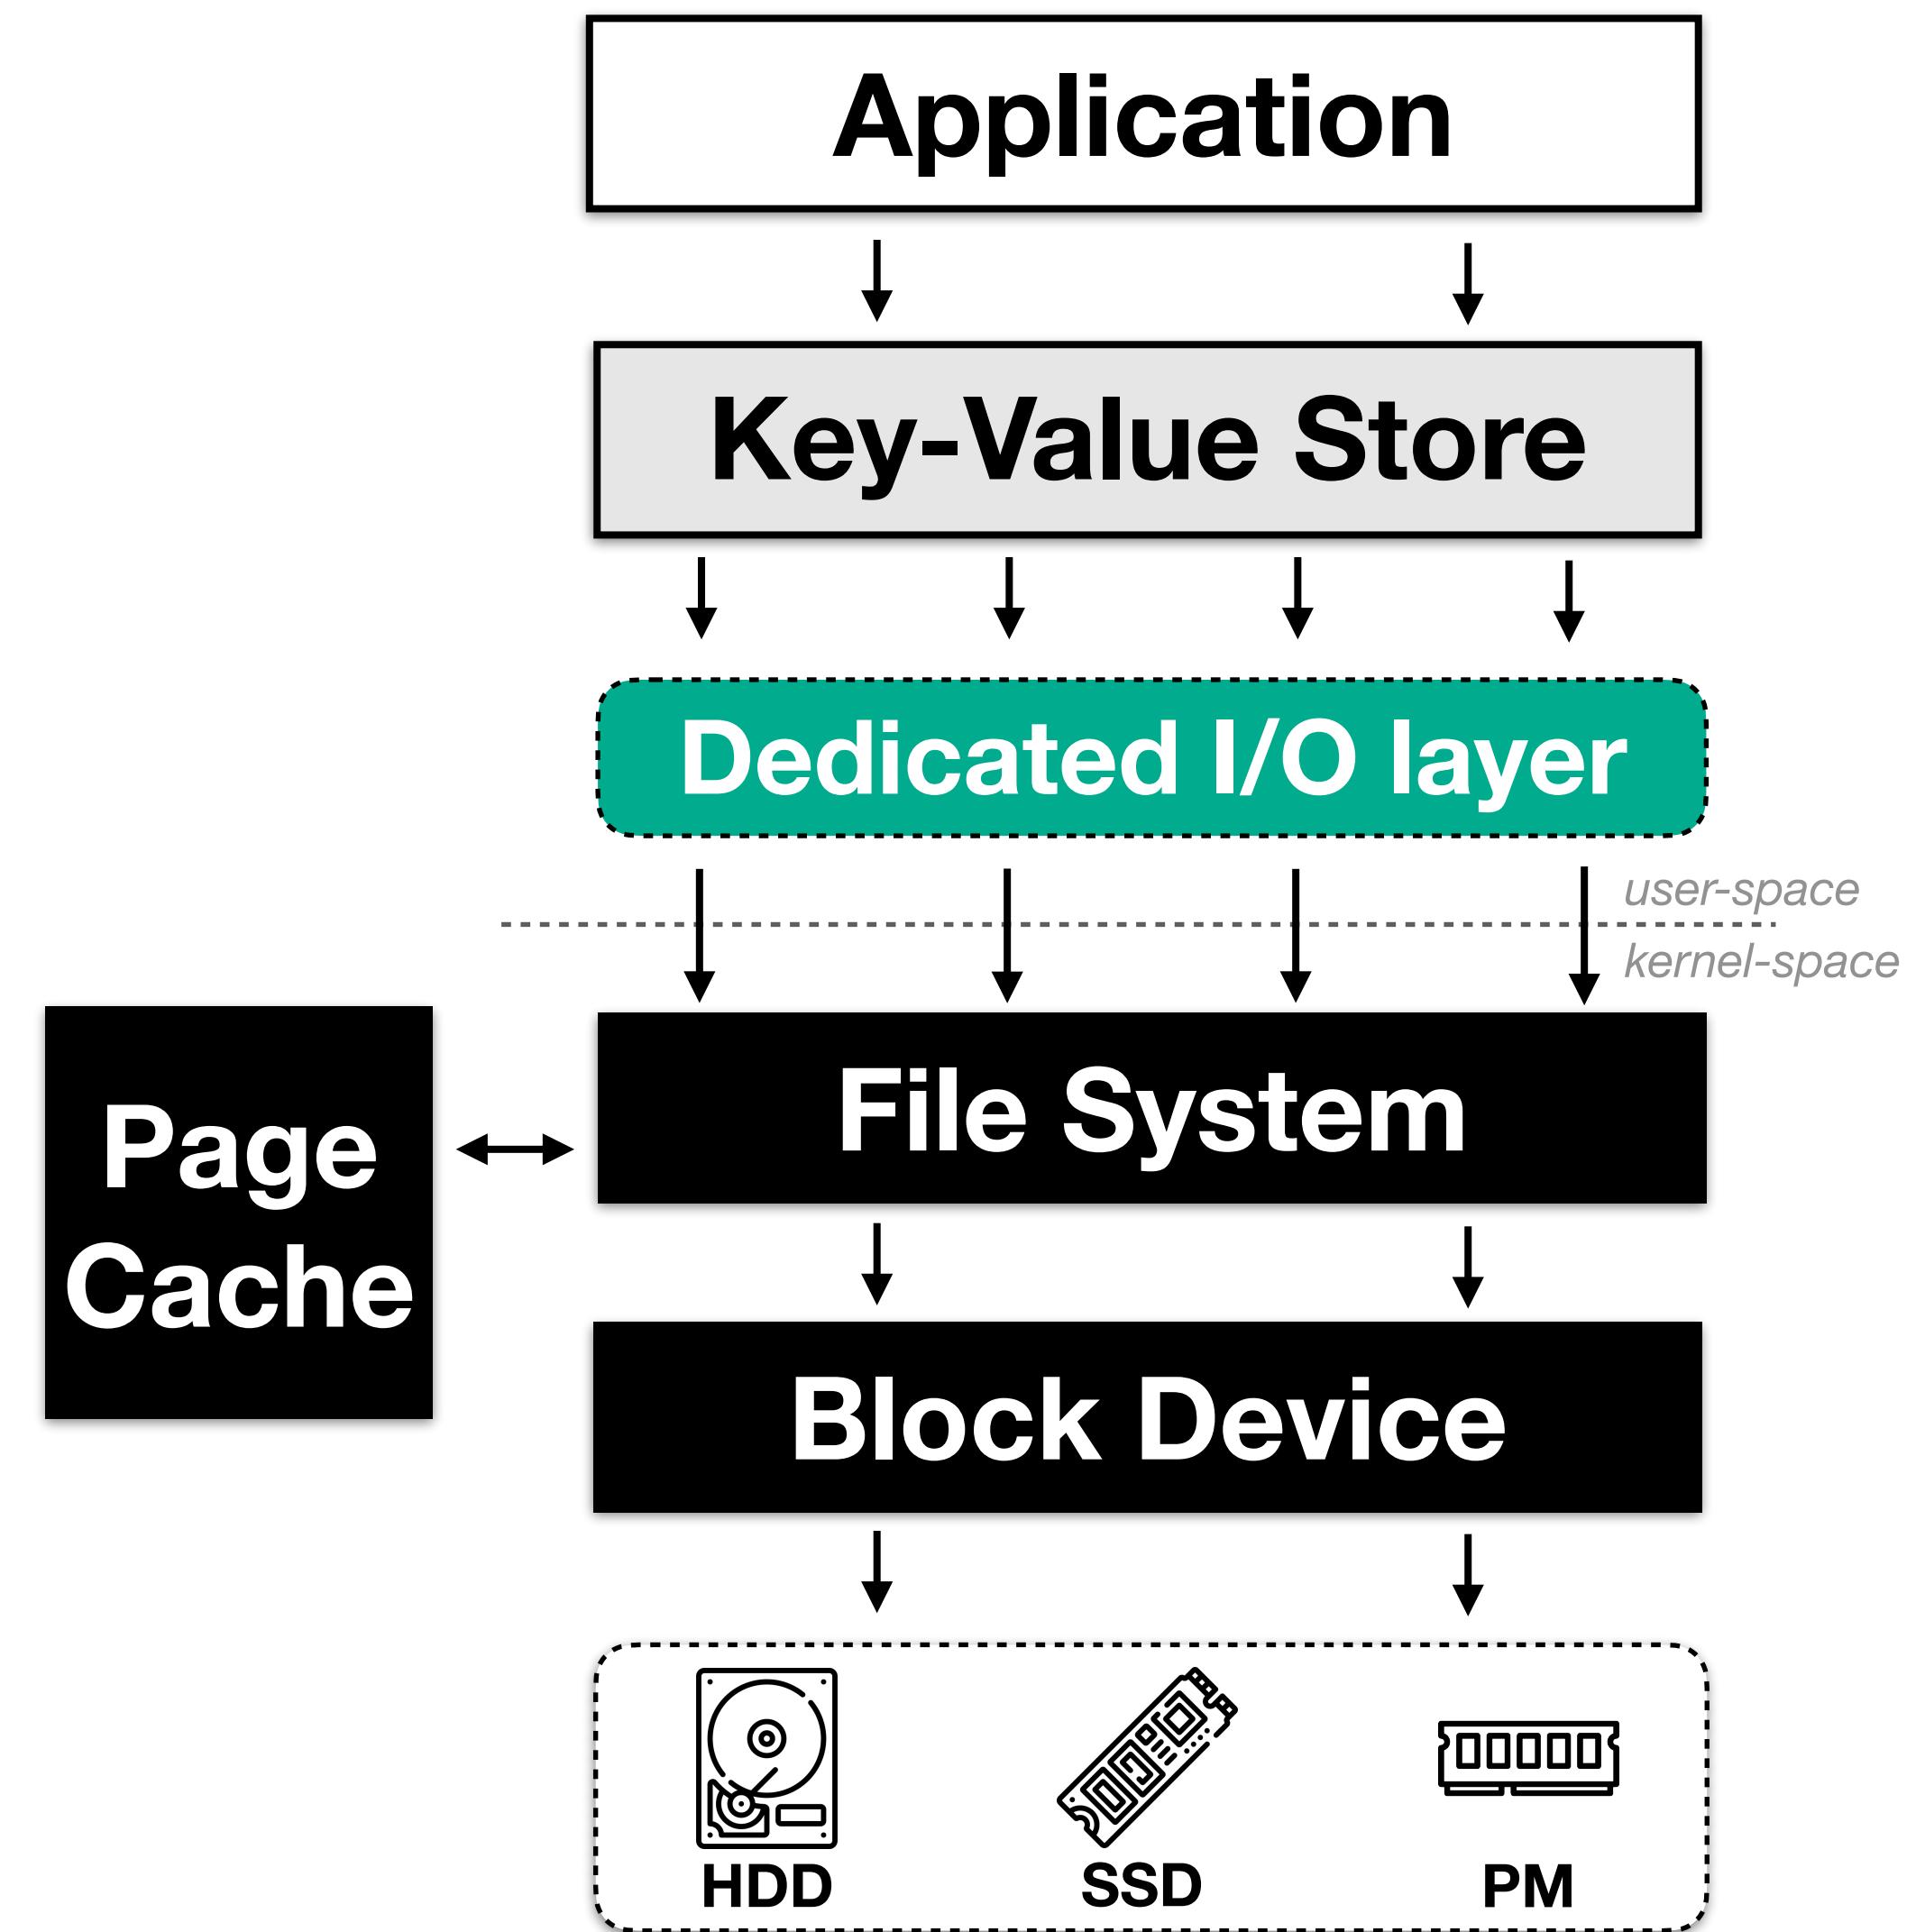




# Challenge #3

## ✓ Actuate at user-level

- Optimizations should be implemented at a dedicated user-level layer
- Promote portability across different systems and layers
- Ease information propagation throughout I/O layers

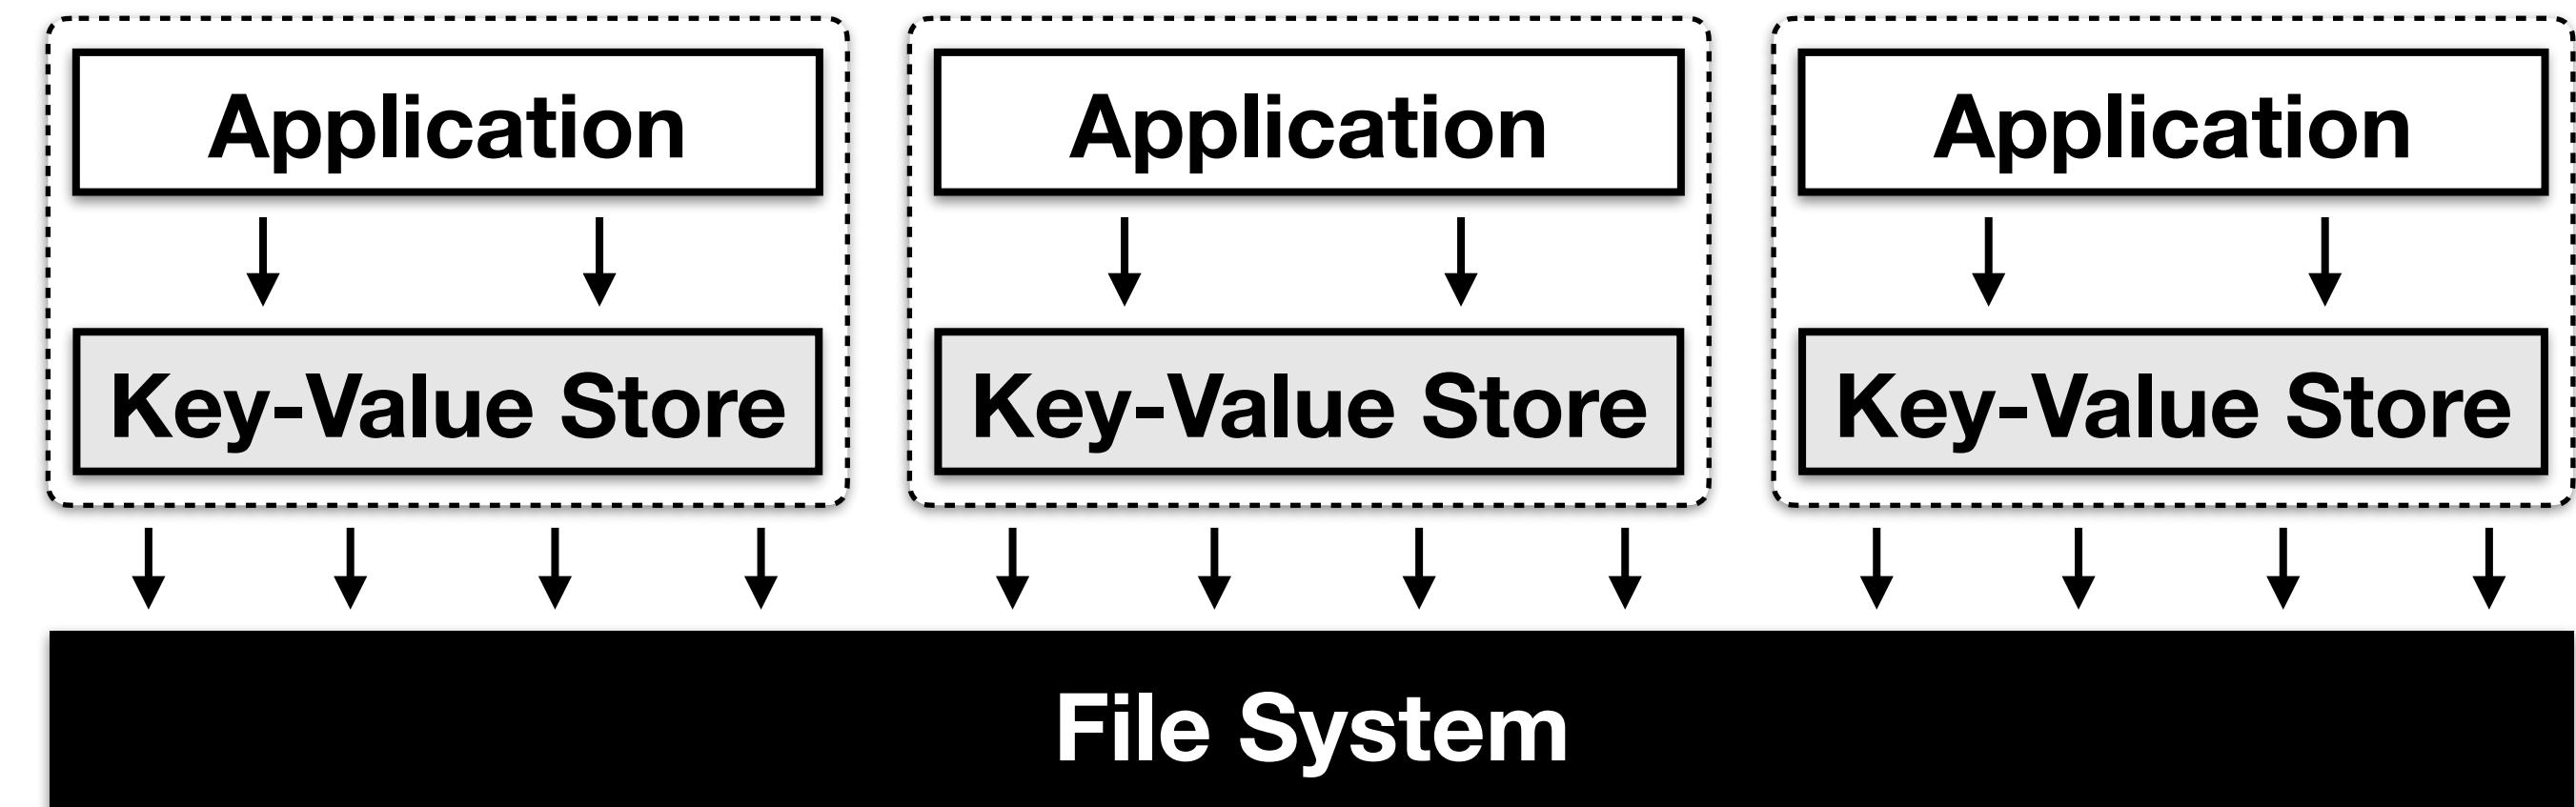




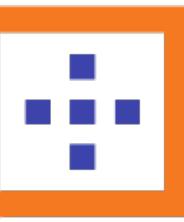
# Challenge #4

## ✖ Partial visibility

- Optimizations are oblivious of other systems
- Lack of coordination
- Conflicting optimizations, I/O contention, and performance variation



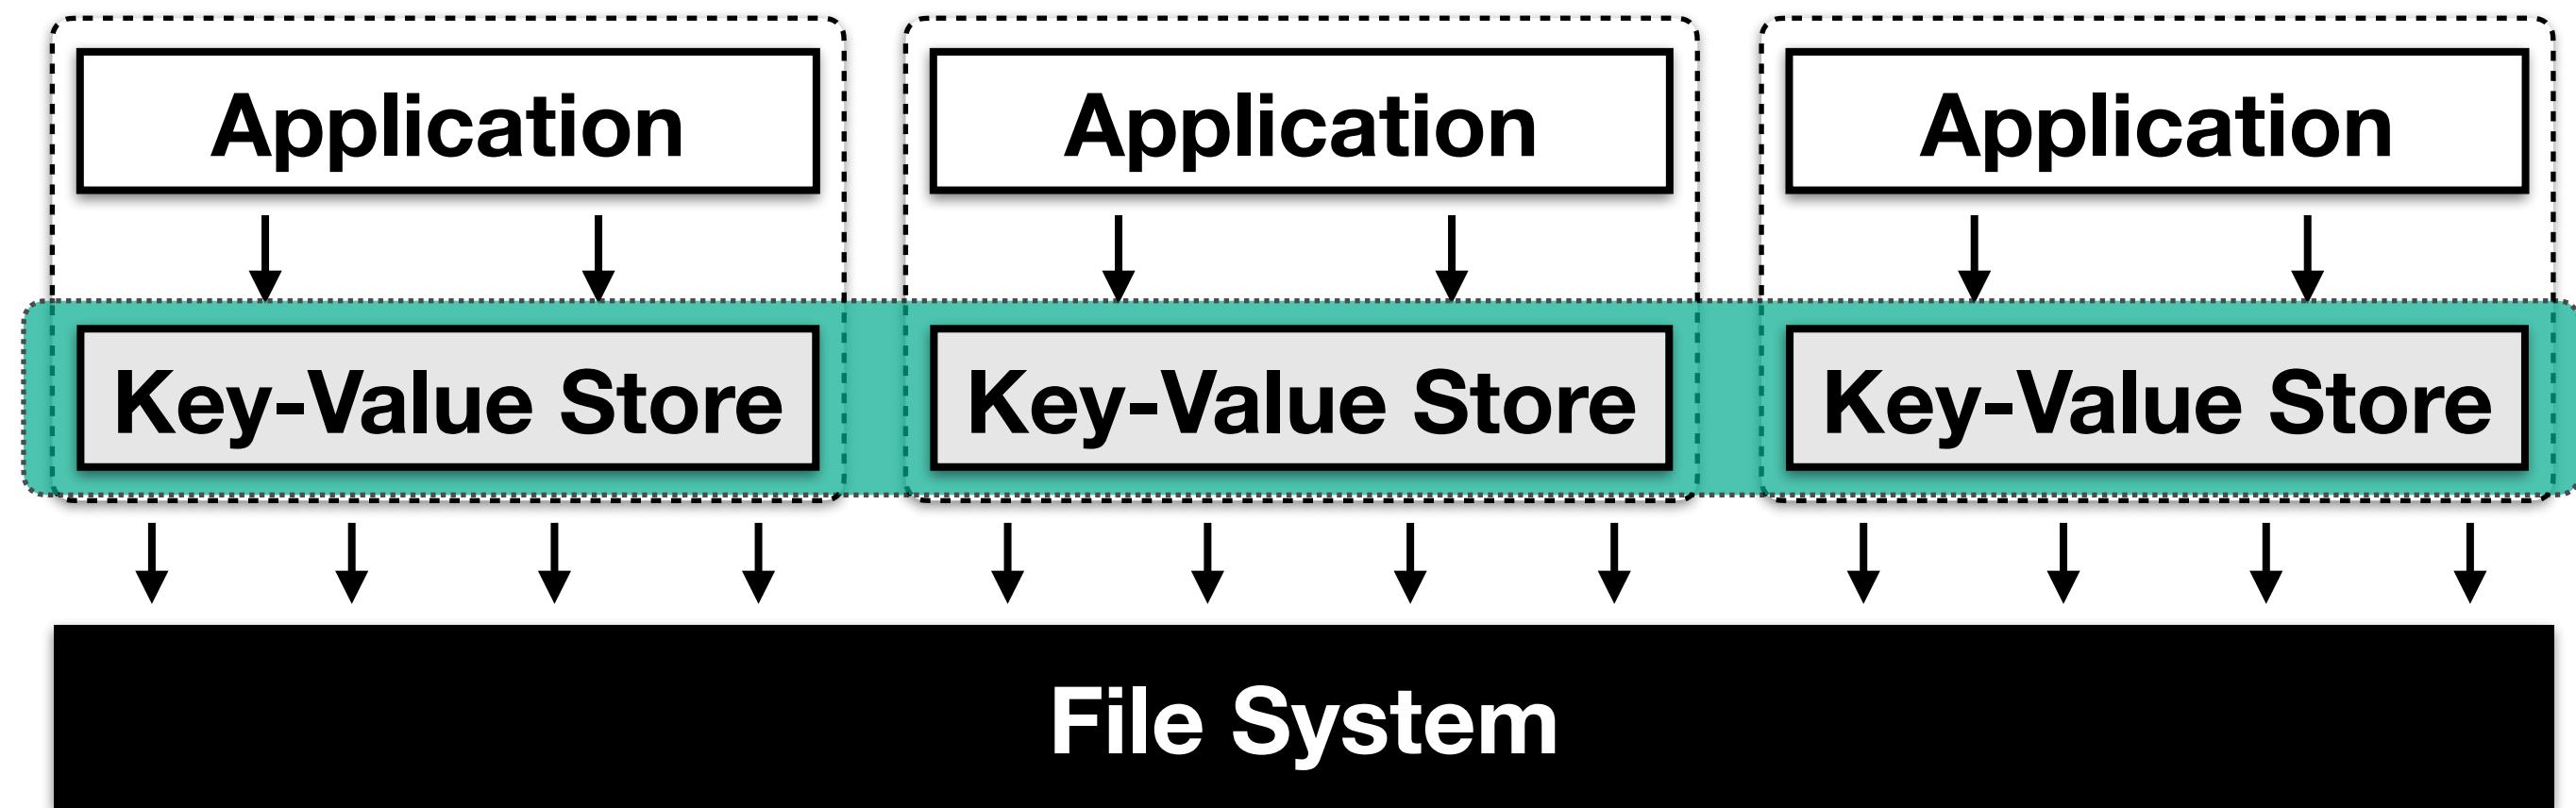
**Note:** the storage backend can either be local (e.g., ext4, xfs) or distributed (e.g., Lustre, GPFS)

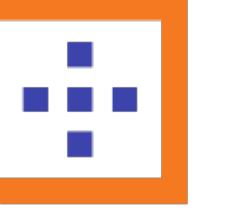


# Challenge #4

## ✓ Global I/O control

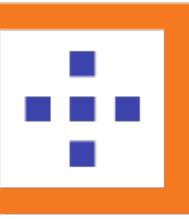
- Optimizations should be aware of the surrounding system stack
- Operate in coordination
- Holistic control of I/O workflows and shared resources





# Part 2

## designing a storage data plane framework



# PAIO

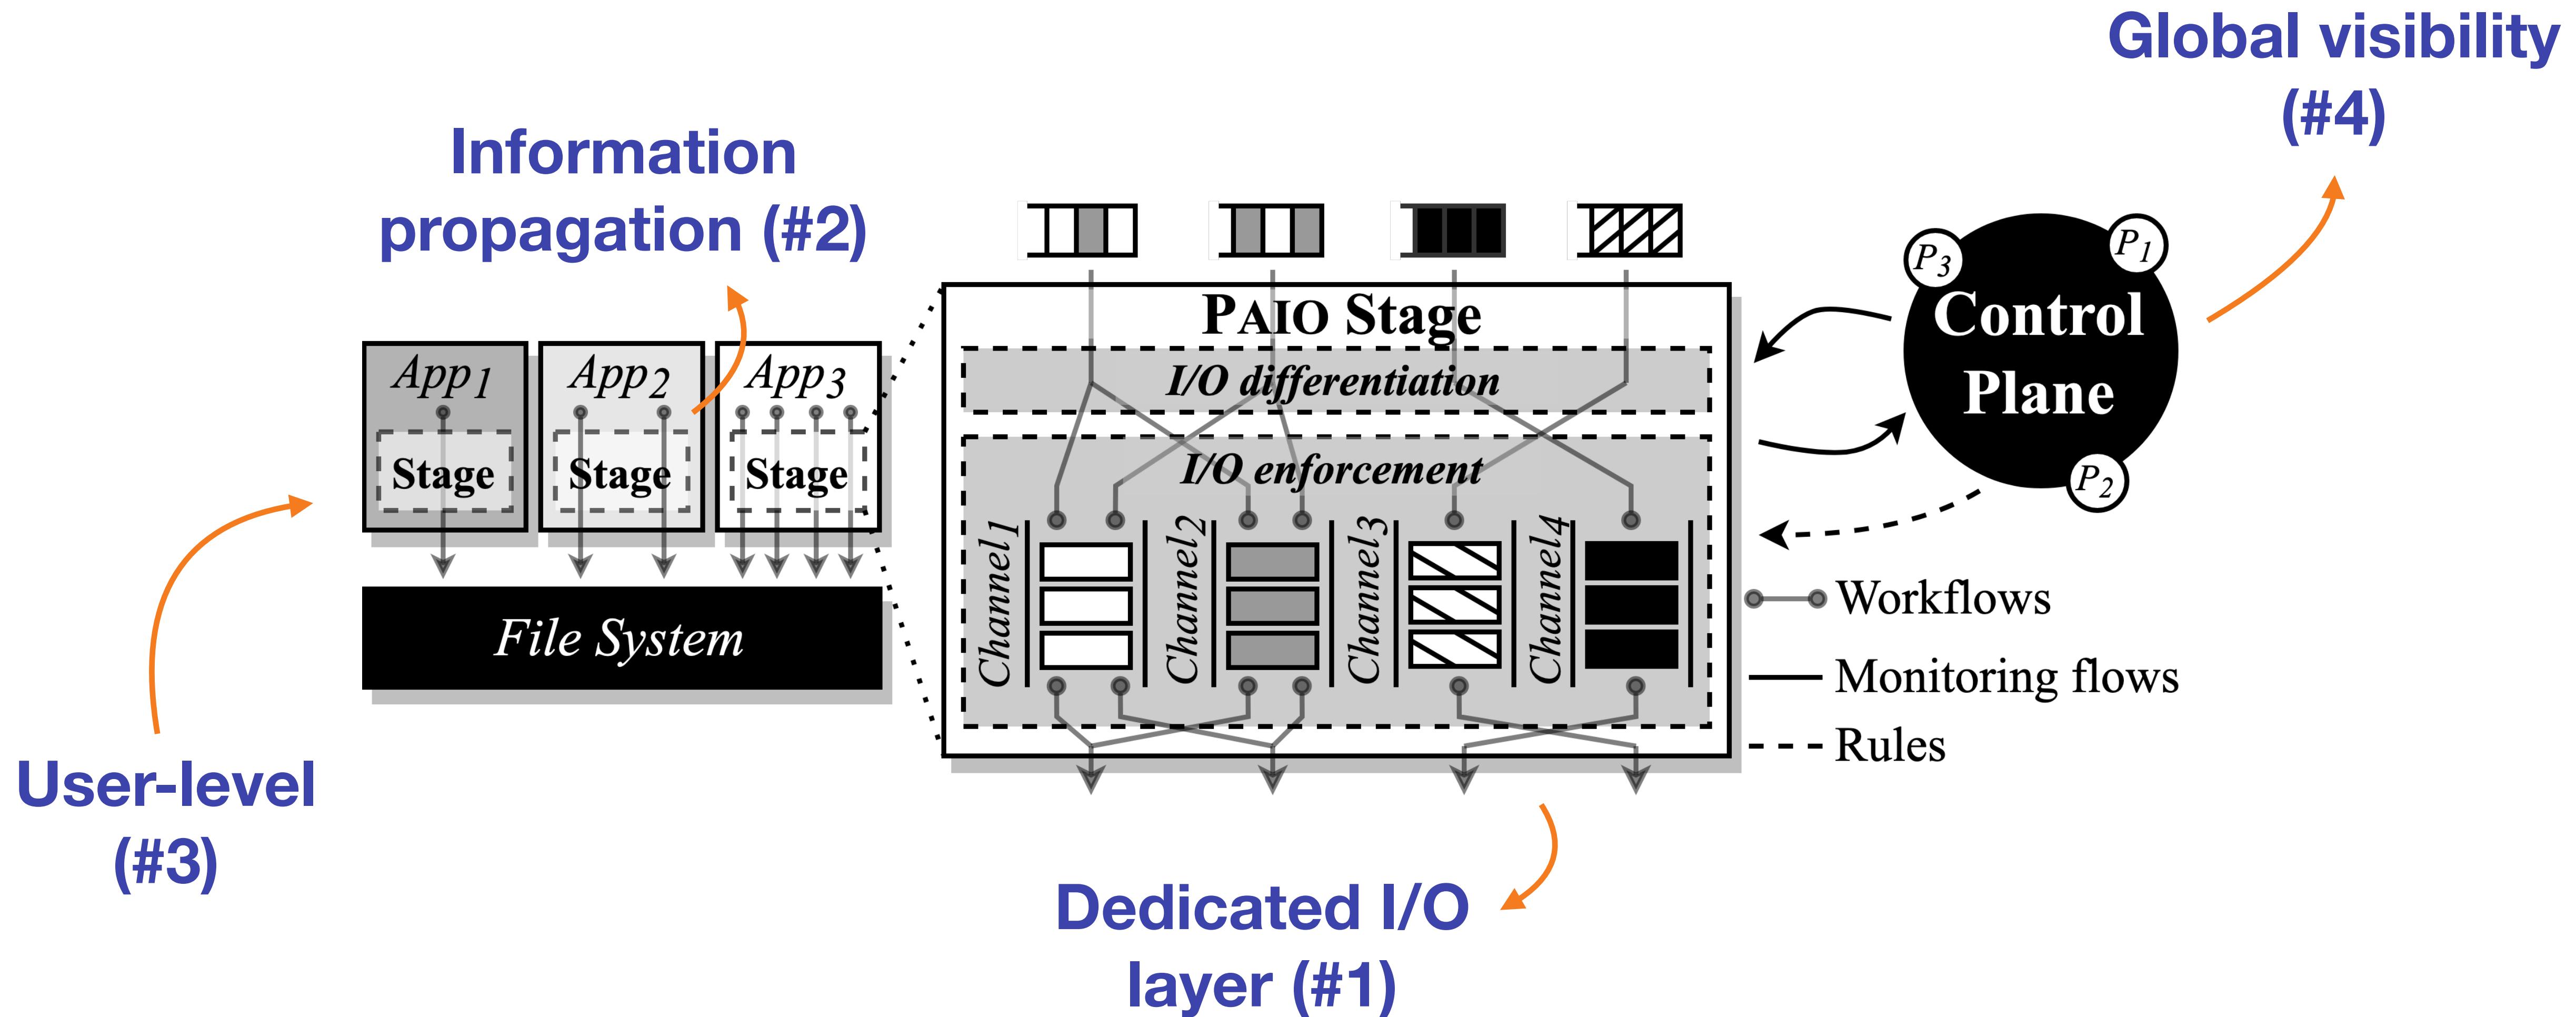
- **User-level** framework for building **portable** and **generally applicable** optimizations
- Adopts ideas from **Software-Defined Storage** [6]
  - I/O optimizations are implemented **outside** applications as **data plane stages**
  - **Stages** are controlled through a **control plane** for coordinated access to resources
- Enables the propagation of application-level information through **context propagation**
- Porting I/O layers to use PAIO requires **none to minor** code changes

[5] “PAIO: General, Portable I/O Optimizations with Minor Application Modifications”. Macedo et al. USENIX FAST 2022.

[6] “A Survey and Classification of Software-Defined Storage Systems”. Macedo et al. ACM CSUR 2020.

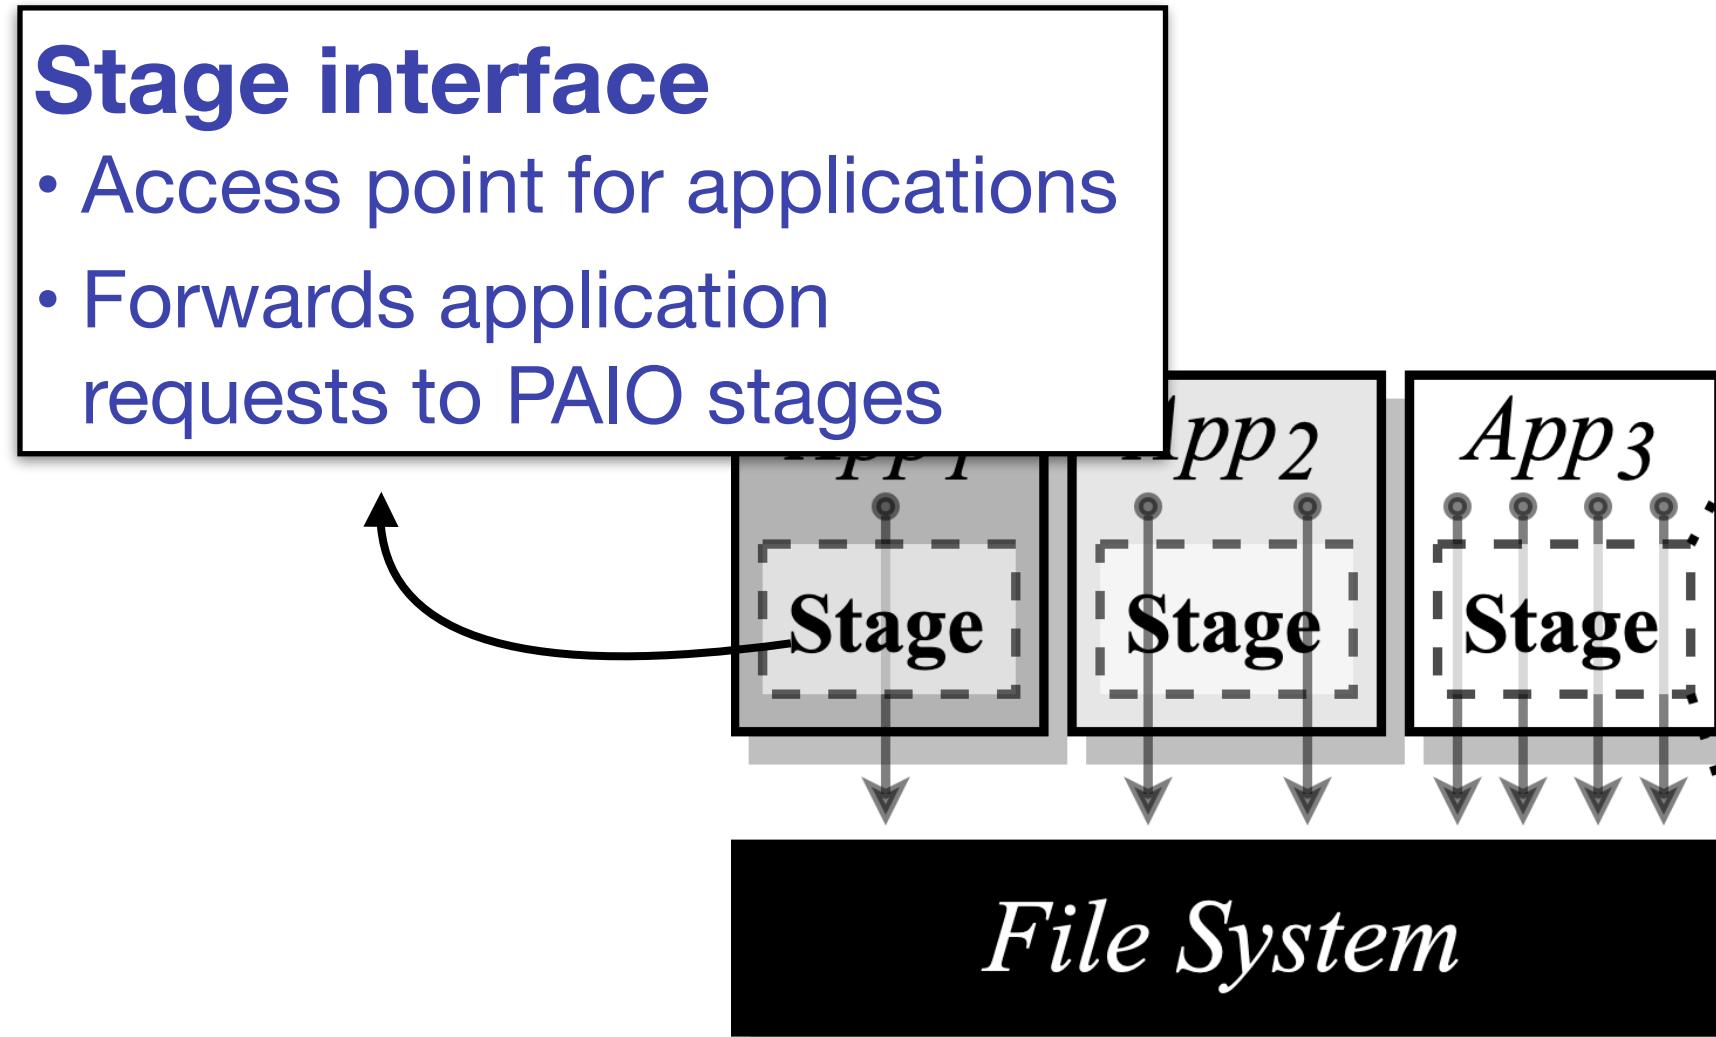


# PAIO design



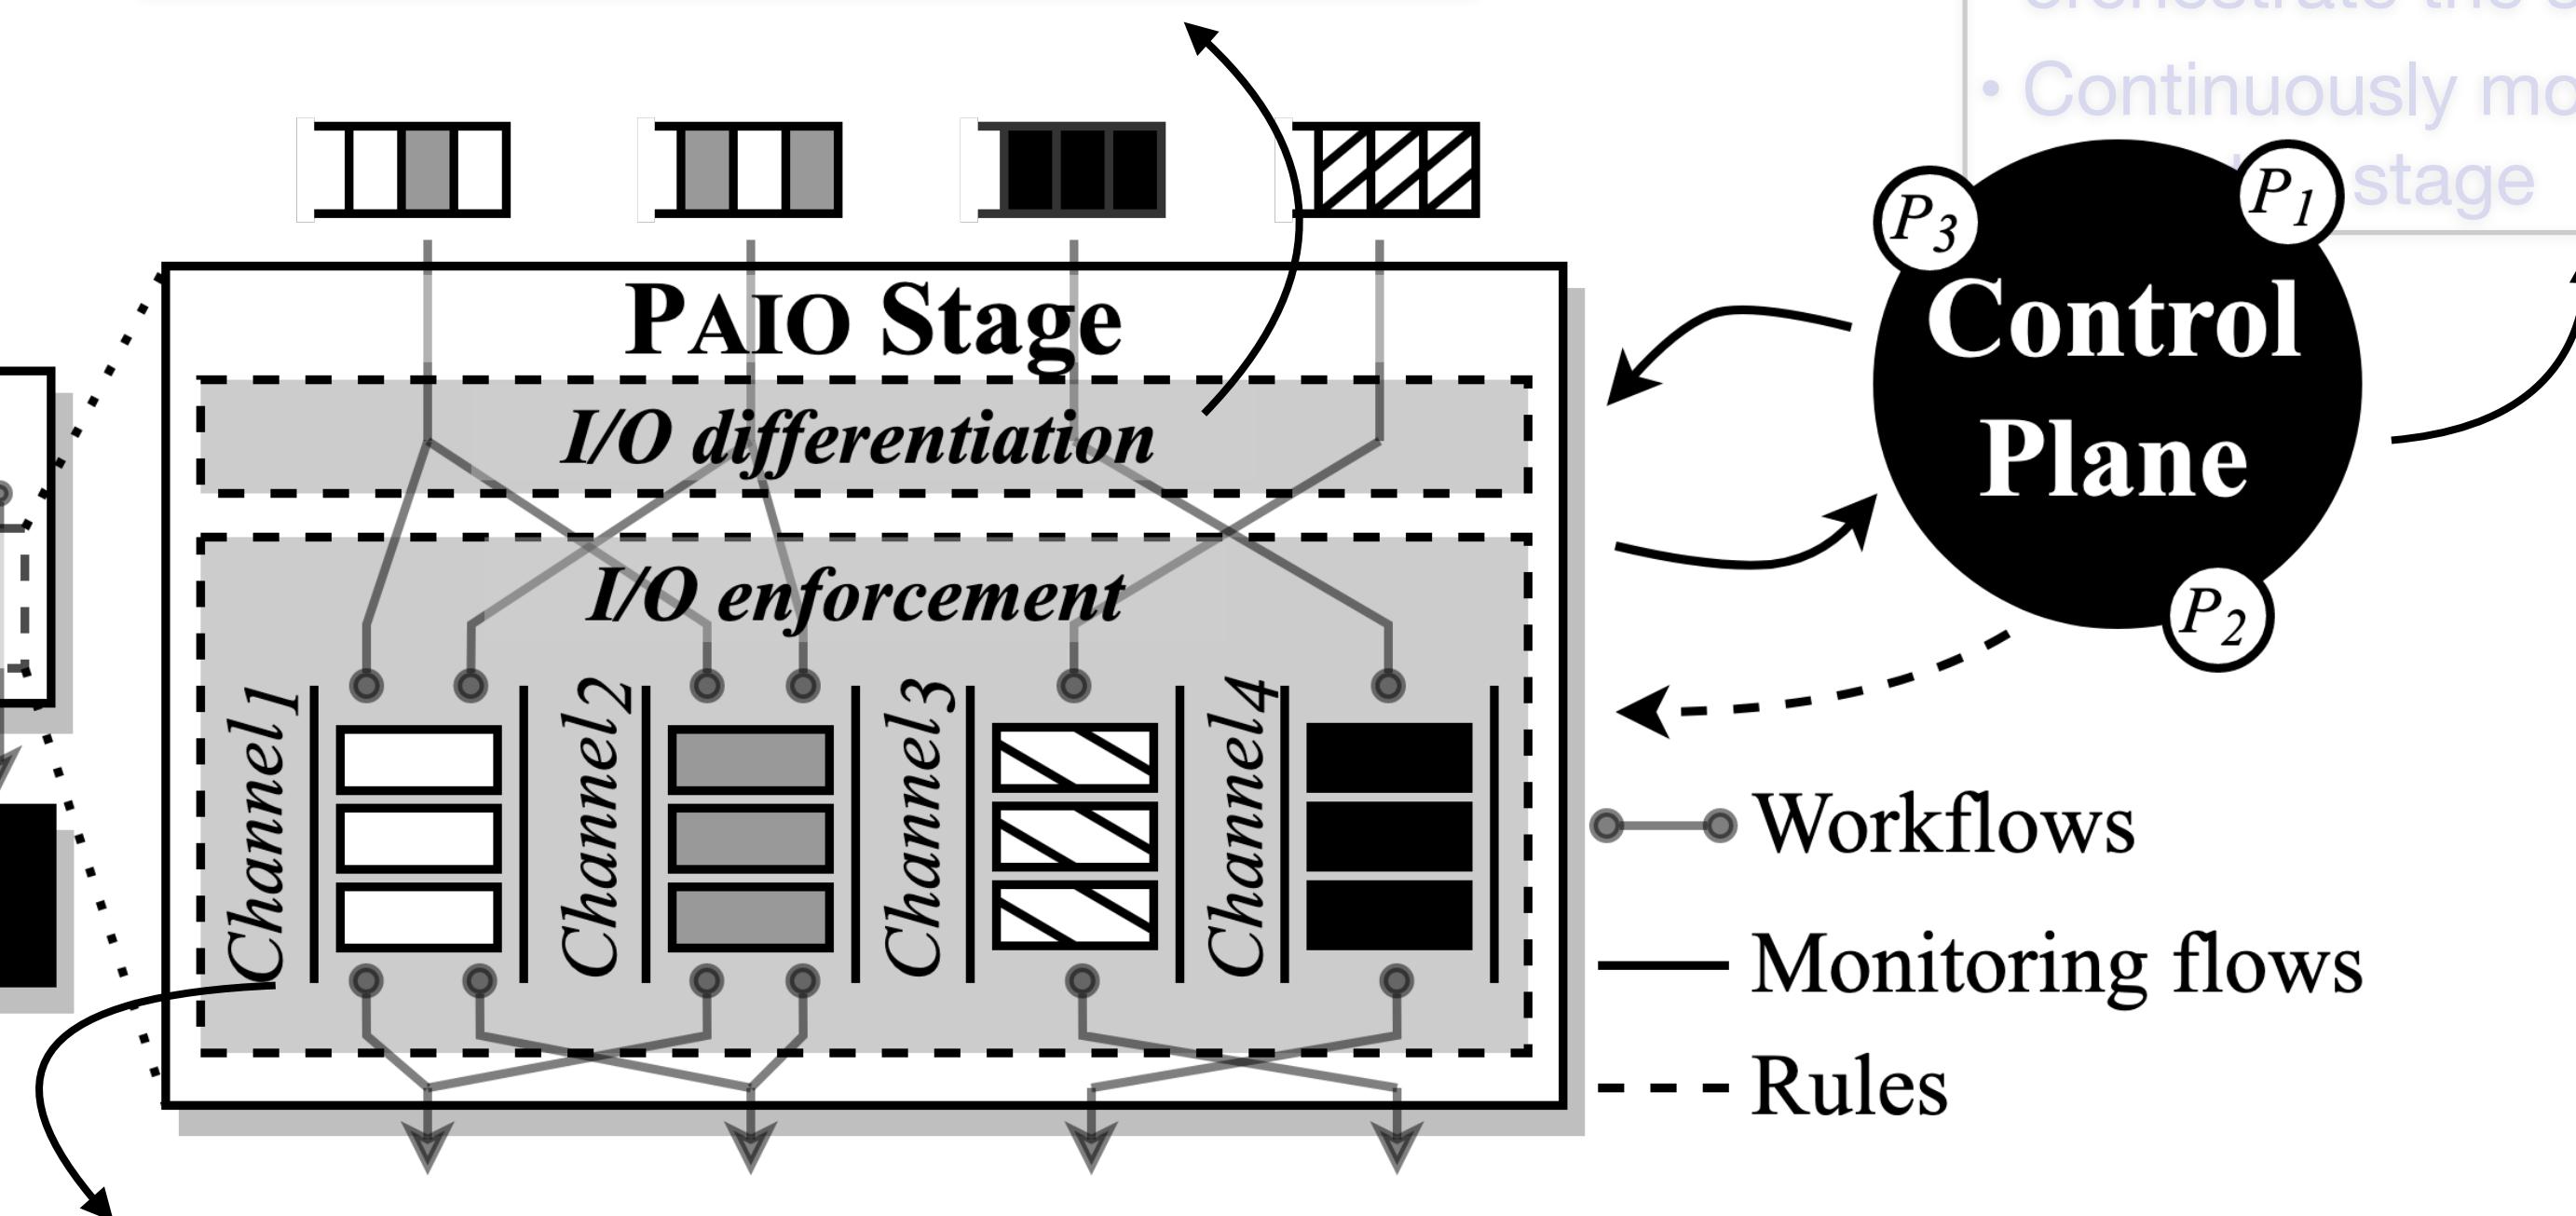


# PAIO design



**Differentiation module**

- Classifies and differentiates requests
- Context propagation
- Different levels of differentiation

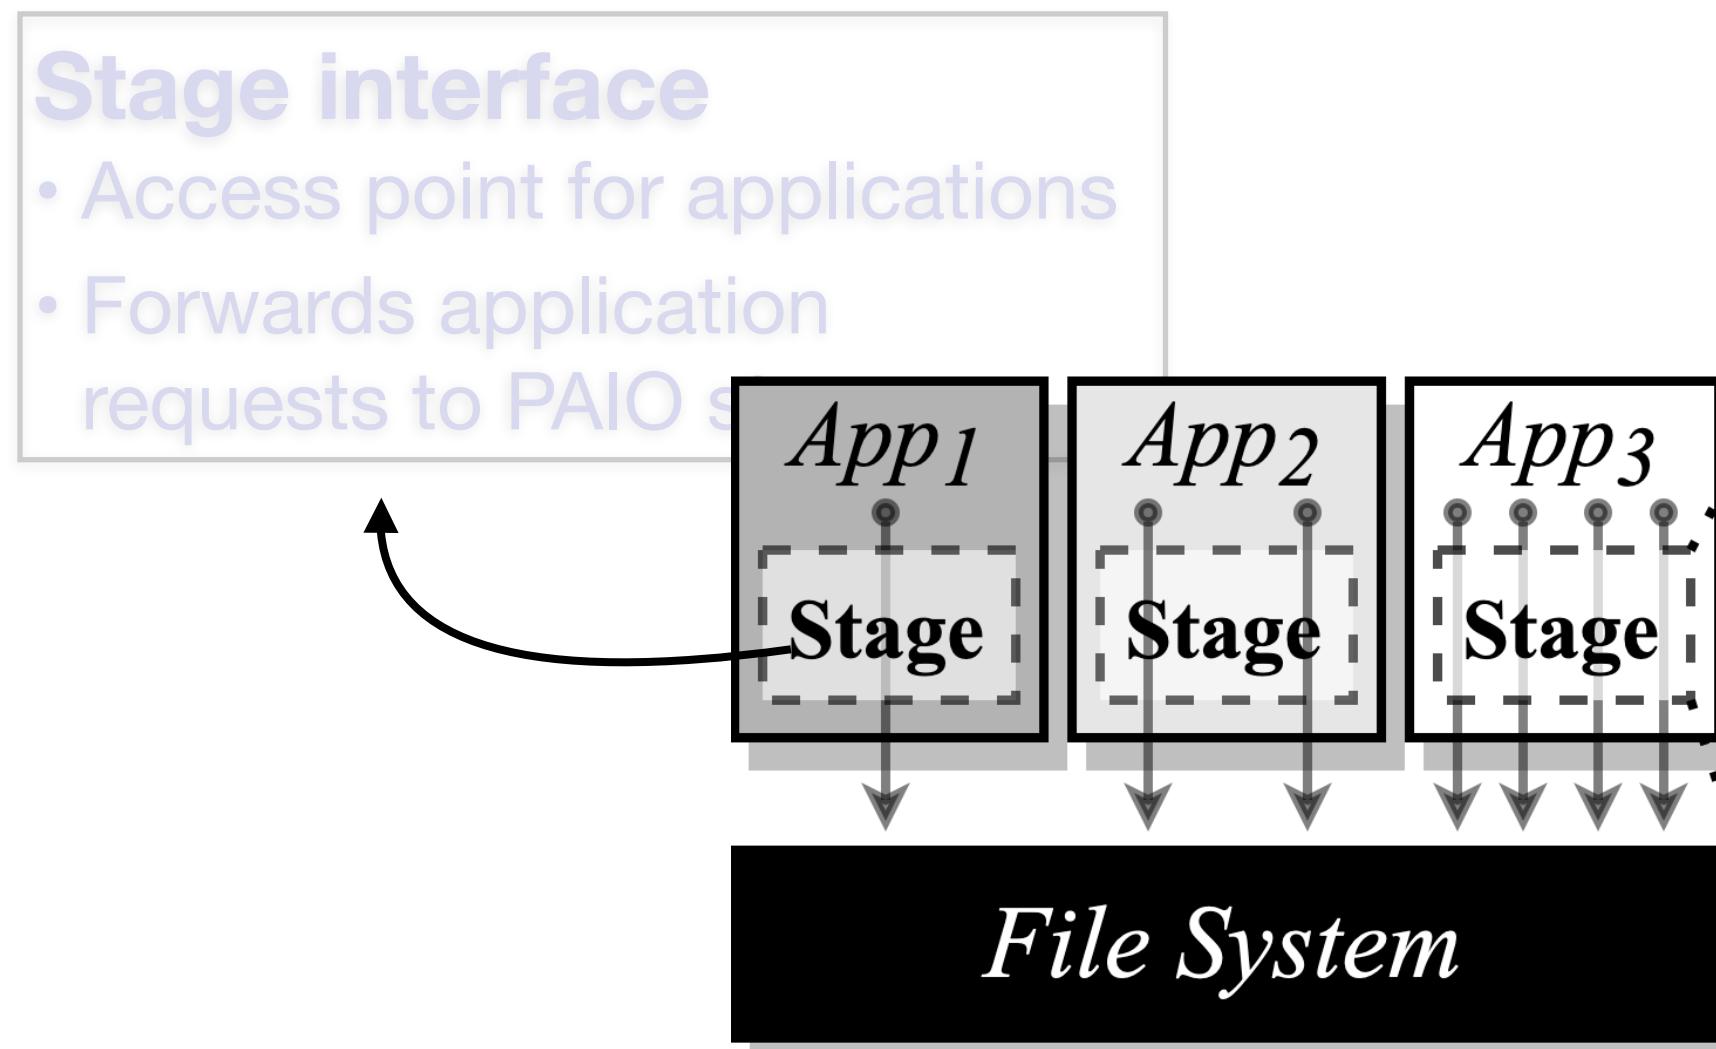


**Enforcement module**

- Applies the actual I/O mechanisms over requests
- Organized in Channels and Enforcement objects

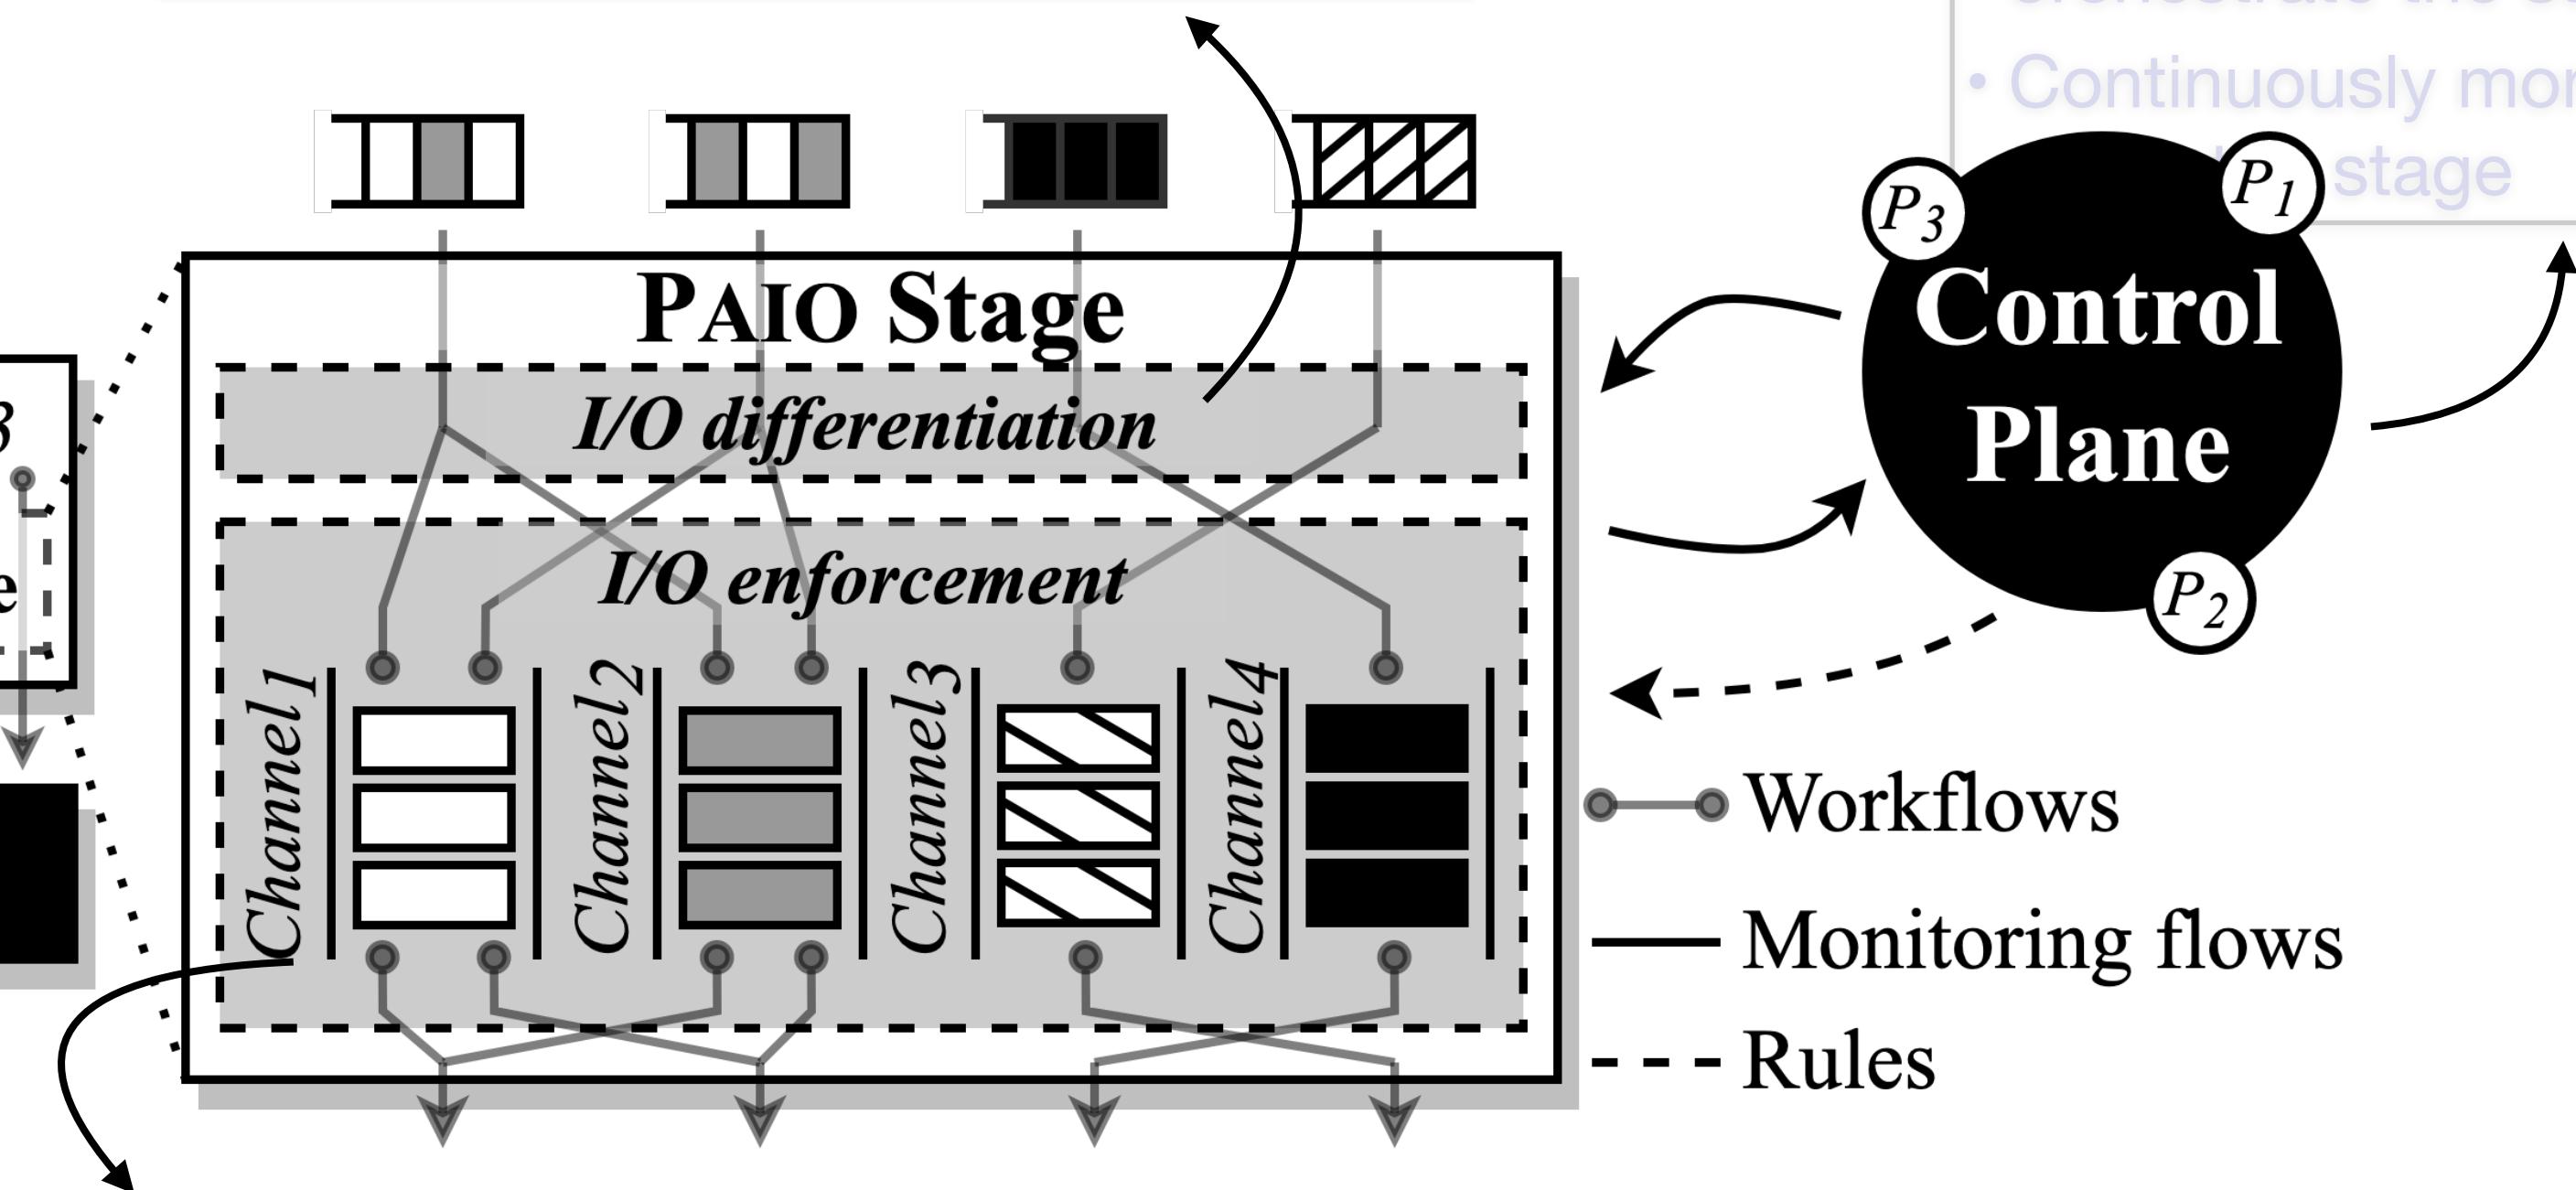


# PAIO design



**Differentiation module**

- Classifies and differentiates requests
- Context propagation
- Different levels of differentiation

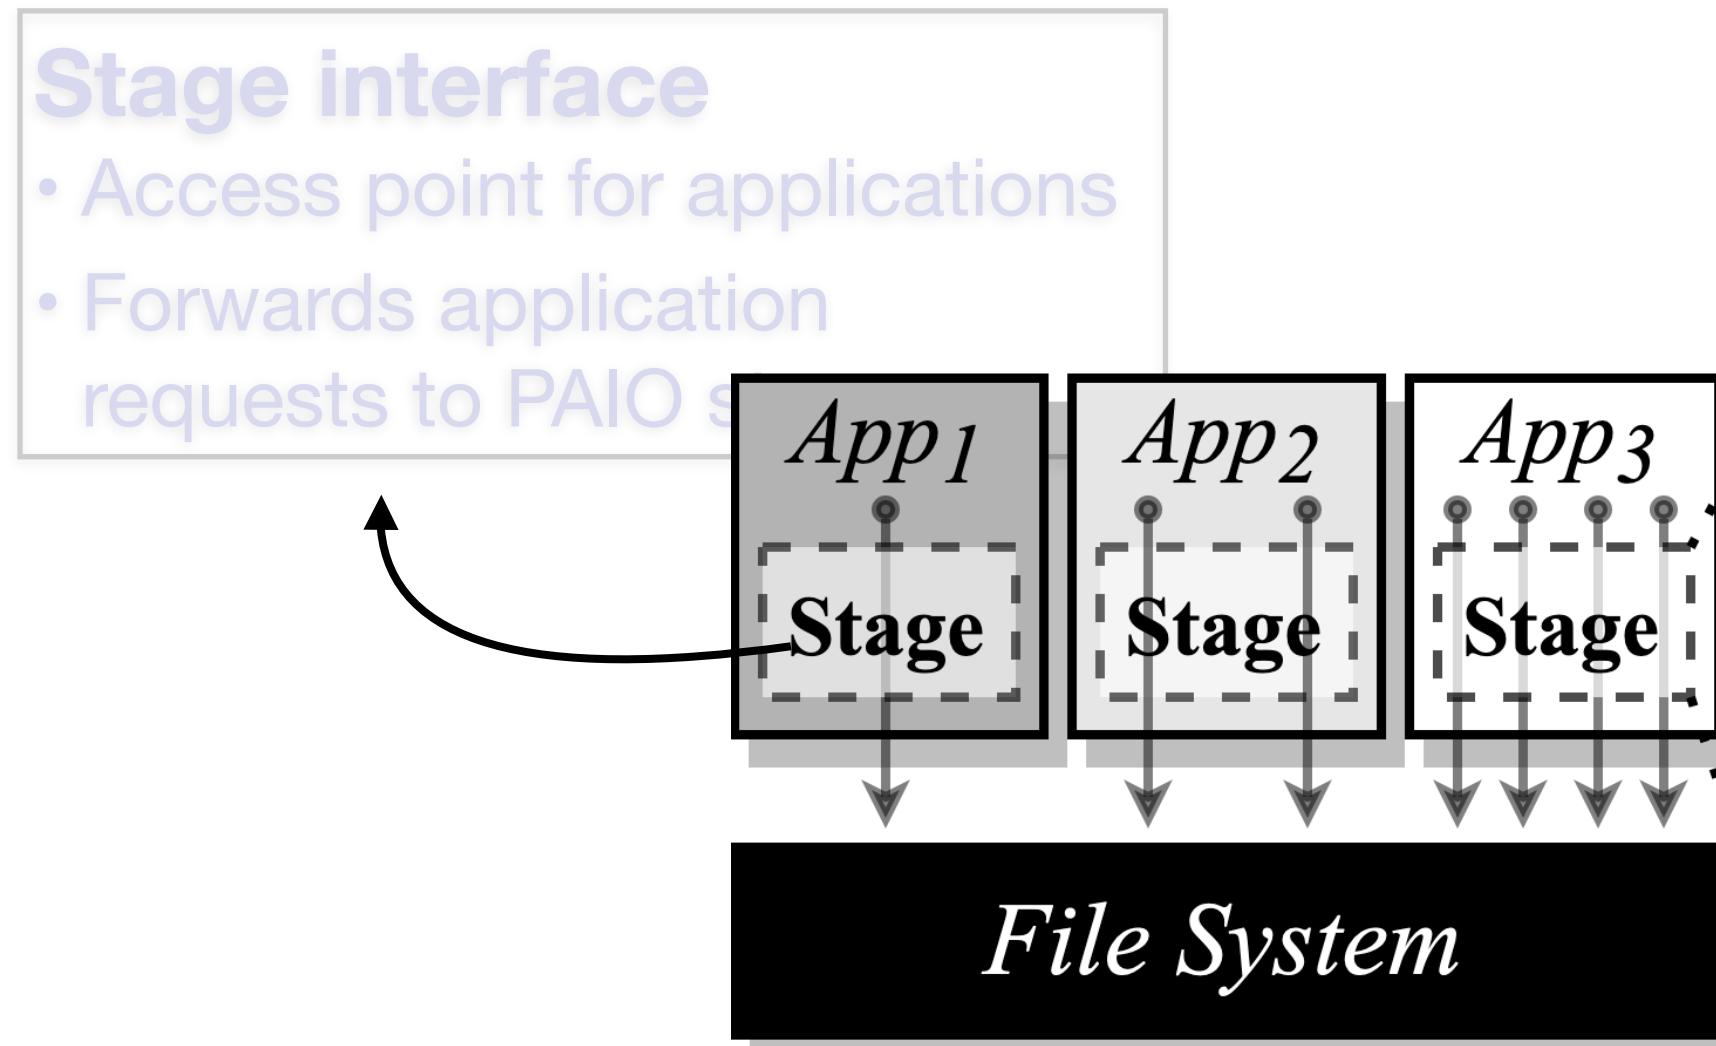


**Enforcement module**

- Applies the actual I/O mechanisms over requests
- Organized in Channels and Enforcement objects

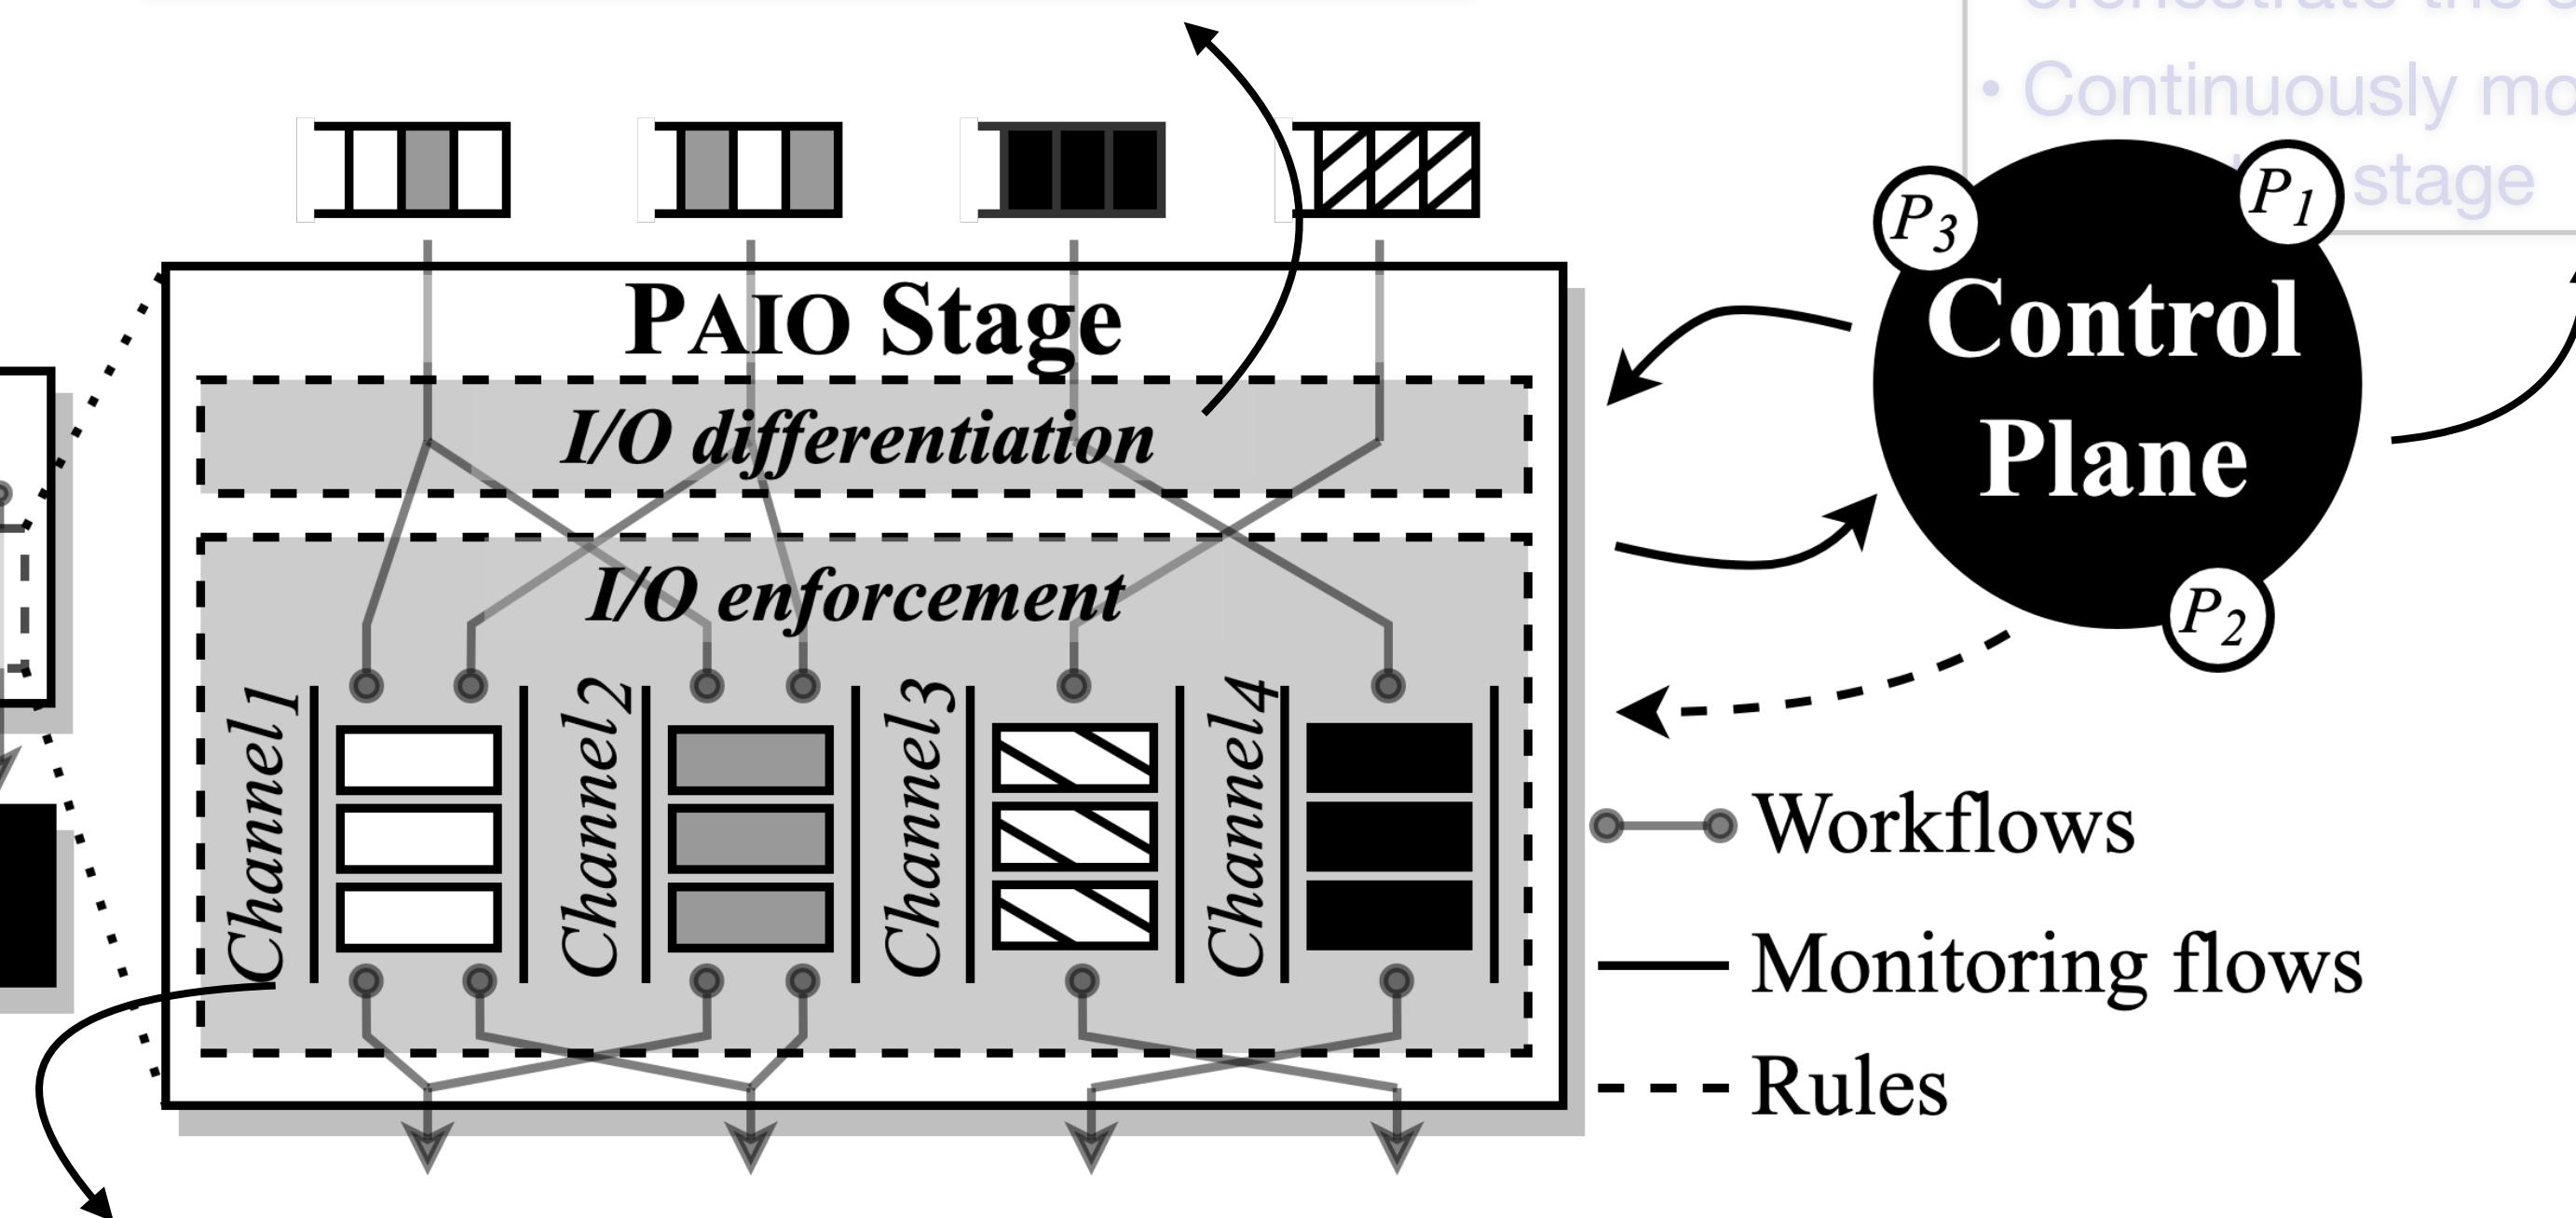


# PAIO design



**Differentiation module**

- Classifies and differentiates requests
- Context propagation
- Different levels of differentiation

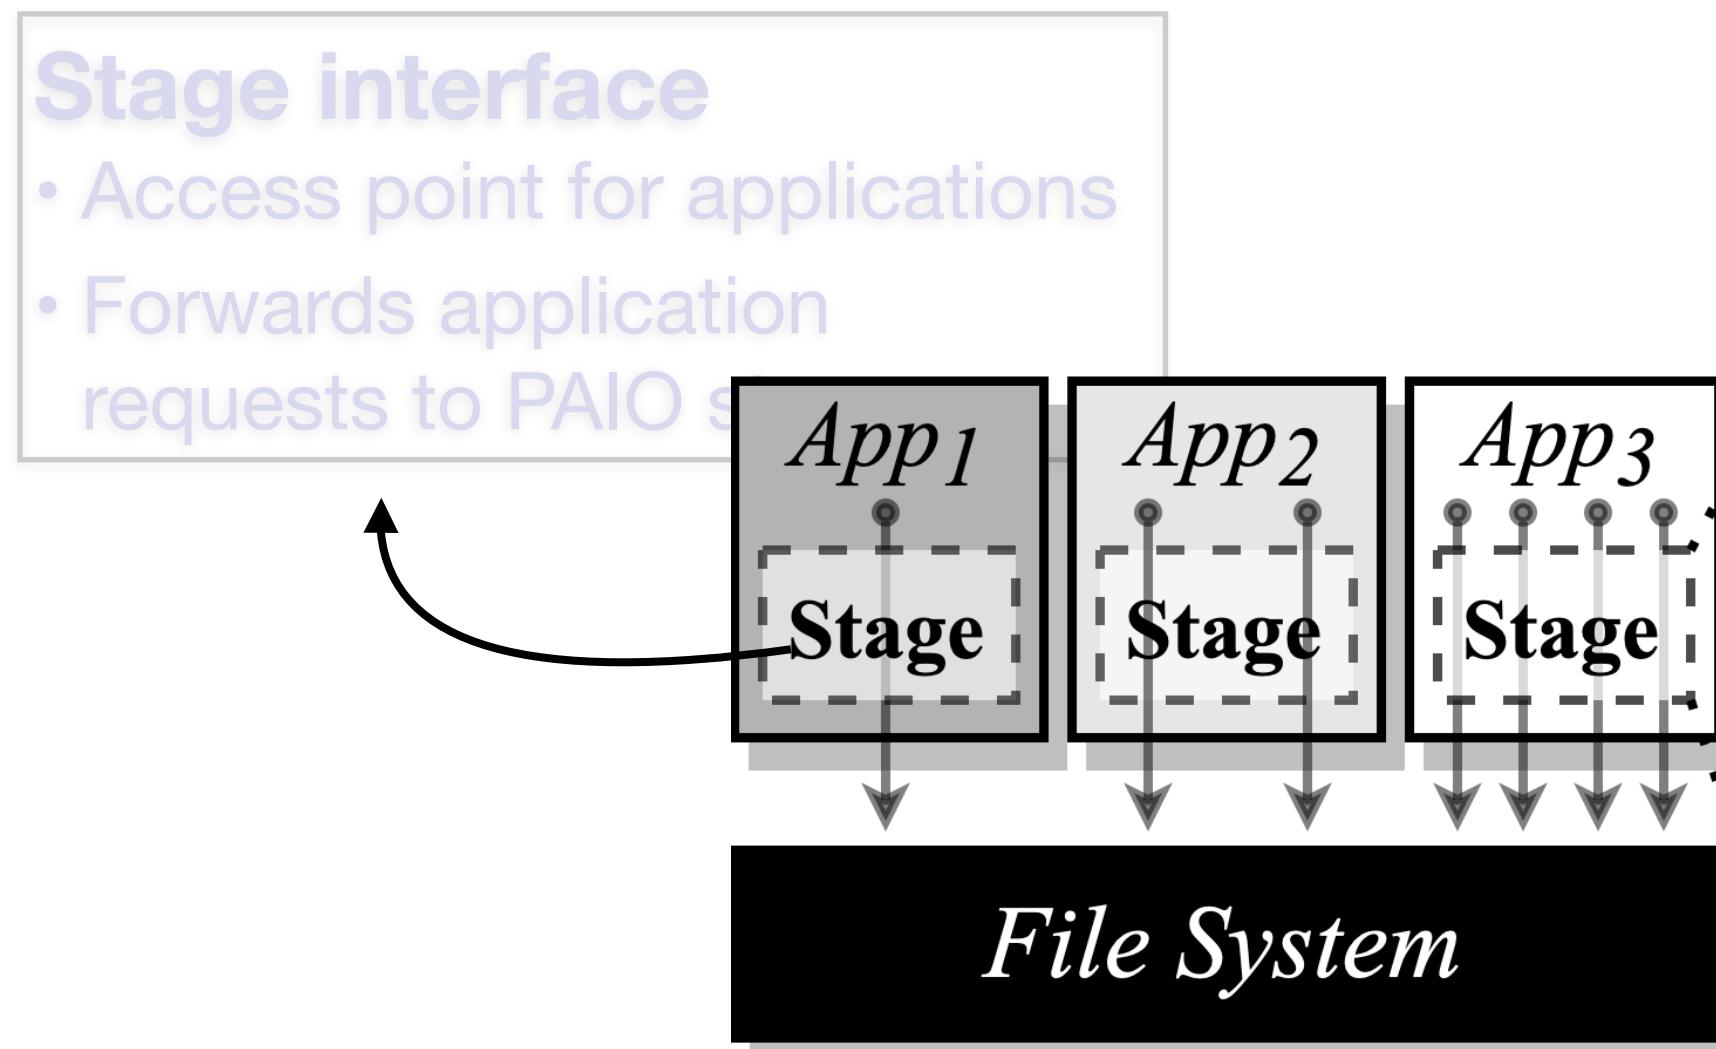


**Enforcement module**

- Applies the actual I/O mechanisms over requests
- Organized in Channels and Enforcement objects

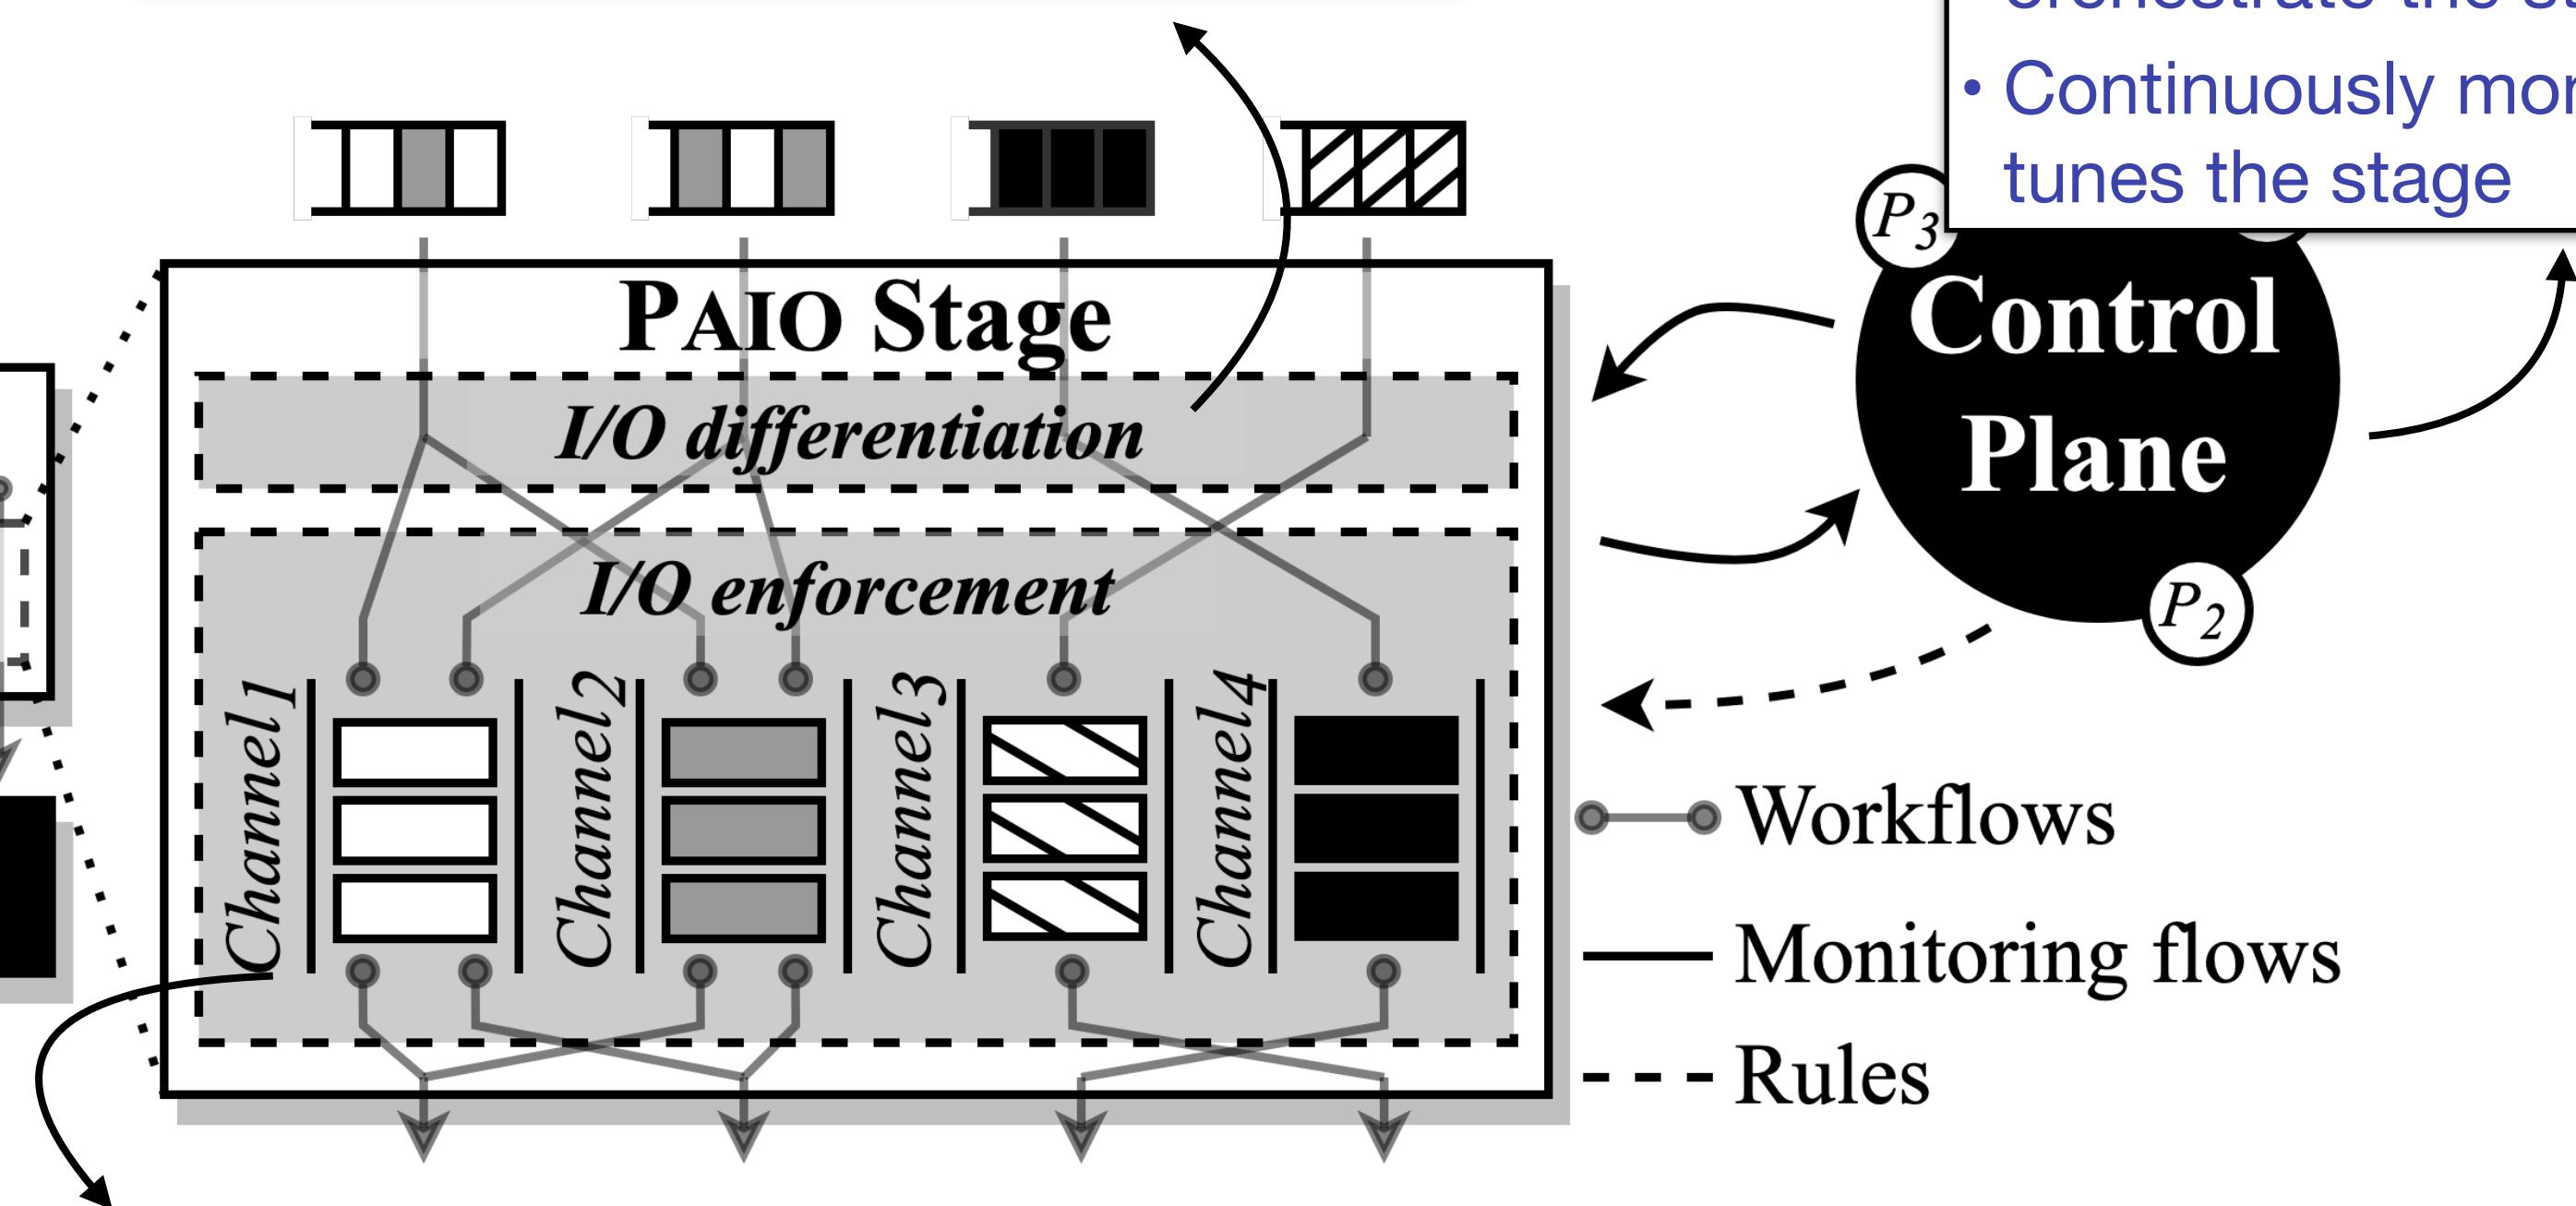


# PAIO design



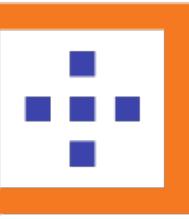
**Differentiation module**

- Classifies and differentiates requests
- Context propagation
- Different levels of differentiation



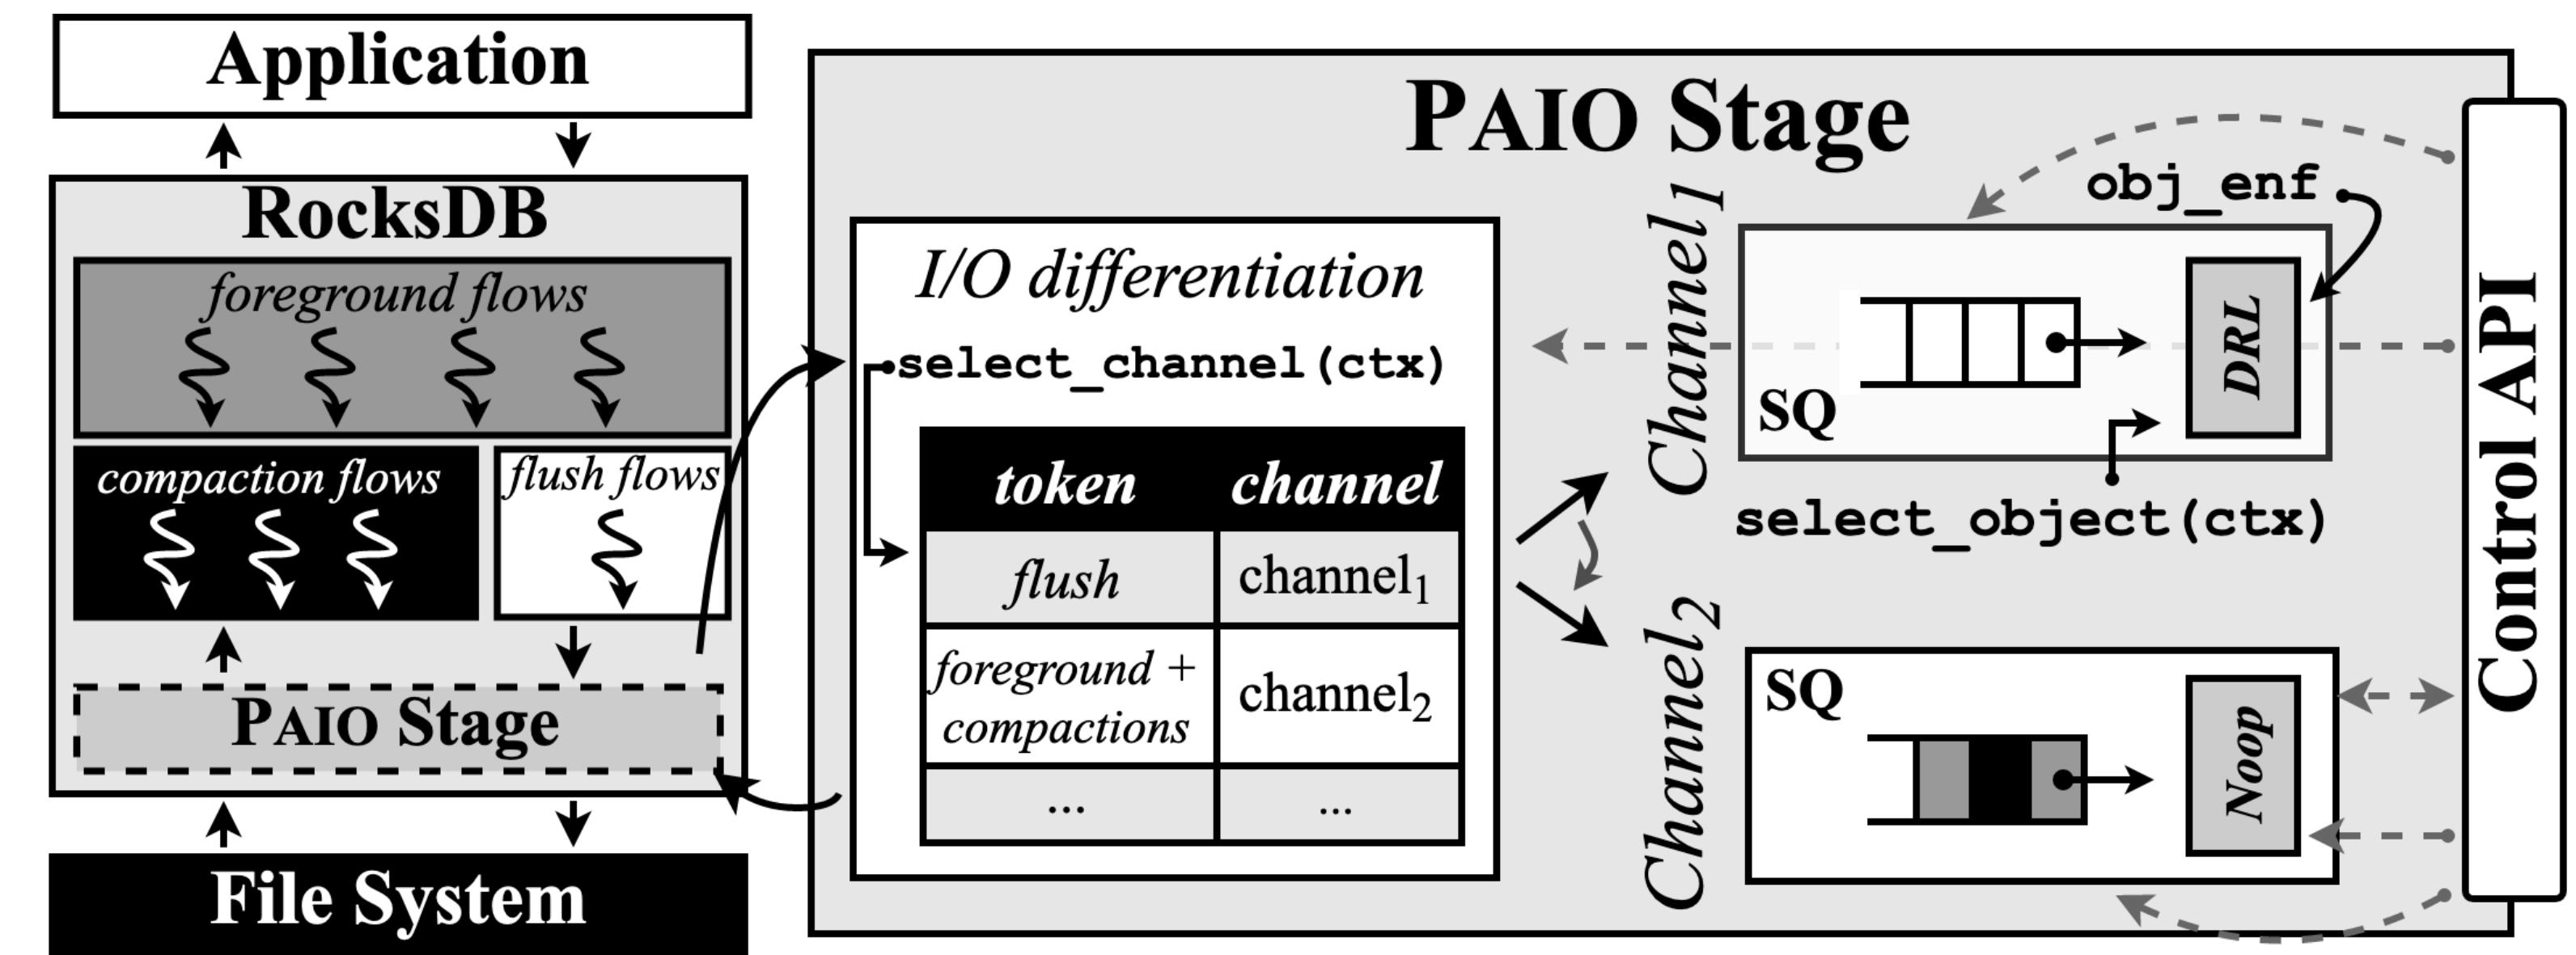
**Enforcement module**

- Applies the actual I/O mechanisms over requests
- Organized in Channels and Enforcement objects



# PAIO design

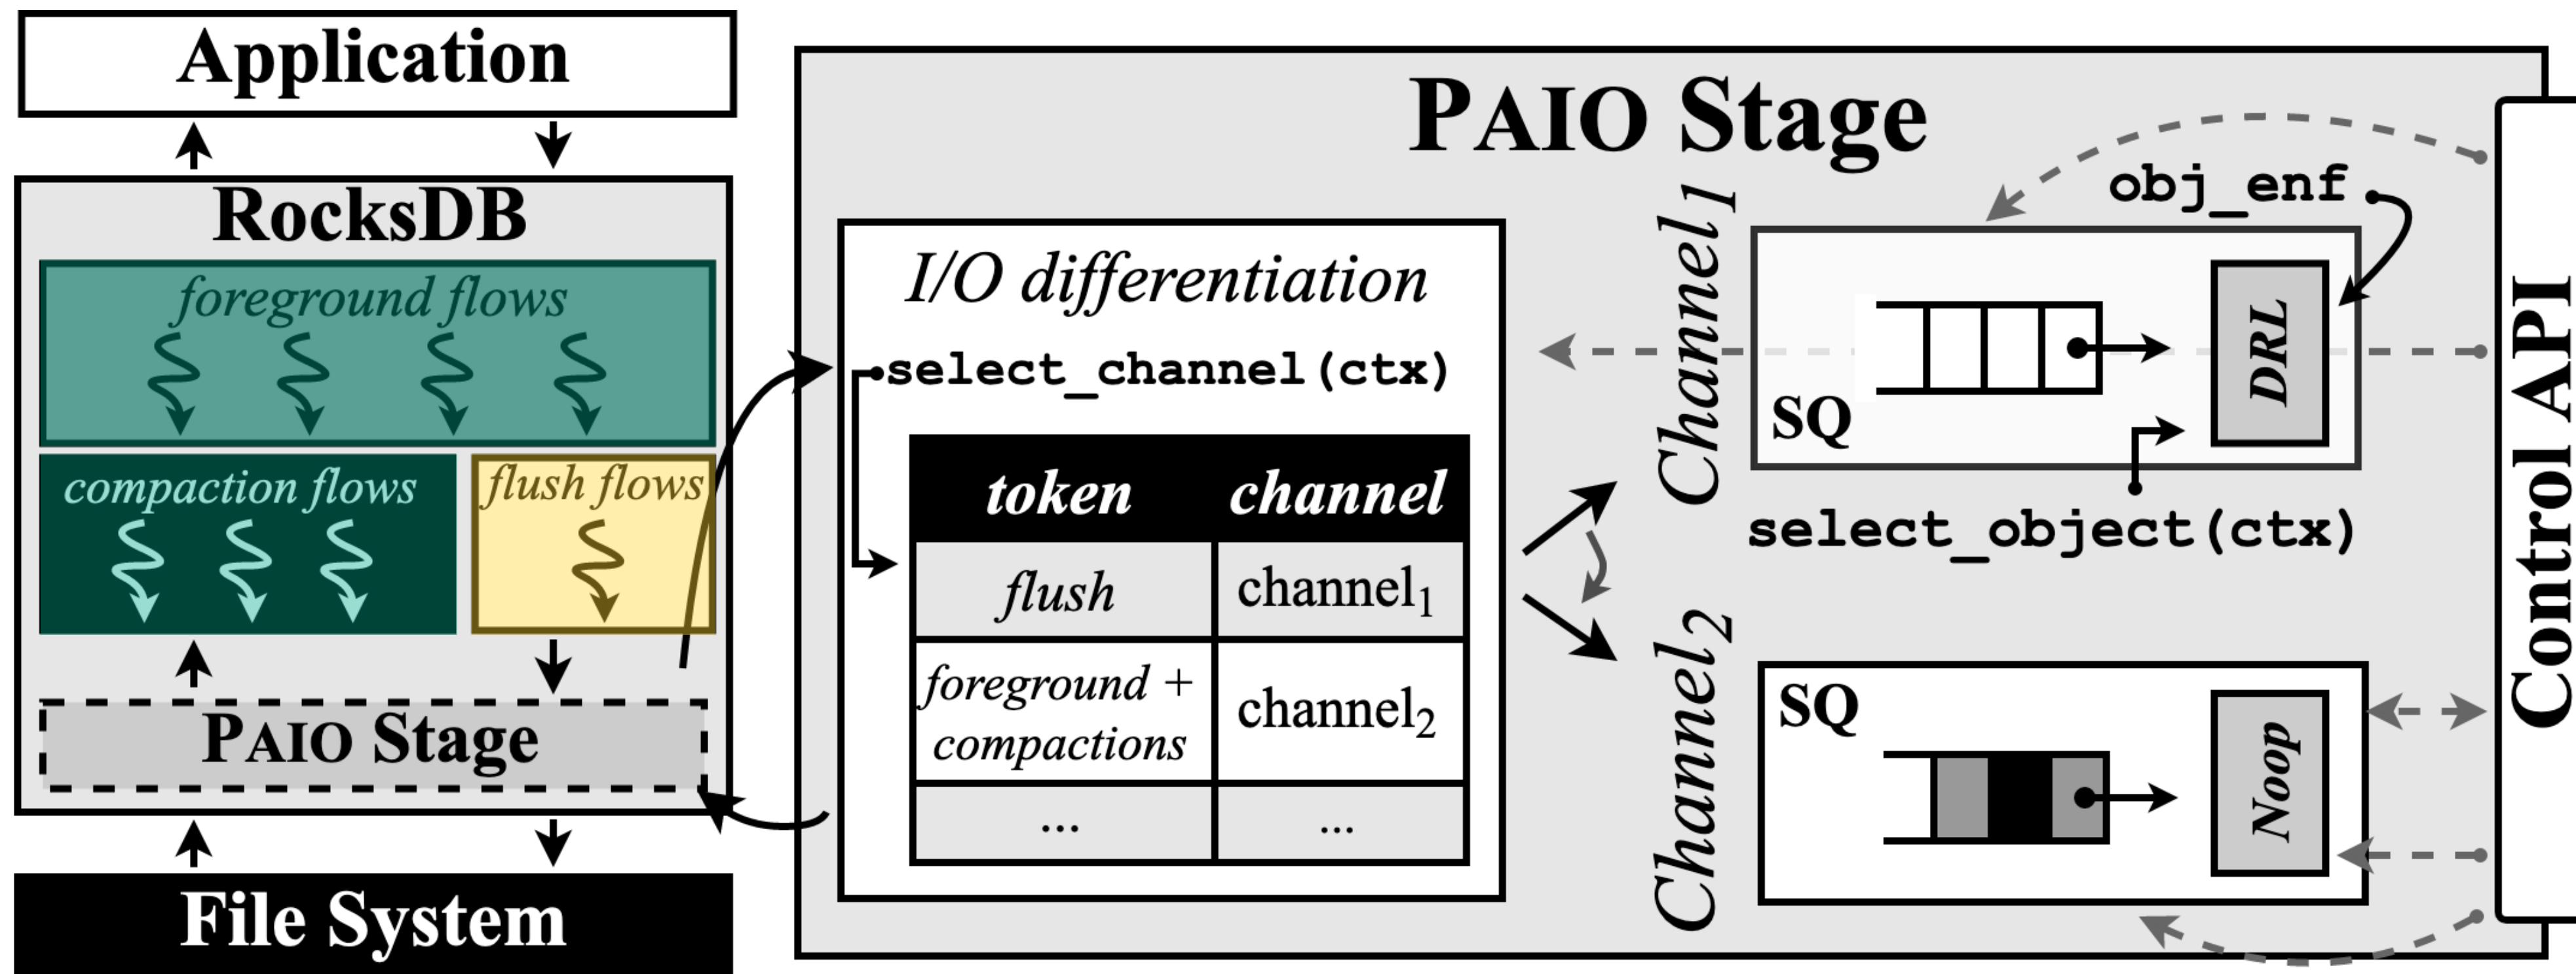
- I/O differentiation
- I/O enforcement
- Control plane interaction



**Policy:** *limit the rate of RocksDB's flush operations to X MiB/s*



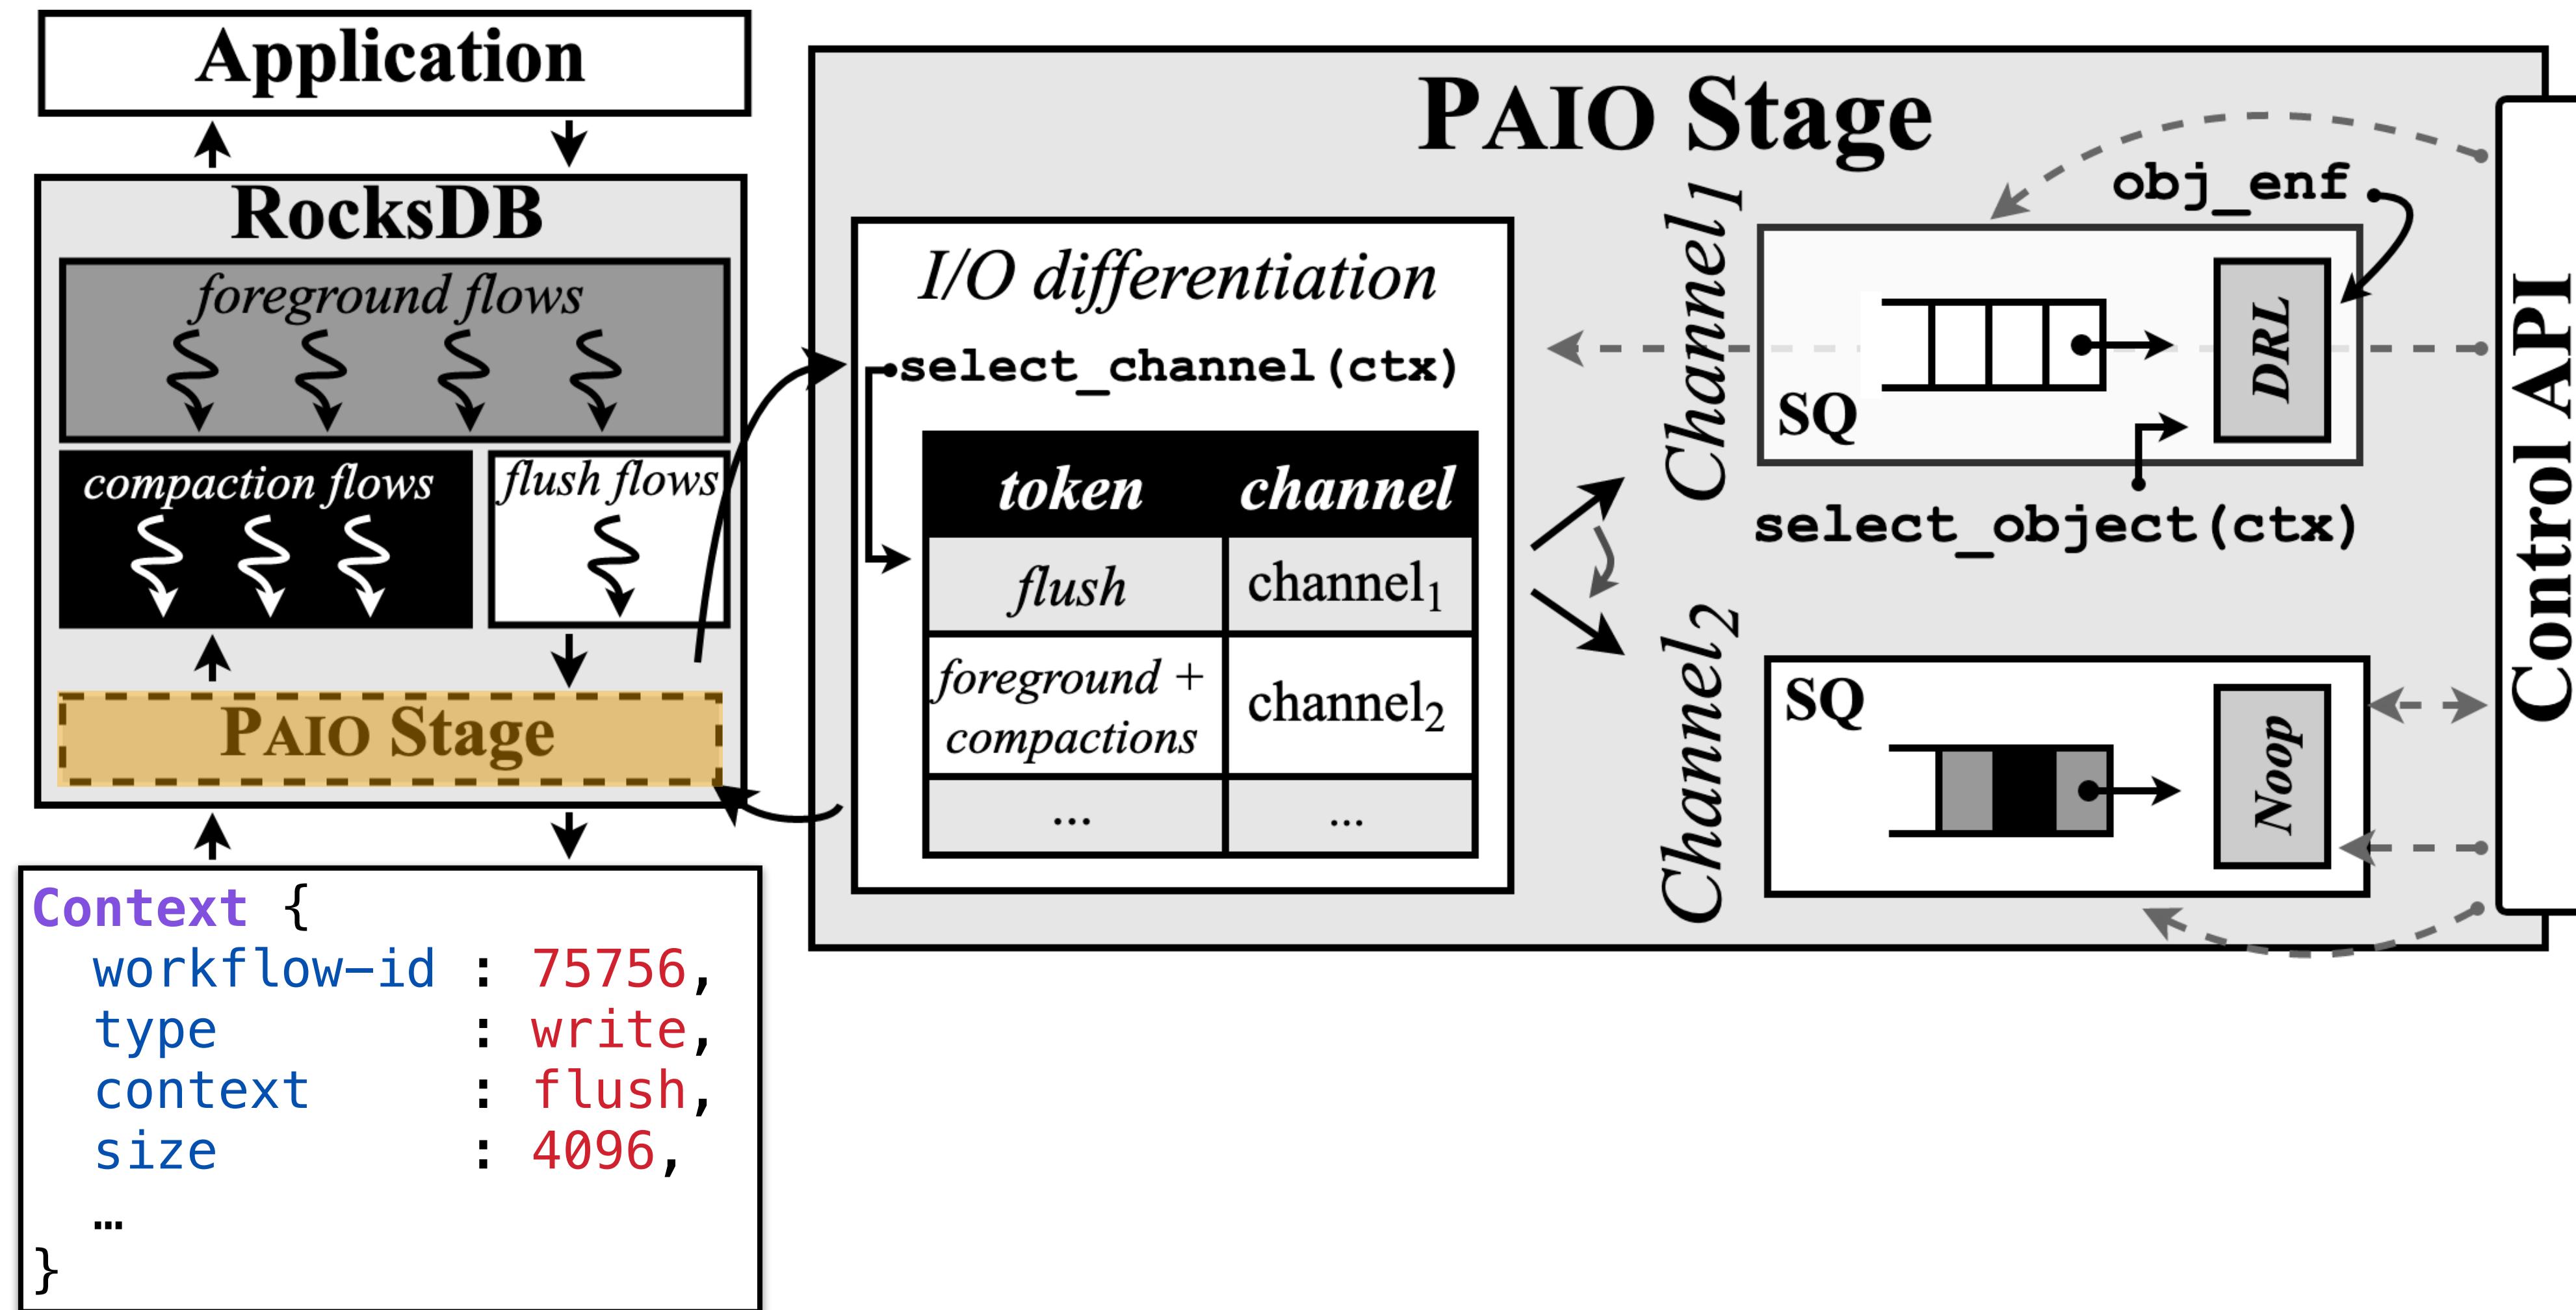
# I/O differentiation



Identify the origin of POSIX operations (i.e., **foreground**, **compaction**, or **flush** operations)

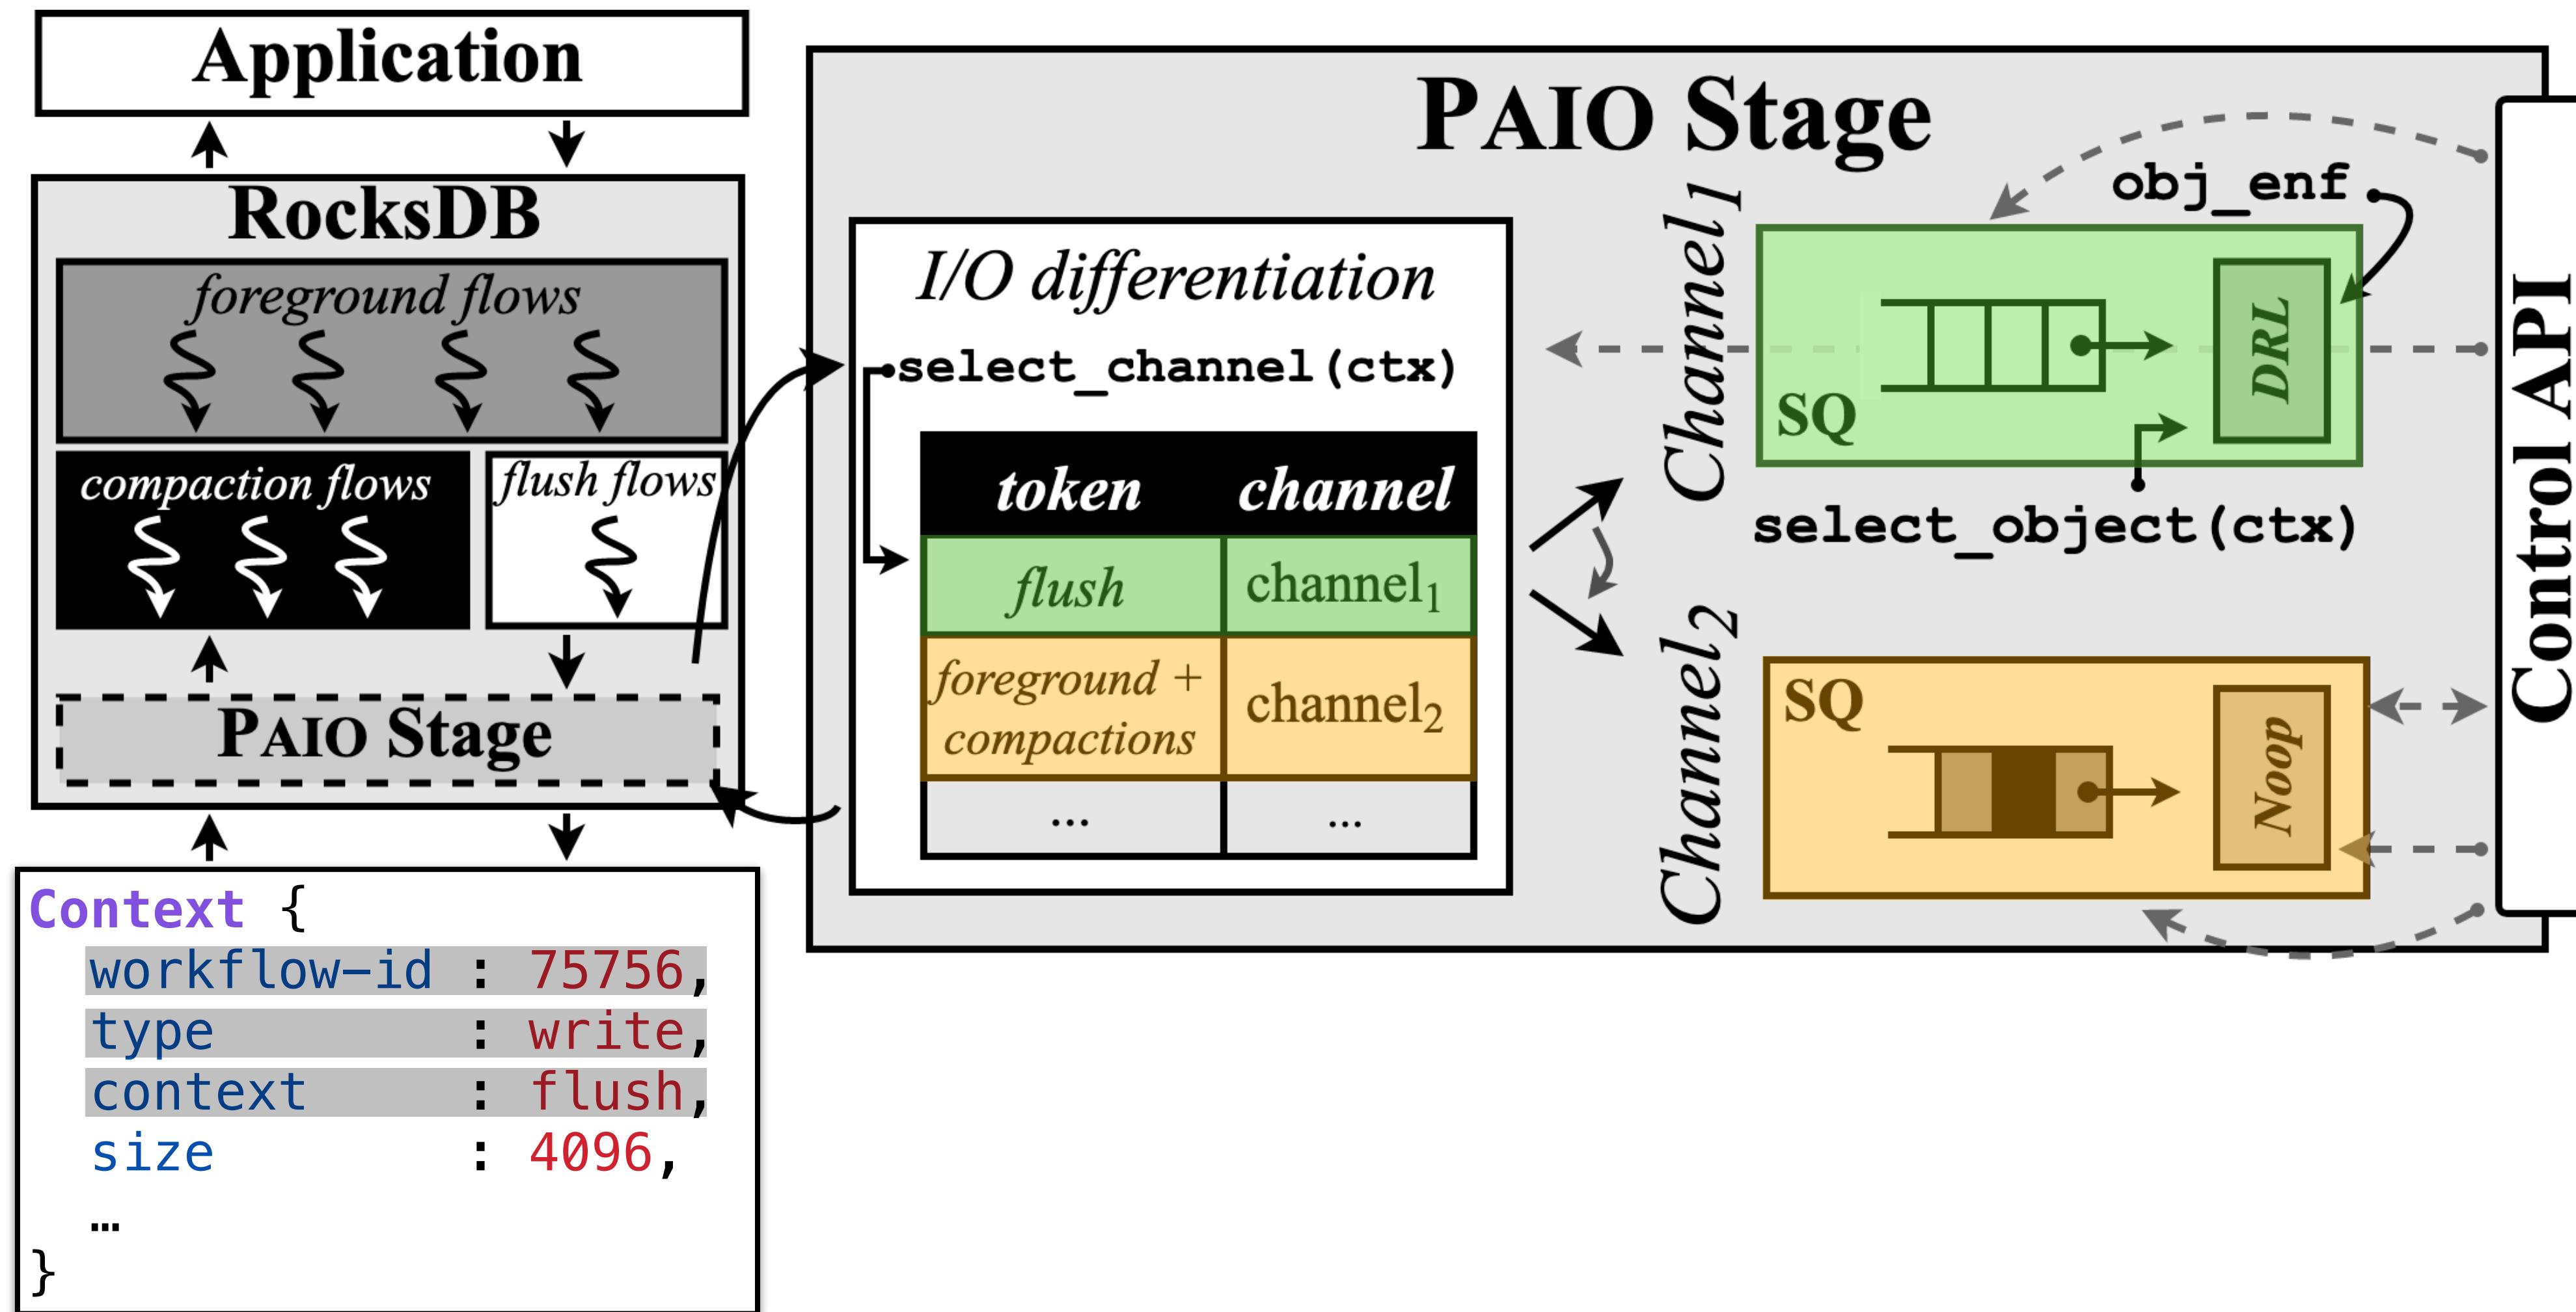


# I/O differentiation



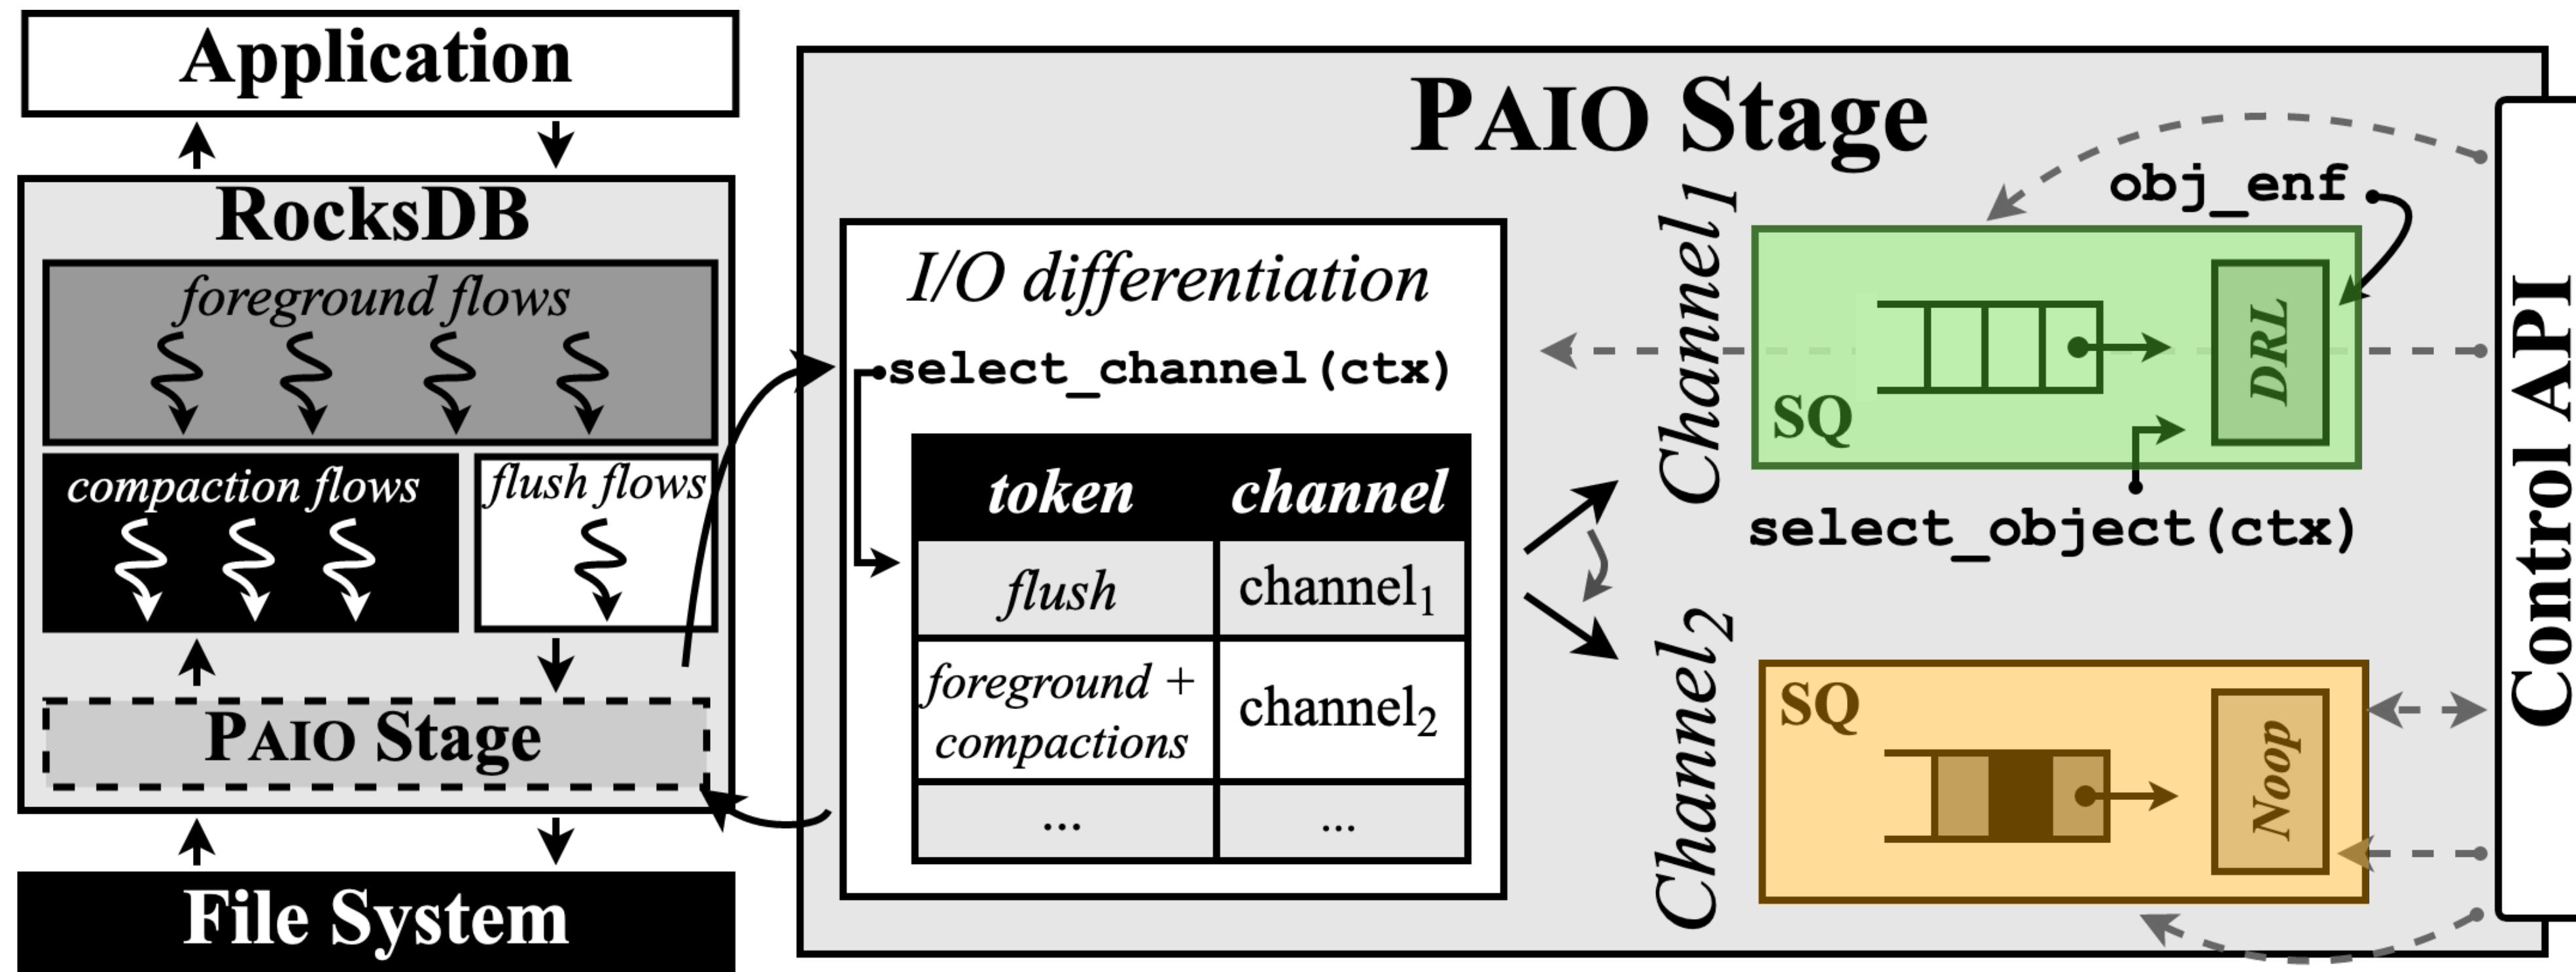


# I/O differentiation





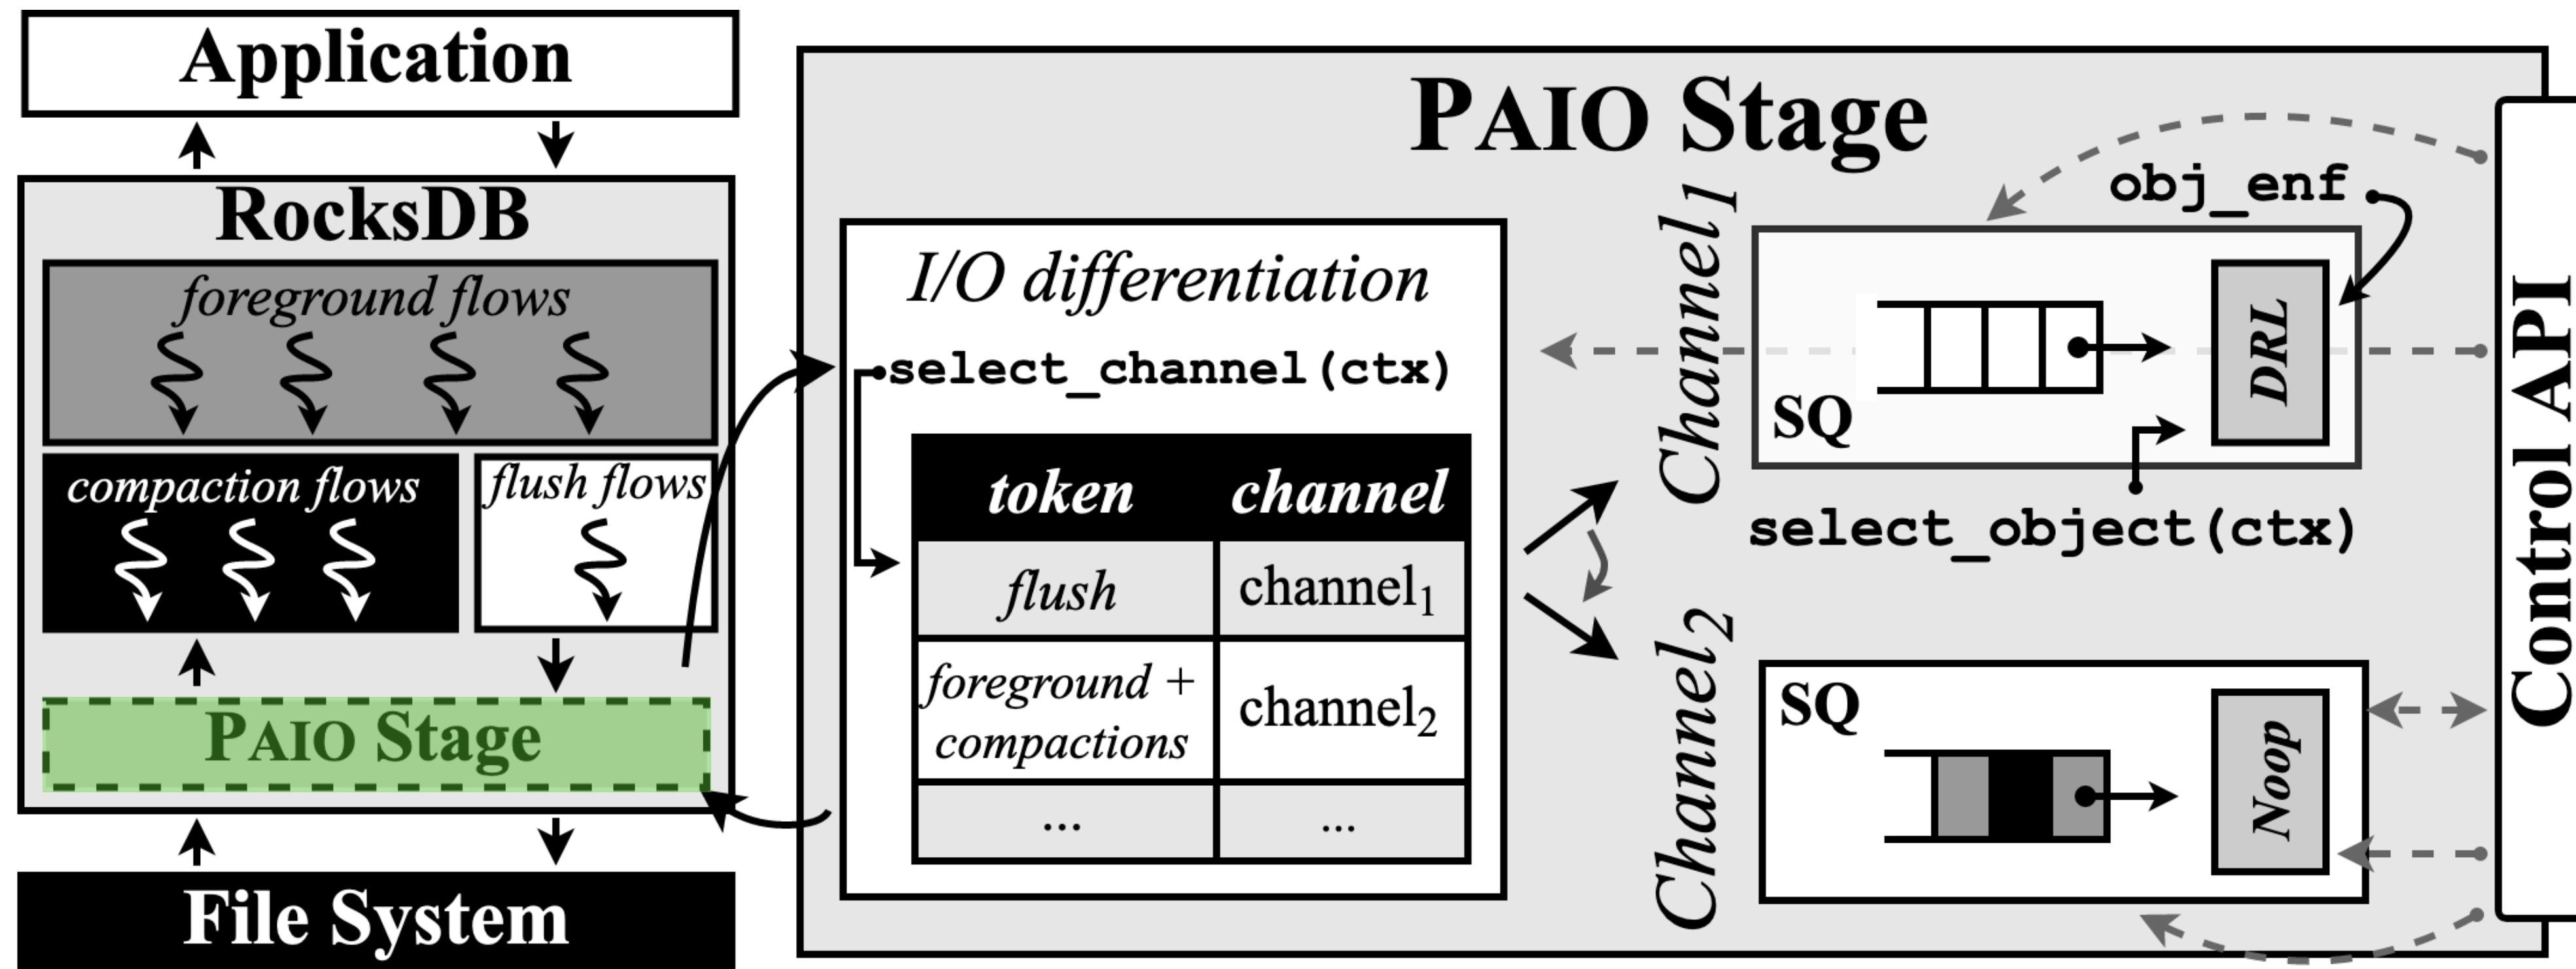
# I/O enforcement



PAIO currently supports **Noop**  
and **DRL** enforcement objects



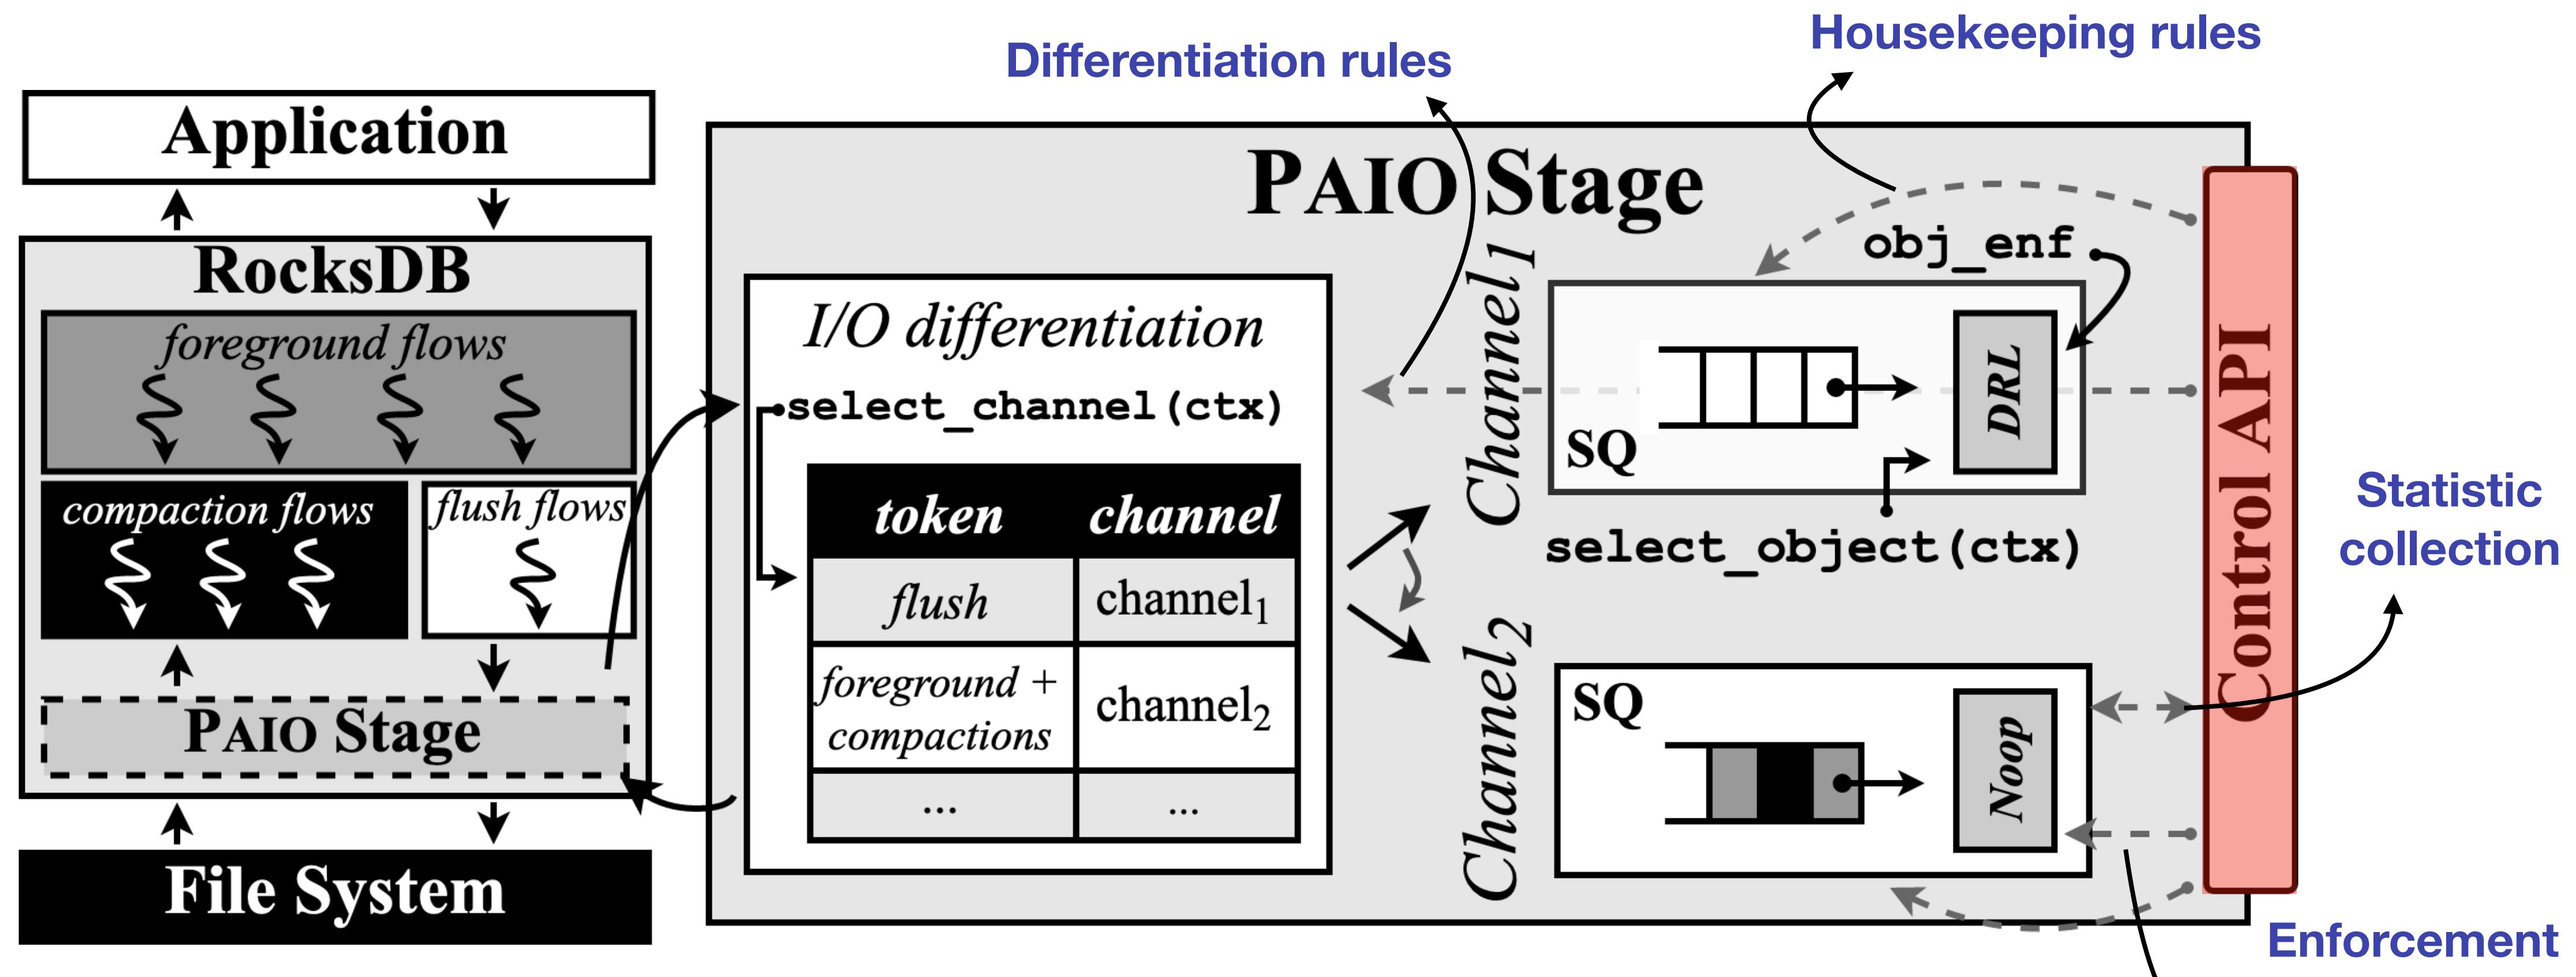
# I/O enforcement



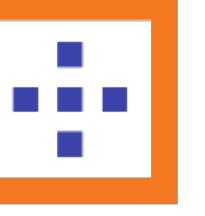
Requests return to their  
original I/O path



# Control plane interaction

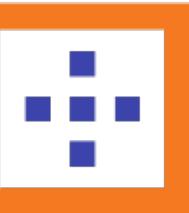


Implements the control algorithms for orchestrating stages (e.g., tail latency control, per-application bandwidth guarantees)



# Part 3

## building storage data planes



# Tail latency control in LSM-based KVS

## RocksDB

- Interference between foreground and background tasks generates high latency spikes
- Latency spikes occur due to L<sub>0</sub>-L<sub>1</sub> compactions and flushes being slow or on hold

## SILK

- I/O scheduler
  - Allocates bandwidth for internal operations when client load is low
  - Prioritizes flushes and low level compactions
  - Preempts high level compactions with low level ones
- Required changing several core modules made of thousands of LoC

## PAIO

- Stage provides the I/O mechanisms for prioritizing and rate limiting background flows
  - Integrating PAIO in RocksDB only required adding 85 LoC
- Control plane provides a SILK-based I/O scheduling algorithm



# Tail latency control in LSM-based KVS

## RocksDB

- Interference between foreground and background t
- Latency spikes occur due to L<sub>0</sub>-L<sub>1</sub> compactions and

## SILK

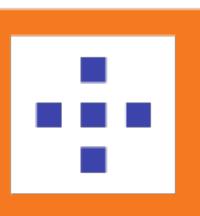
- I/O scheduler
  - Allocates bandwidth for internal operations when client load is low
  - Prioritizes flushes and low level compactions
  - ~~Preempts high level compactions with low level ones~~
- Required changing several core modules made of thousands of LoC

## PAIO

- Stage provides the I/O mechanisms for prioritizing and rate limiting background flows
  - Integrating PAIO in RocksDB only required adding 85 LoC
- Control plane provides a SILK-based I/O scheduling algorithm

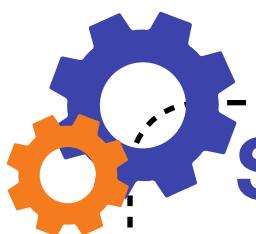


**Note:** By propagating application-level information to the stage, PAIO can enable similar control and performance as system-specific optimizations



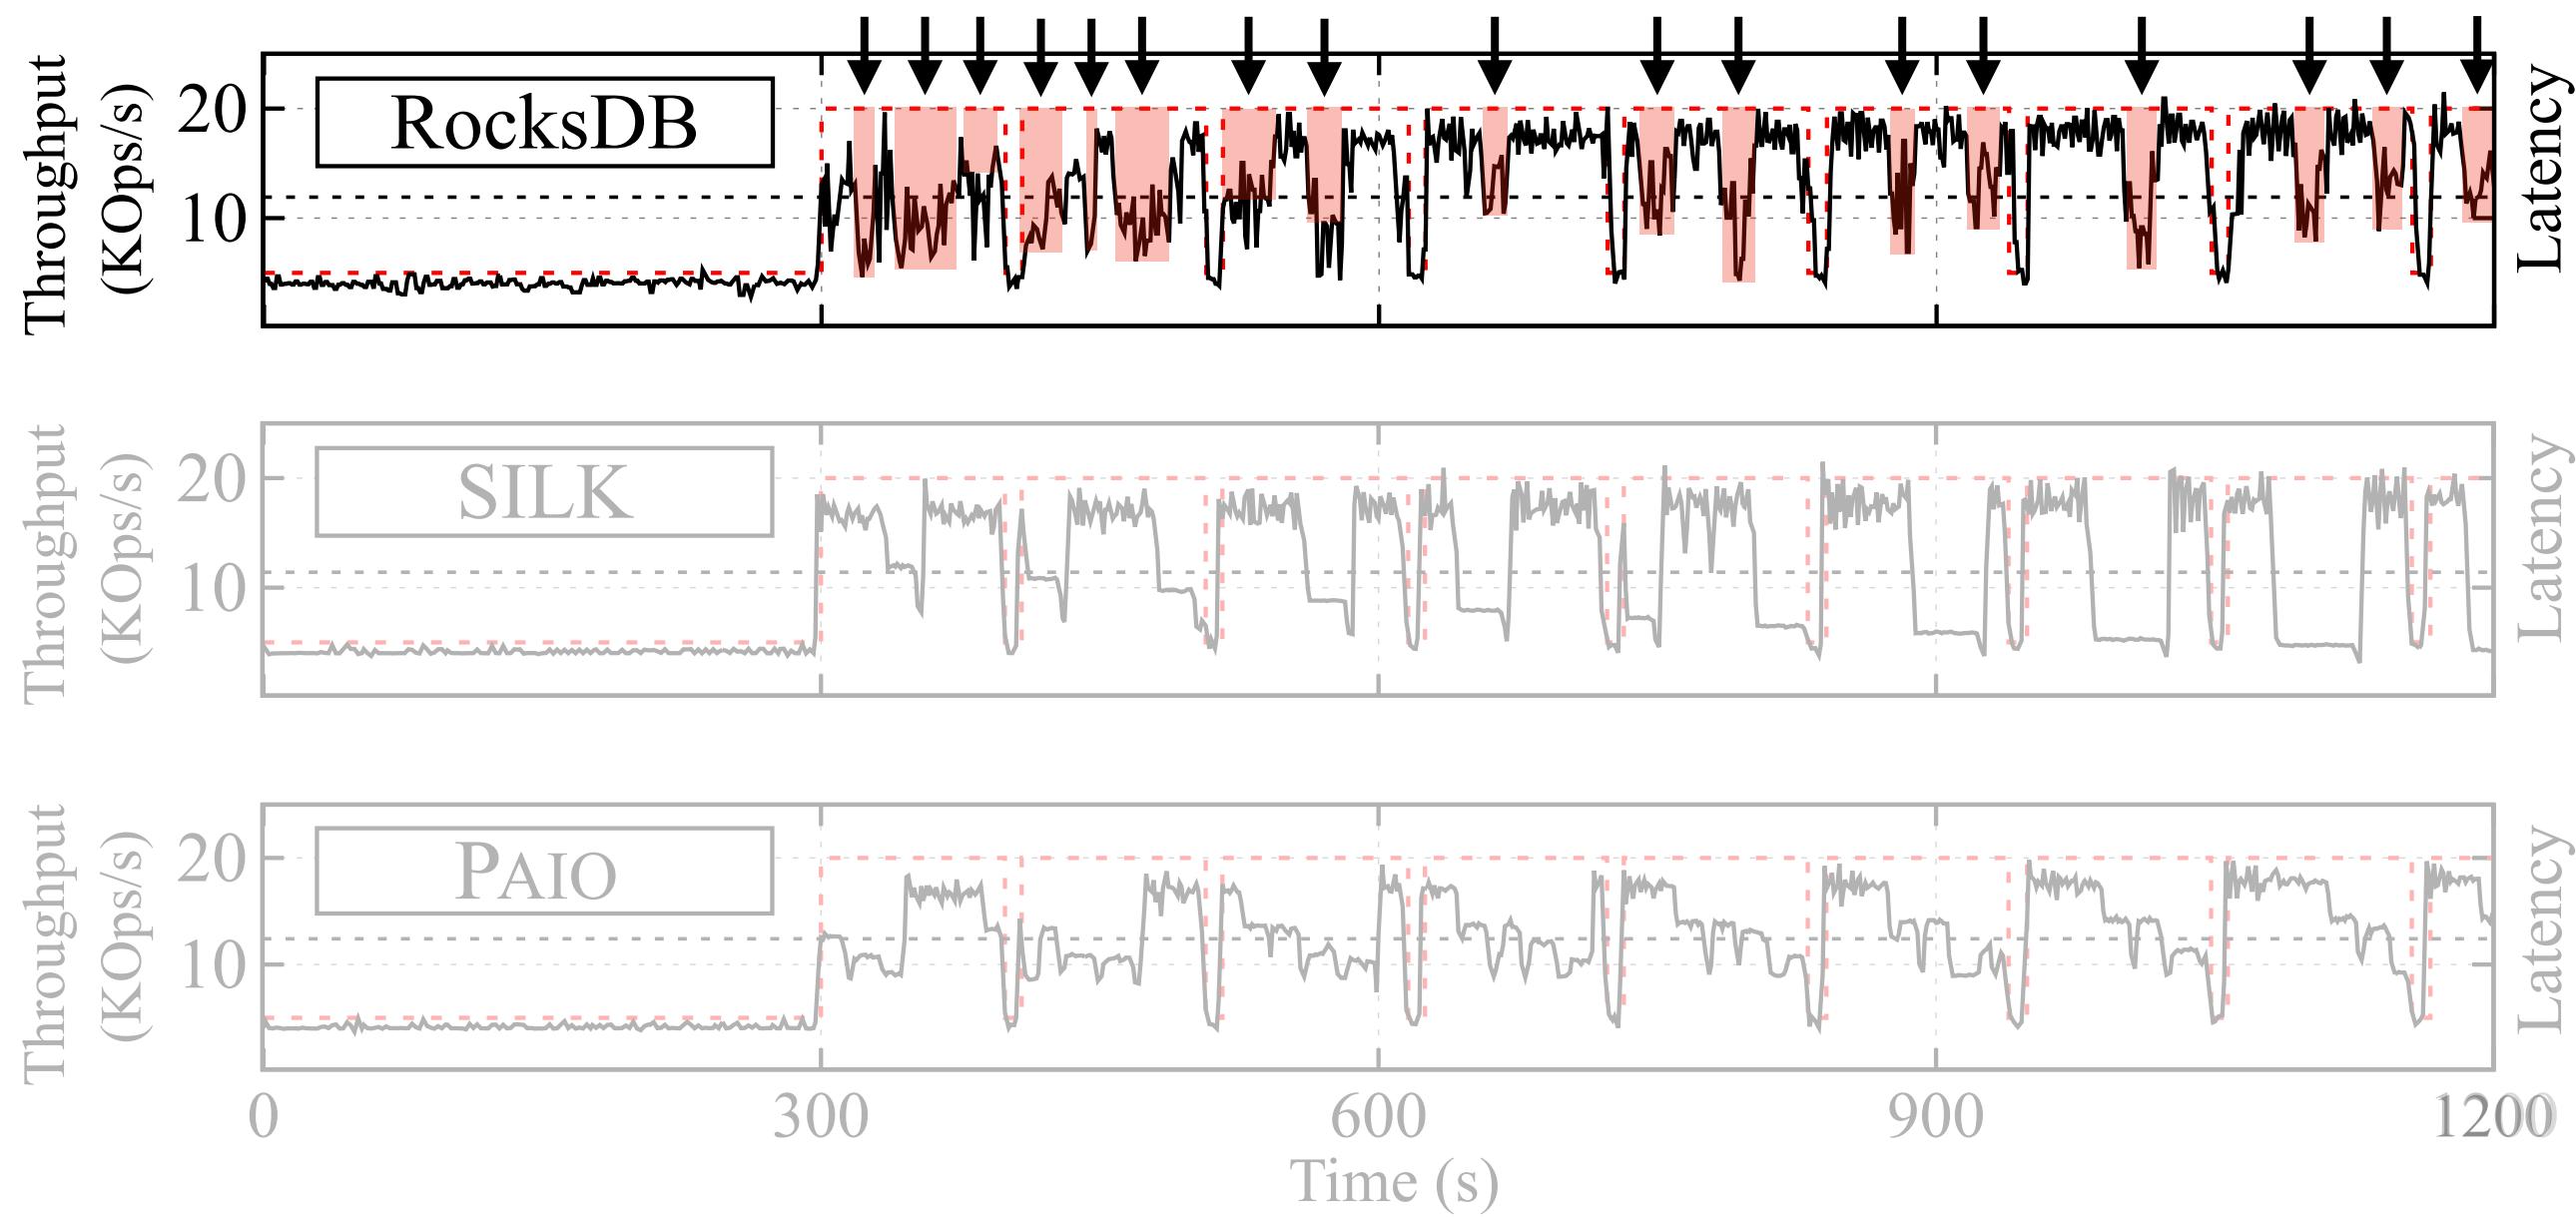
# Mixture workload

50% read 50% write

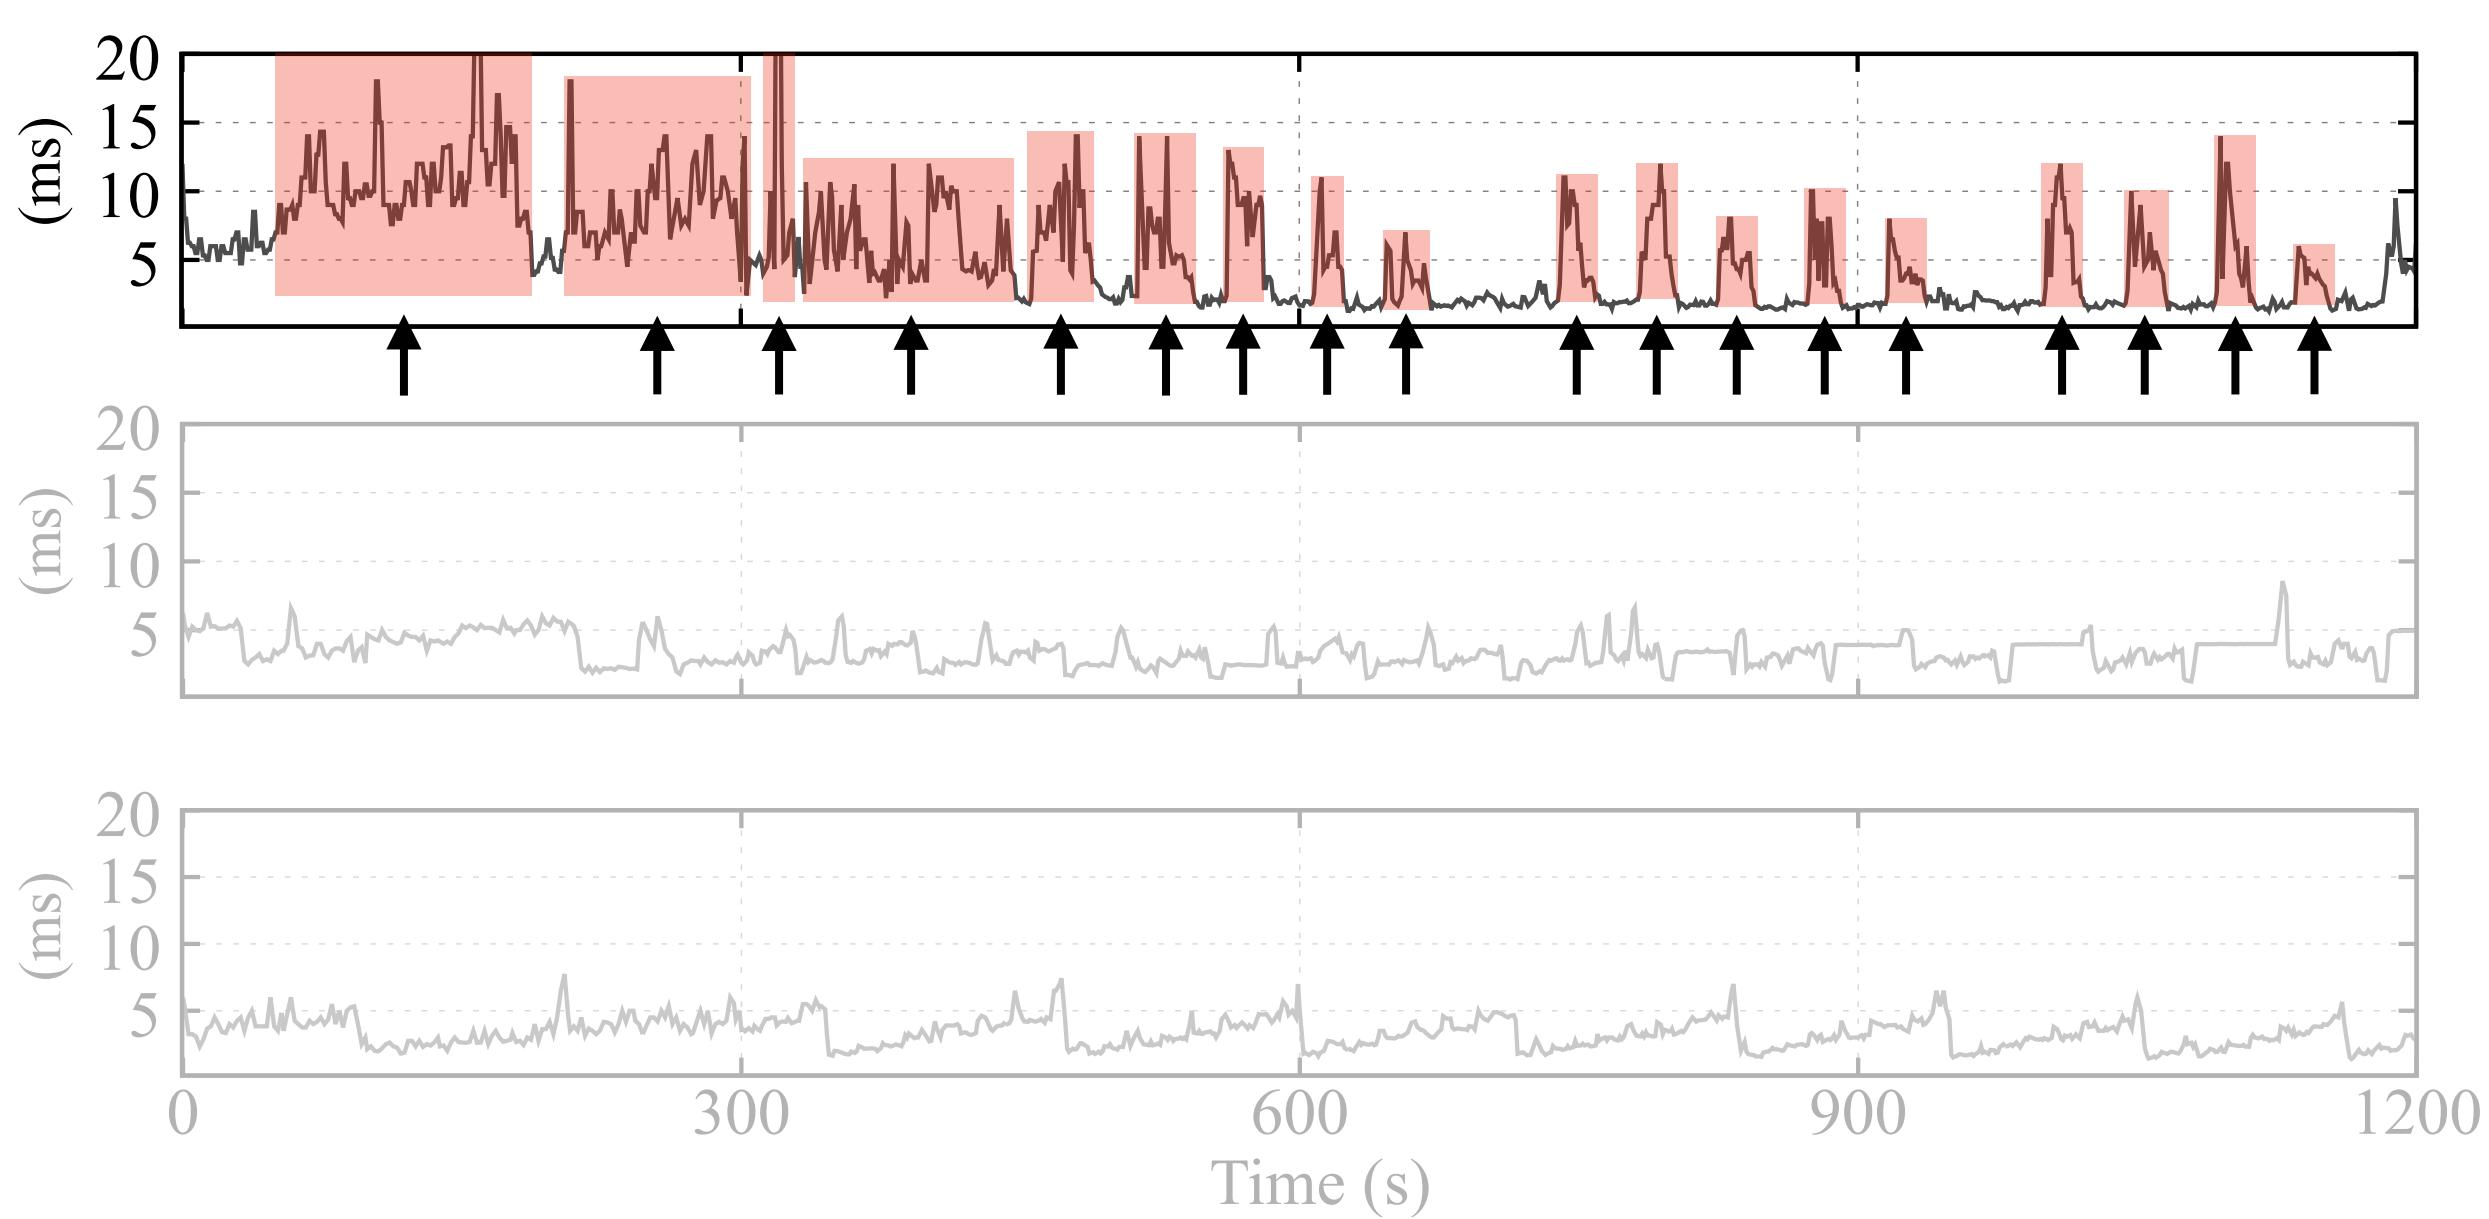


## System configuration and workload

- 8 client threads and 8 background threads
- Memory limited to 1GB and I/O BW to 200MB/s
- Bursty workload with peaks and valleys



**Throughput:** high variability due to constant flushes and compactions

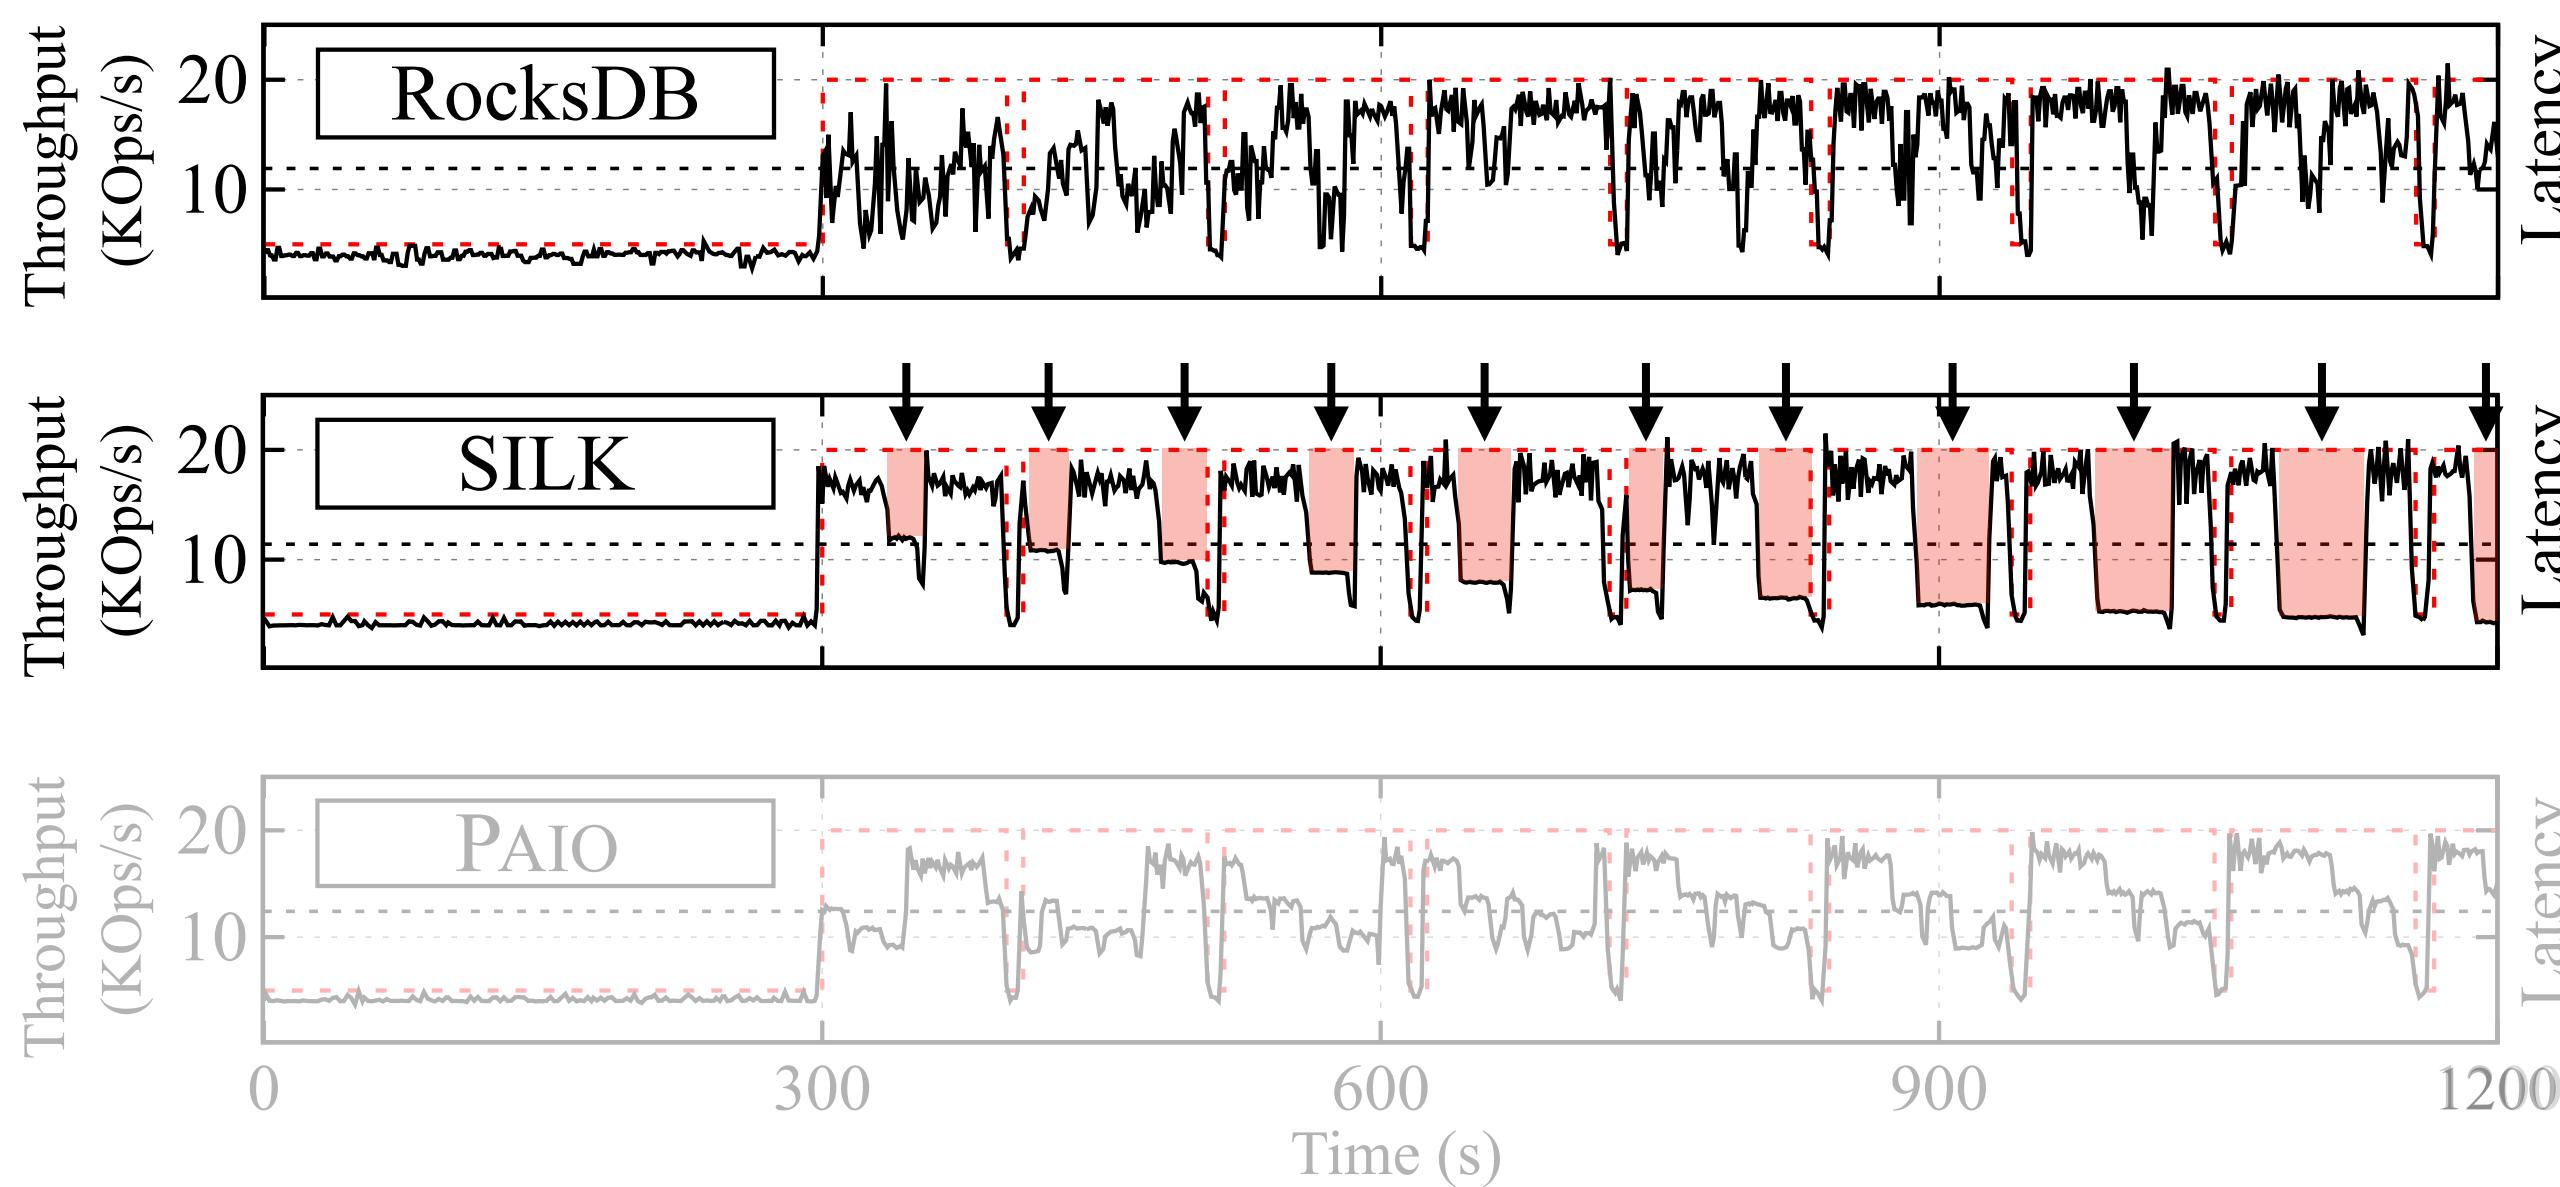


**99<sup>th</sup> latency:** high tail latency with peaks with an average range between 3 and 15 ms

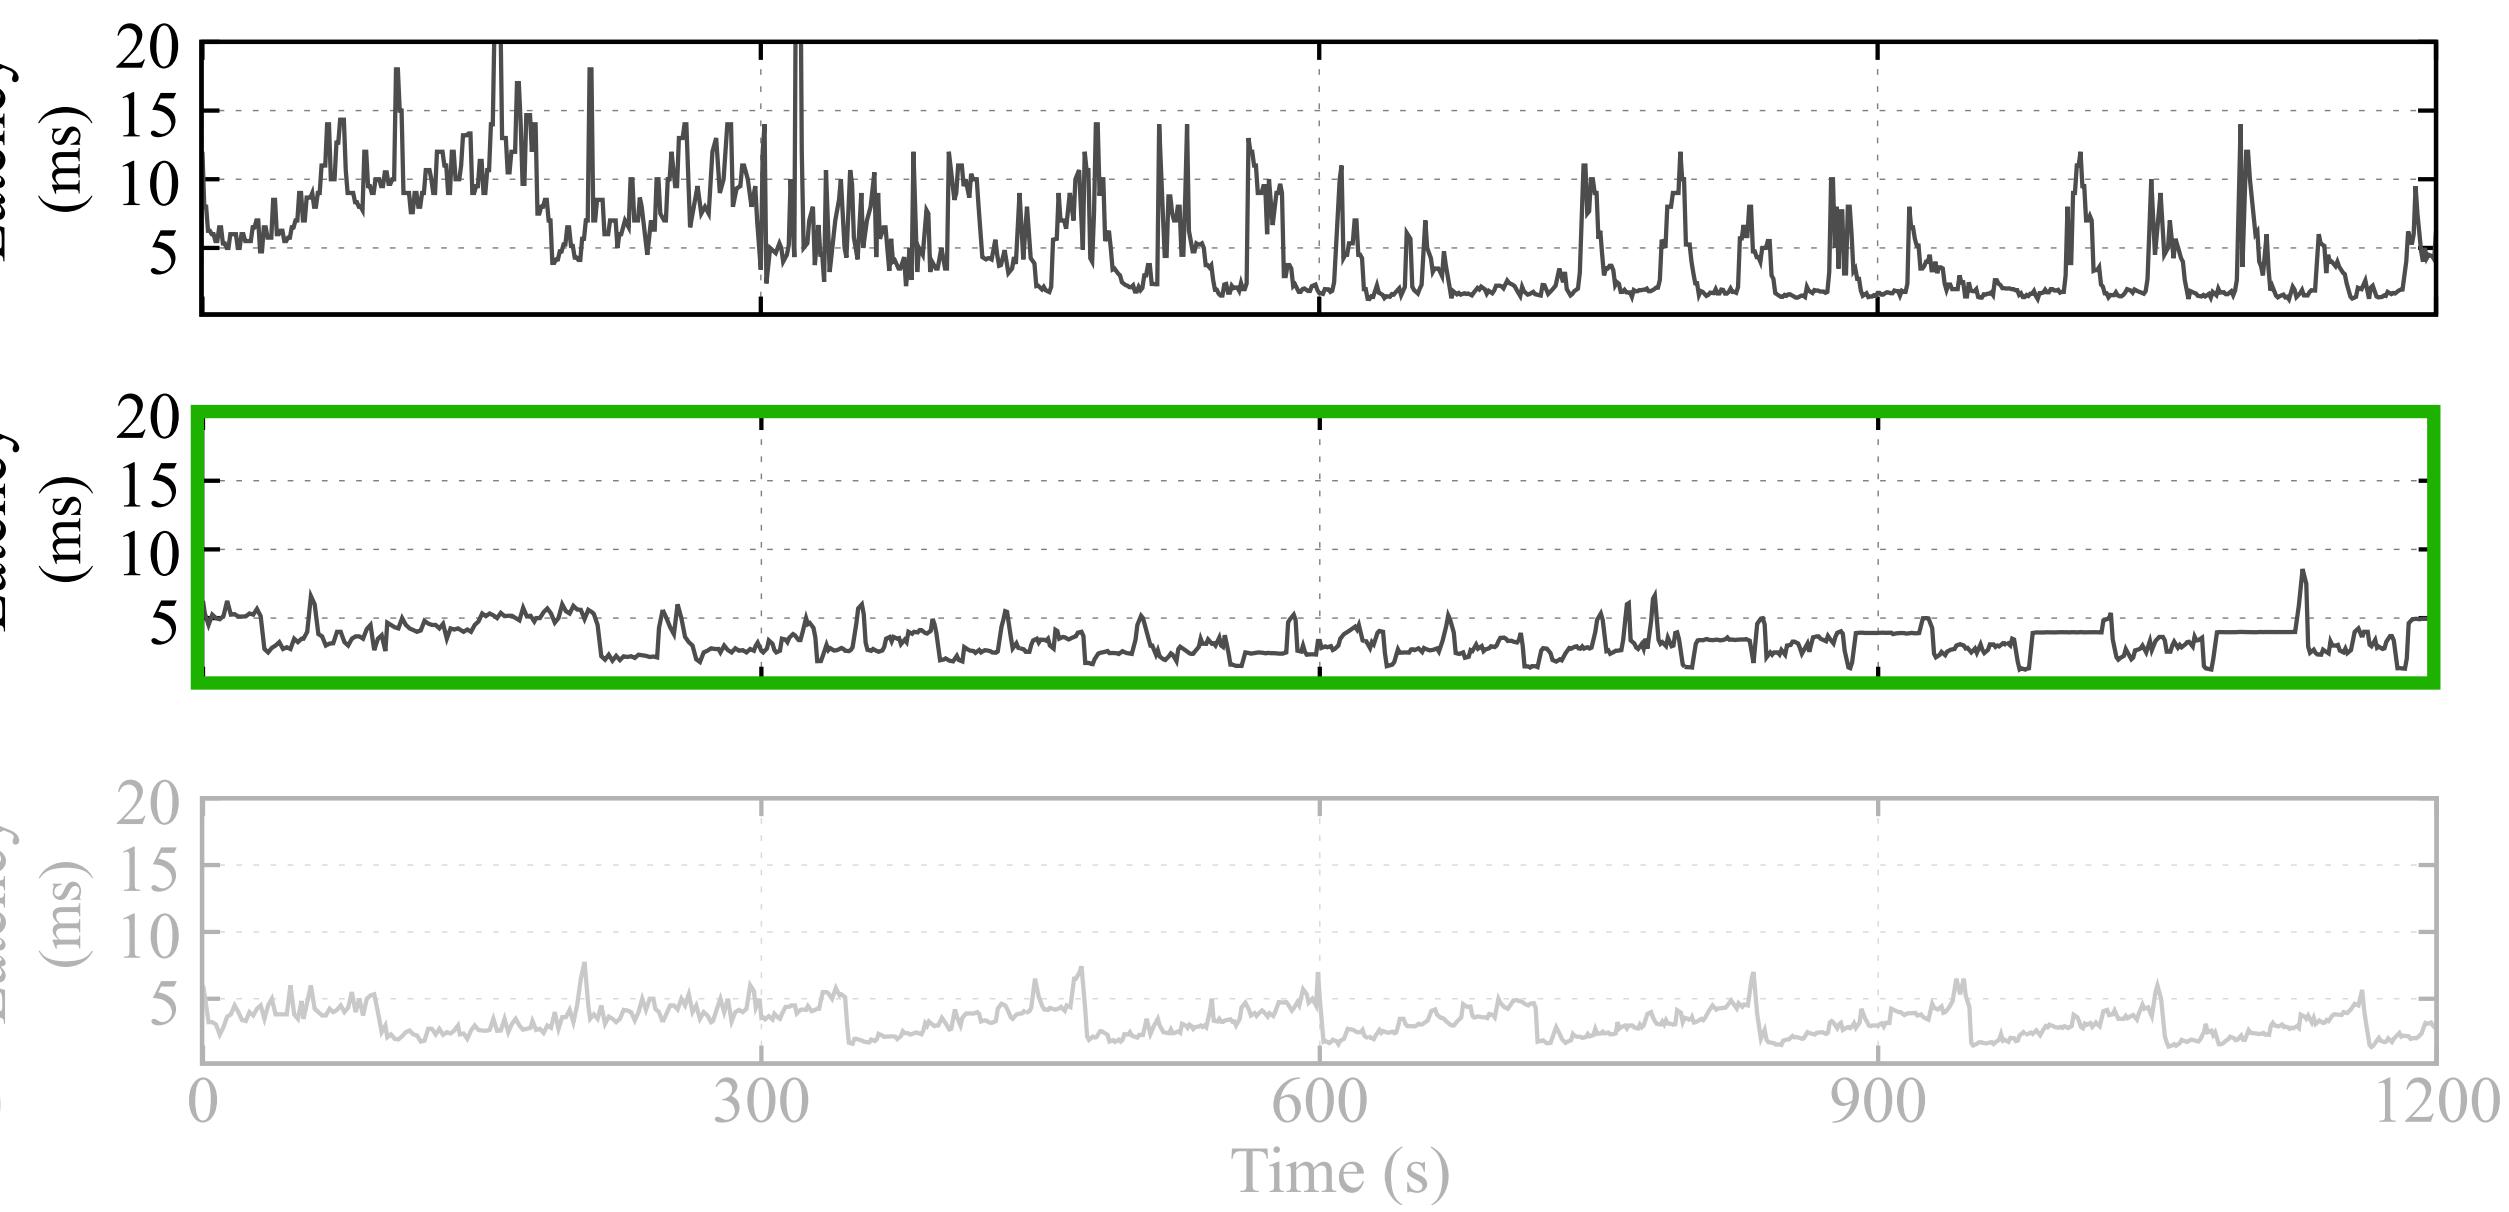


# Mixture workload

50% read 50% write



**Throughput:** suffers periodic throughput drops due to accumulated backlog

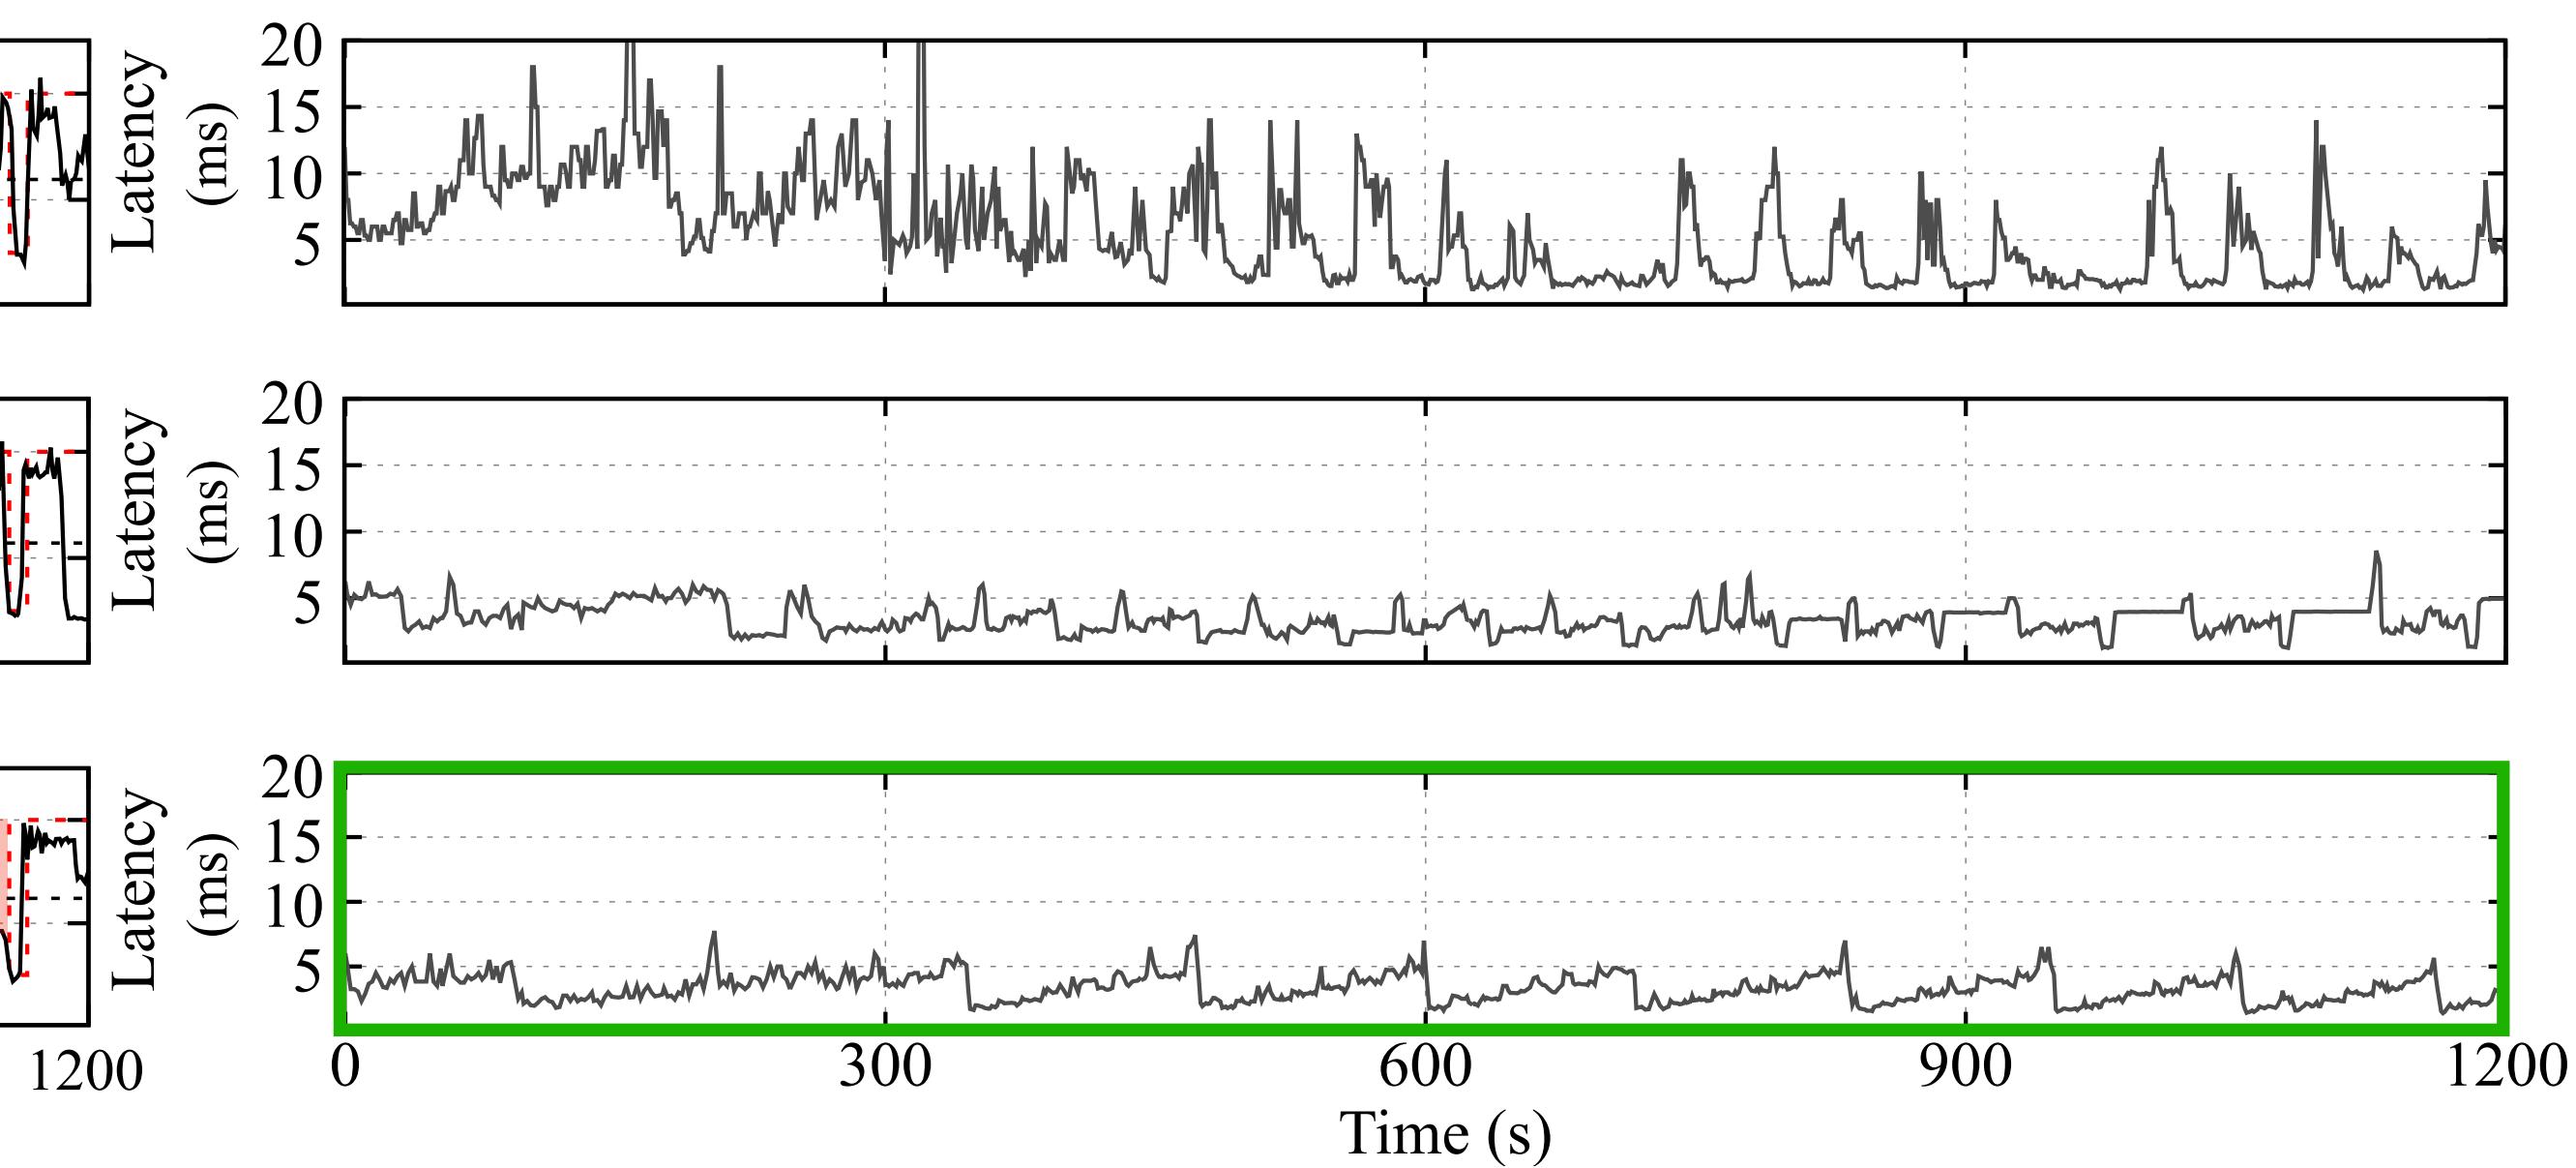
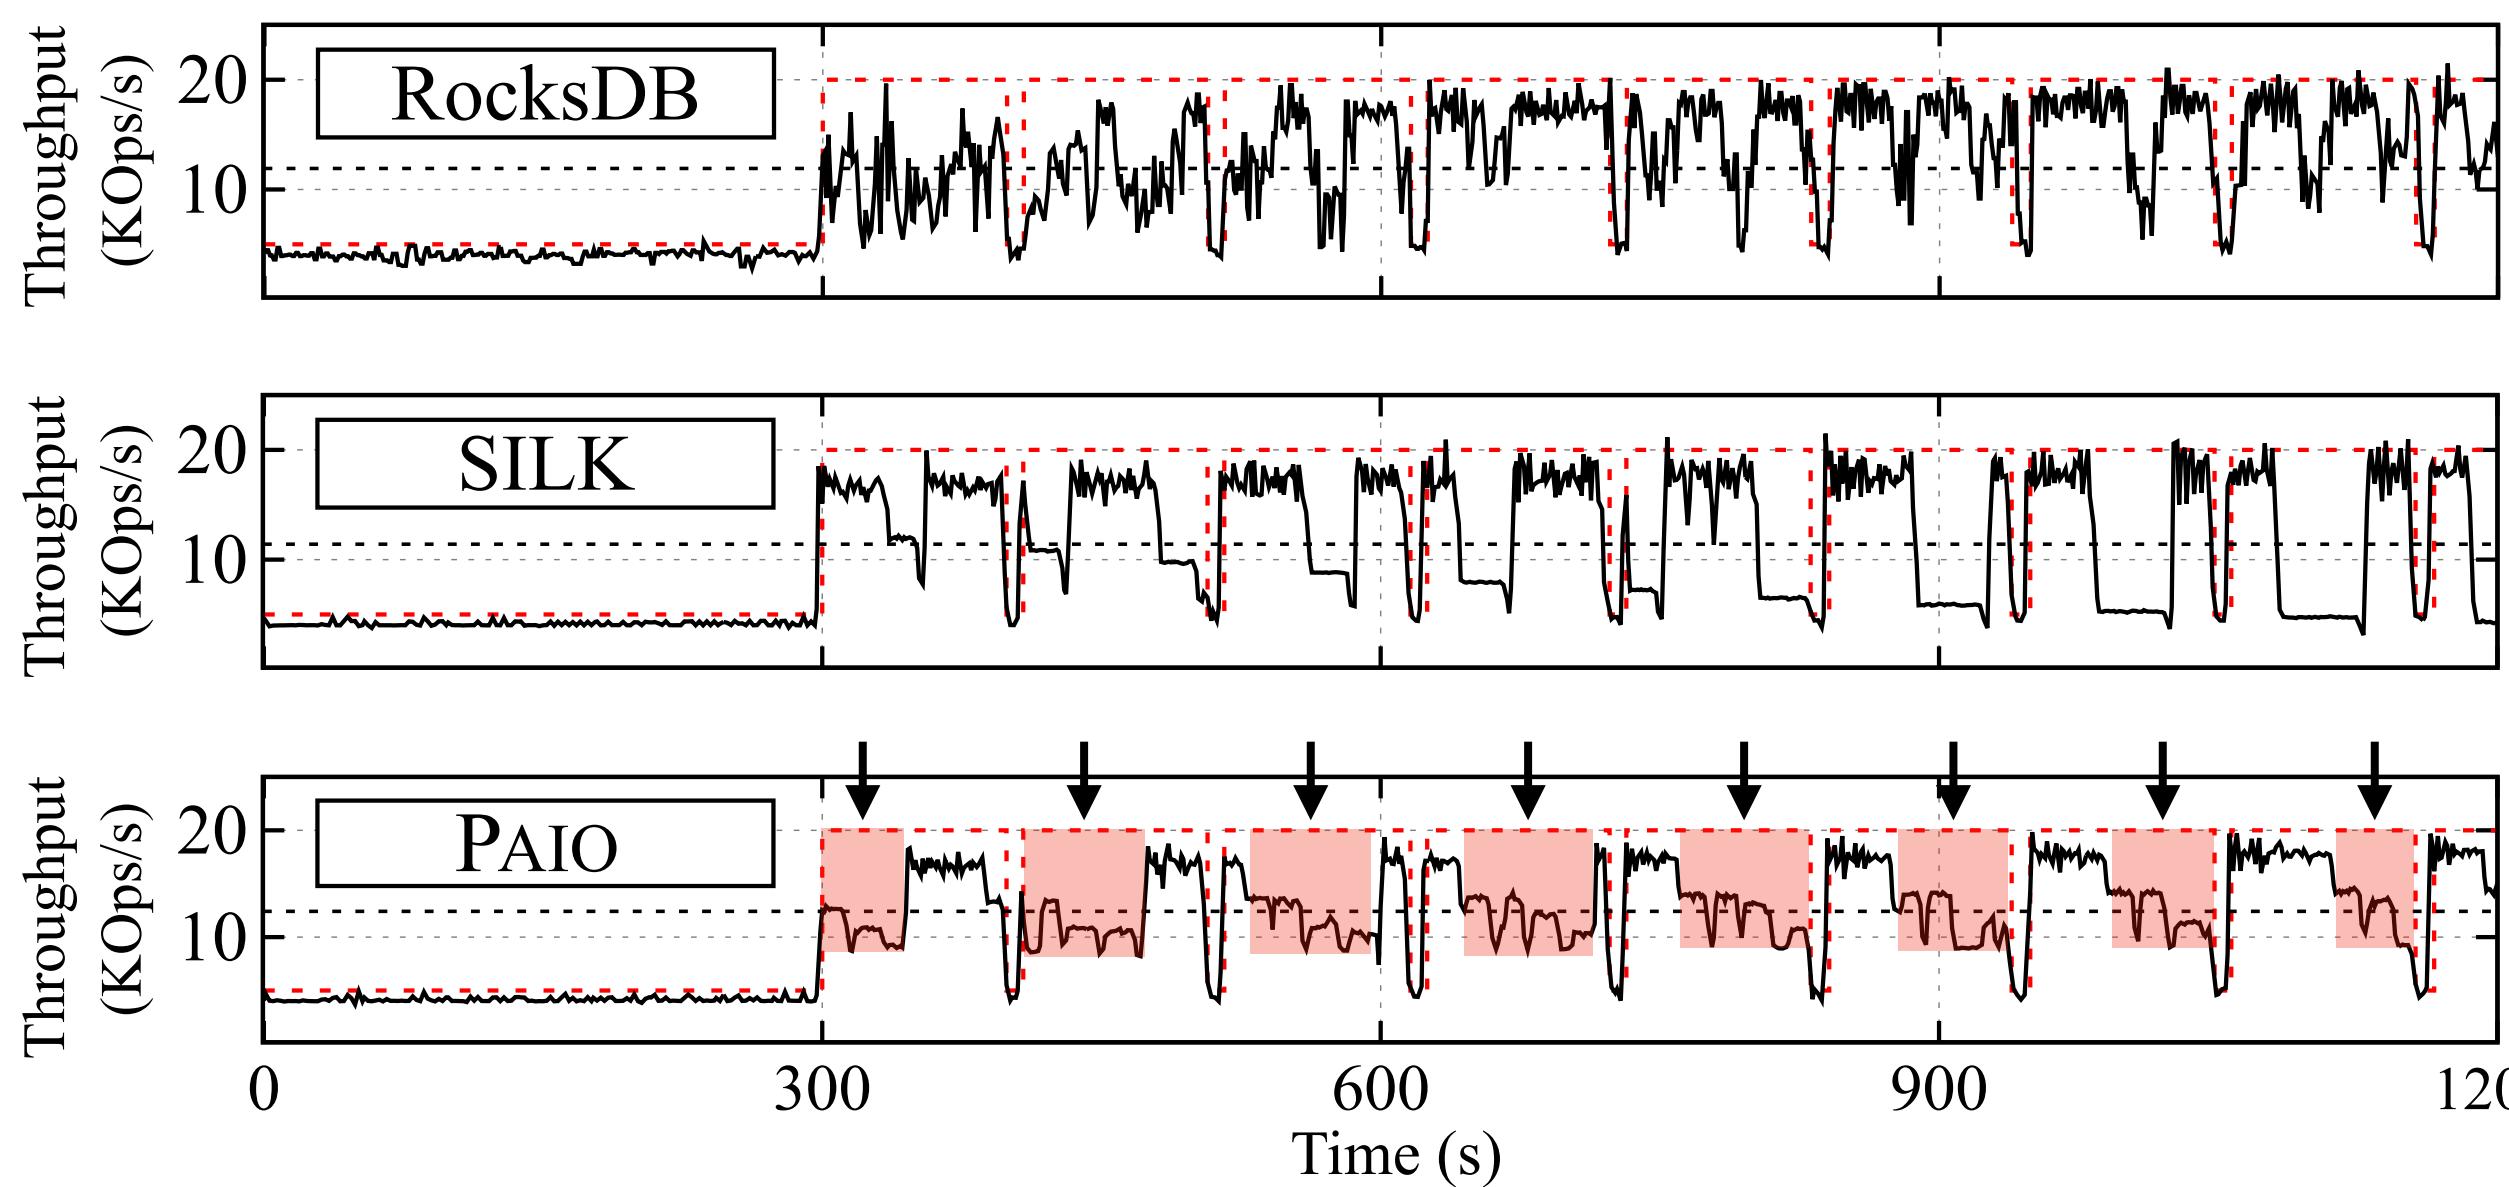


**99<sup>th</sup> latency:** low and sustained tail latency



# Mixture workload

50% read 50% write



PAIO and SILK observe a 4x decrease in absolute tail latency



# Per-application bandwidth control

## ABCI supercomputer

- Jobs can be co-located in the same compute node
- Each job runs with dedicated CPU cores, memory, GPU, and storage quota
- Local disk bandwidth is still shared, leading to I/O interference and performance variation

## BLKIO

- cgroup's block I/O controller allows static rate limiting read and write operations
- Adjusting the rate requires stopping and restarting jobs
- Cannot leverage from leftover bandwidth

## PAIO

- Stage provides the I/O mechanisms to dynamically rate limit workflows at each instance
  - Integrating PAIO in TensorFlow did not required any code changes (`LD_PRELOAD`)
- Control plane provides a proportional sharing algorithm to ensure per-application bandwidth QoS guarantees



# Per-application bandwidth control

## ABCI supercomputer

- Jobs can be co-located in the same compute node
- Each job runs with dedicated CPU cores, memory, GPU, and storage quota
- Local disk bandwidth is shared leading to I/O interference and performance variation

## BLKIO

- cgroup's block I/O controller allows **static rate limiting** read and write operations
- Adjusting the rate requires **stopping** and **restarting jobs**
- **Cannot leverage** from **leftover bandwidth**

## PAIO

- Stage provides the I/O mechanisms to dynamically rate limit workflows at each instance
  - Integrating PAIO in TensorFlow did not required any code changes (LD\_PRELOAD)
- Control plane provides a proportional sharing algorithm to ensure per-application bandwidth QoS guarantees



# Per-application bandwidth control

## ABCI supercomputer

- Jobs can be co-located in the same compute node
- Each job runs with dedicated CPU cores, memory, GPU, and storage quota
- Local disk bandwidth is shared leading to I/O interference and performance variation

## BLKIO

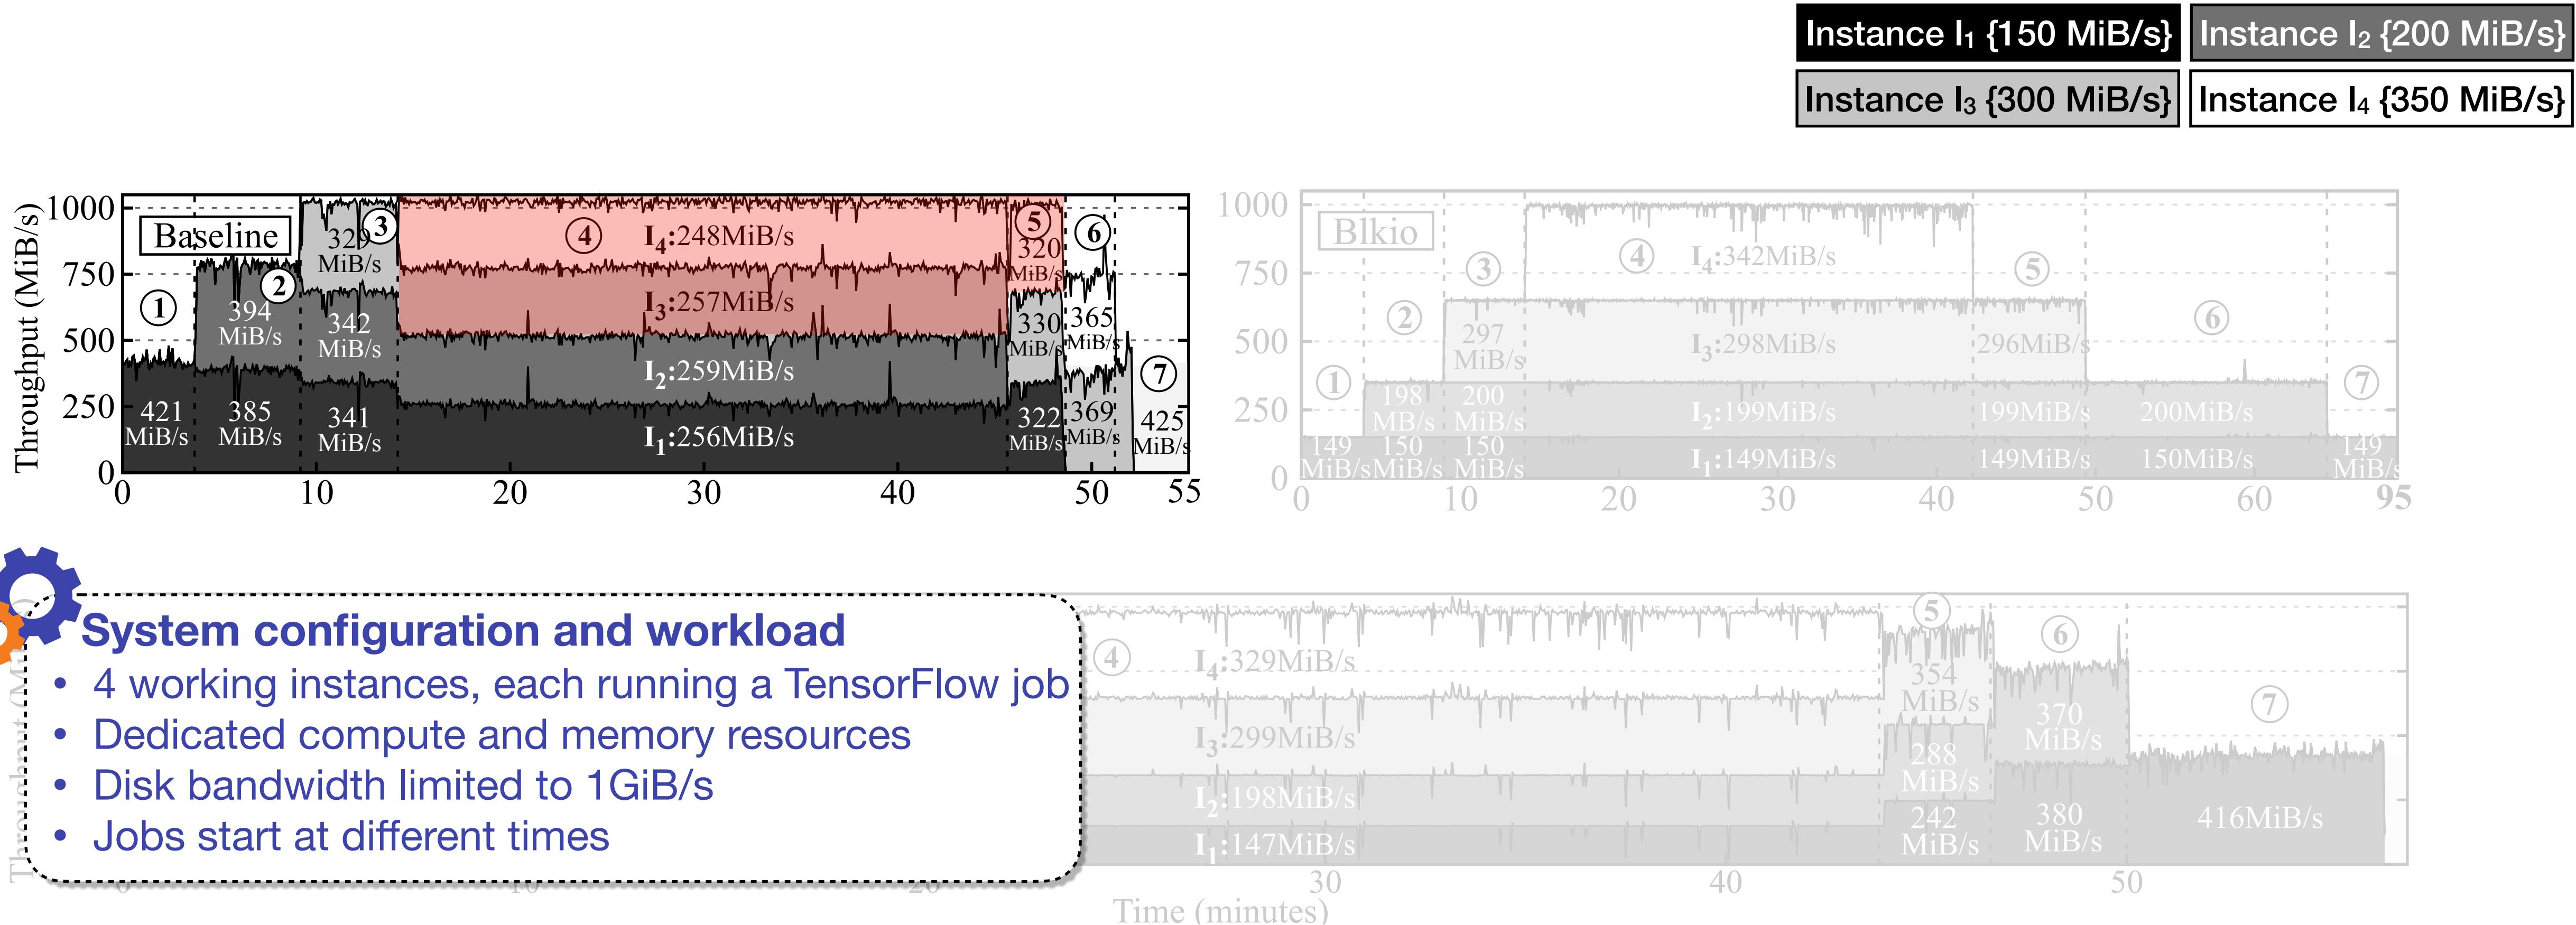
- cgroup's block I/O controller allows static rate limiting read and write operations
- Adjusting the rate requires stopping and restarting jobs
- Cannot leverage from leftover bandwidth

## PAIO

- Stage provides the I/O mechanisms to **dynamically rate limit workflows** at each instance
  - Integrating PAIO in TensorFlow **did not required any code changes** (`LD_PRELOAD`)
- Control plane provides a **proportional sharing algorithm** to ensure per-application bandwidth QoS guarantees



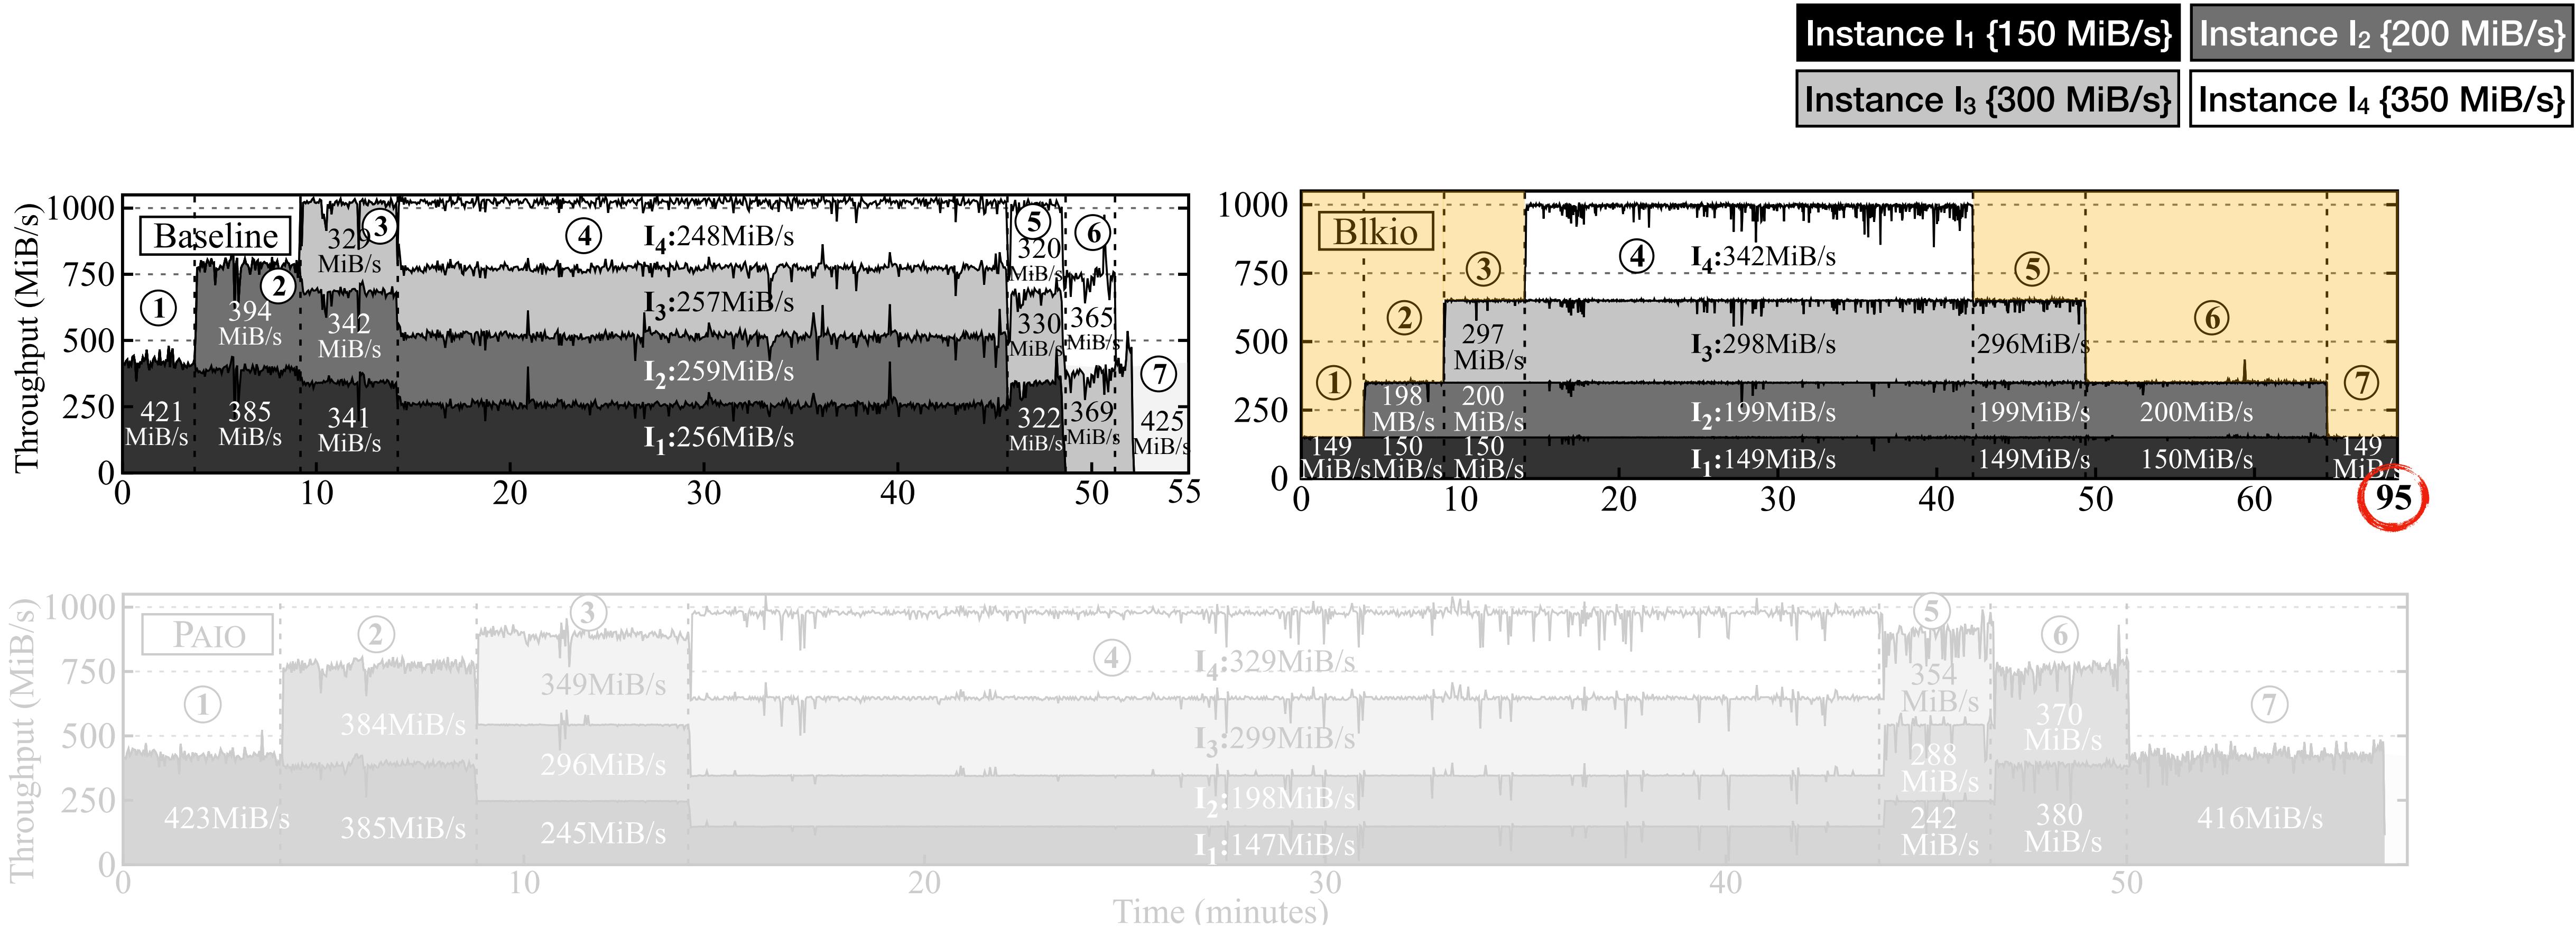
# Per-application bandwidth control



I<sub>3</sub> and I<sub>4</sub> cannot meet their bandwidth targets during 31 and 34 minutes



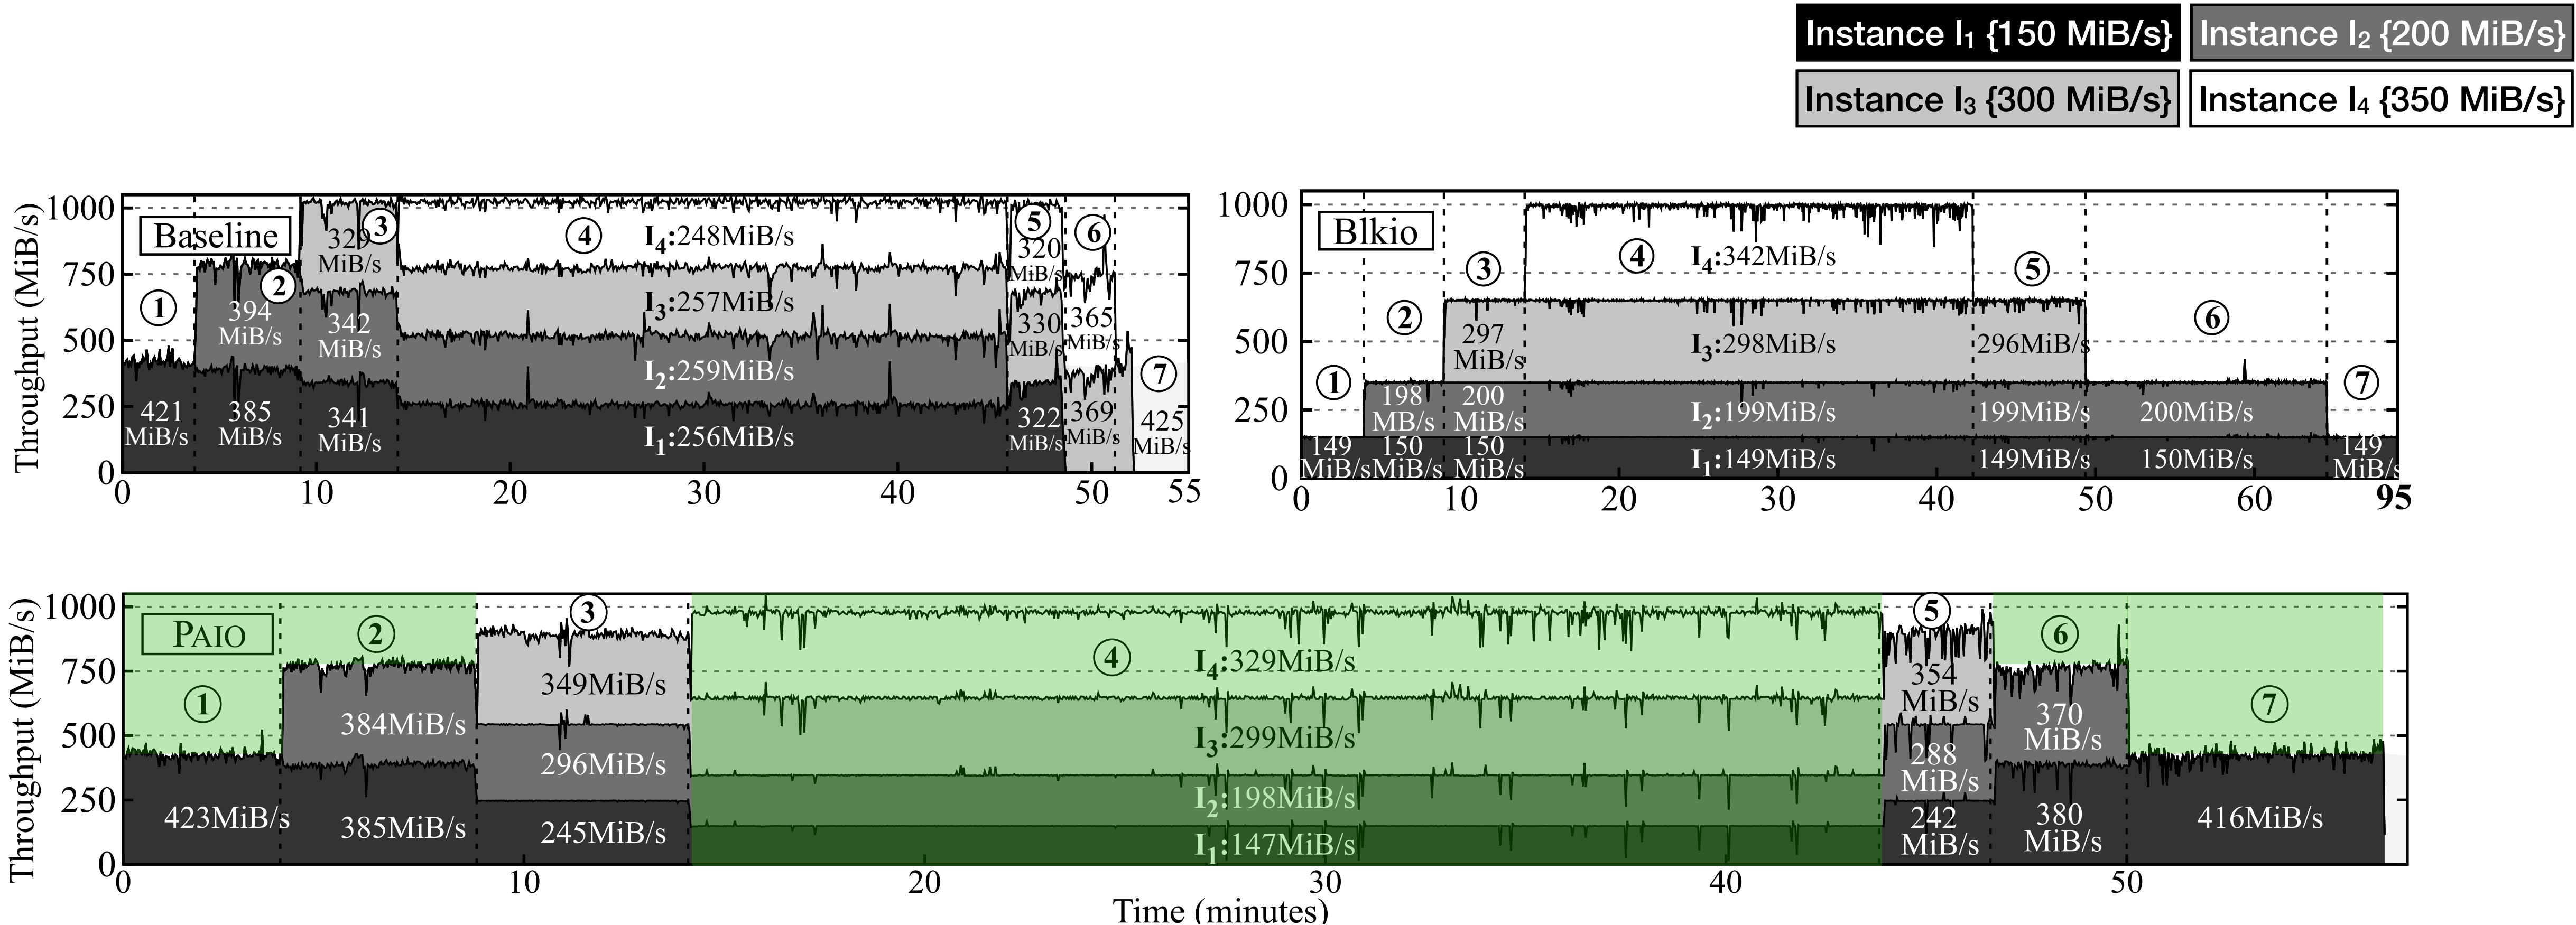
# Per-application bandwidth control



Instances cannot be dynamically provisioned with available disk bandwidth



# Per-application bandwidth control



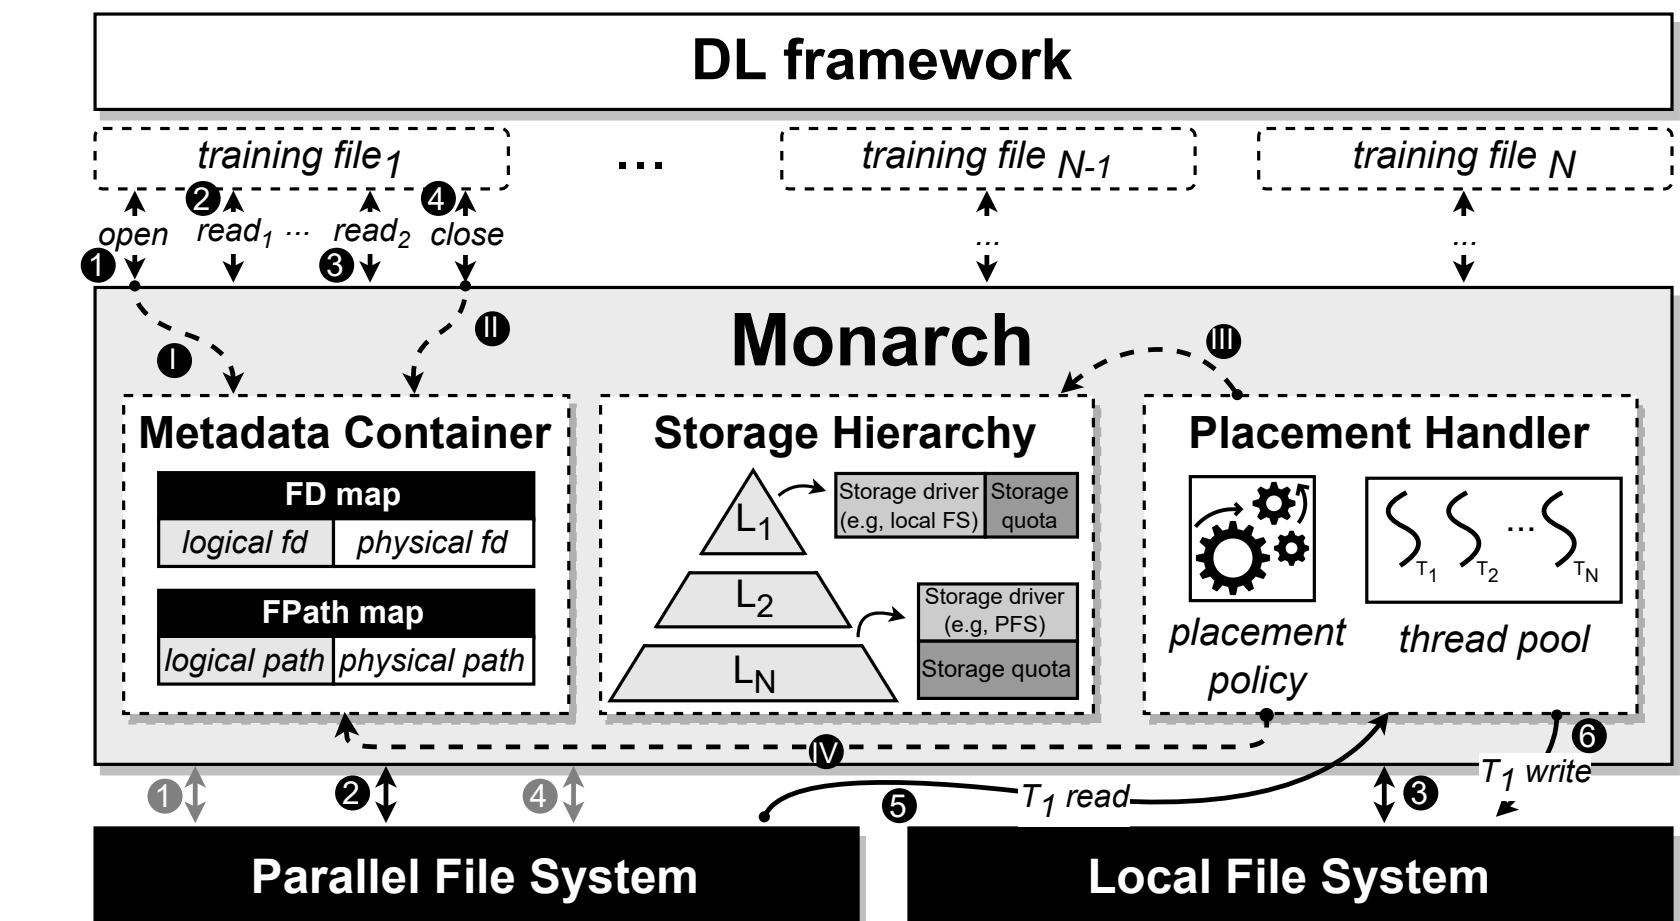
PAIO ensures that policies are met at all times, and whenever leftover bandwidth is available, PAIO shares it across active instances



# Storage data planes for deep learning

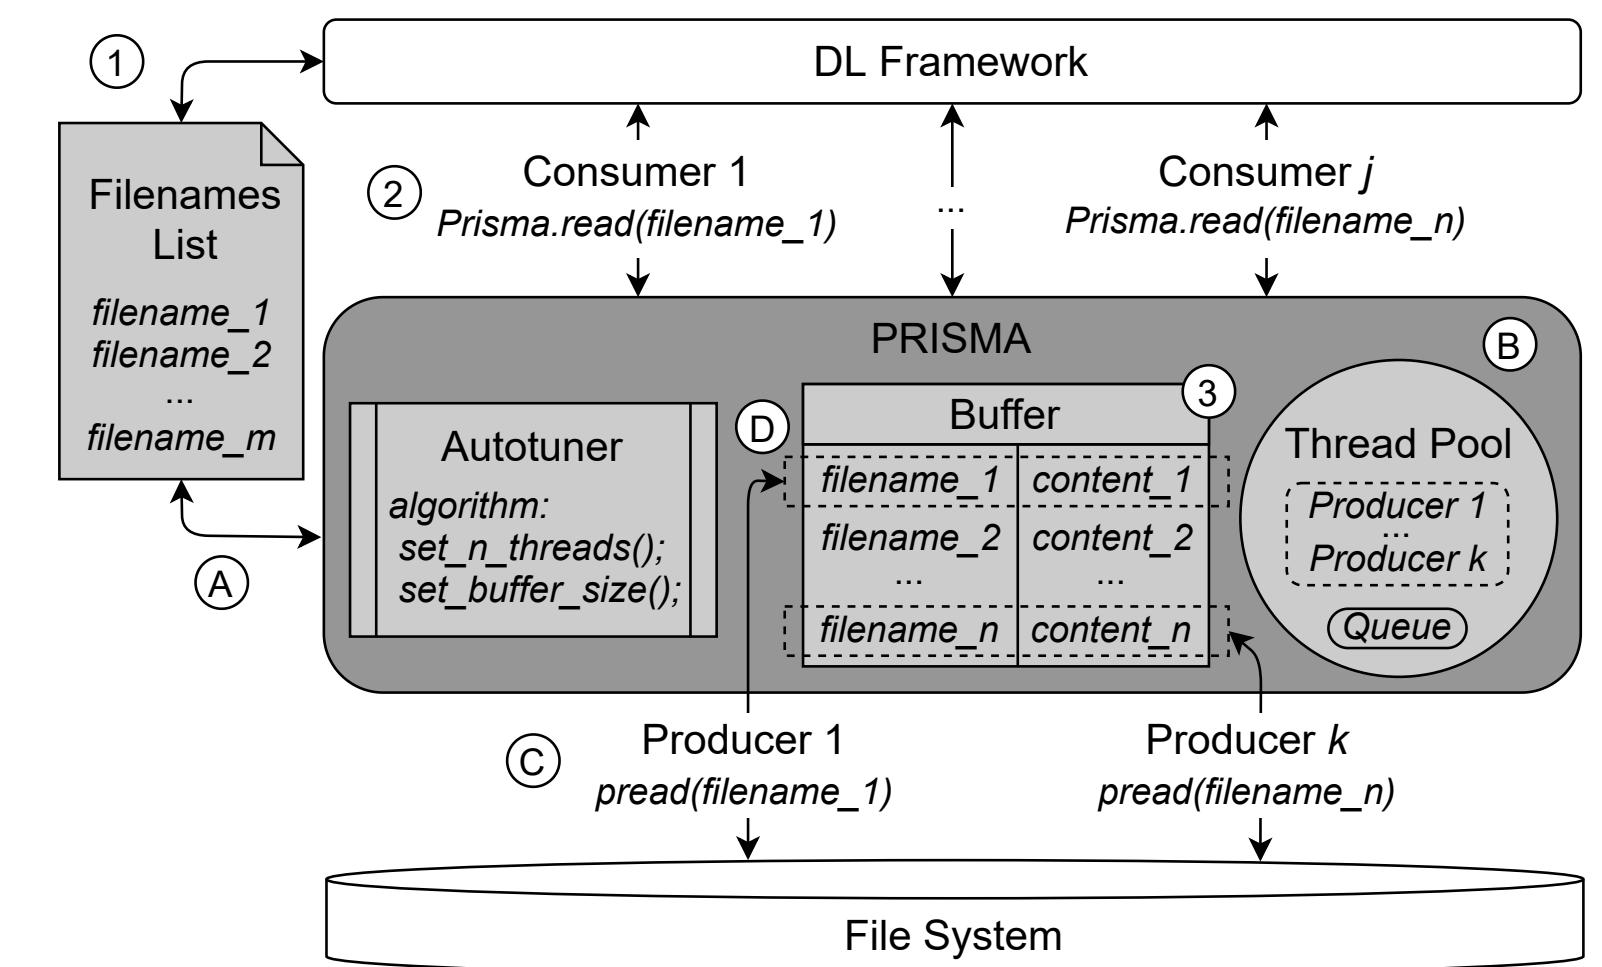
## Storage tiering (Monarch)

- Framework-agnostic storage middleware
- Leverages existing storage tiers of supercomputers
- Accelerates DL training time by up to 28% and 37% in TensorFlow and PyTorch
- Decreases the operations submitted to the PFS



## Parallel data prefetching (Prisma)

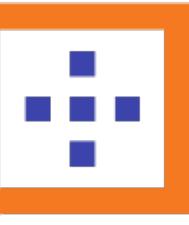
- Data plane for prefetching training data samples
- Significantly outperforms baseline PyTorch and TensorFlow configurations
- Achieves similar performance as carefully engineered I/O optimizations in TensorFlow



[7] "Accelerating Deep Learning Training Through Transparent Storage Tiering". Dantas et al. ACM/IEEE CCGrid 2022.

[8] "Monarch: Hierarchical Storage Management for Deep Learning Frameworks". Dantas et al. IEEE Cluster@Rex-IO 2021.

[9] "The Case for Storage Optimization Decoupling in Deep Learning Frameworks". Macedo et al. IEEE Cluster@Rex-IO 2021.



# Summary and takeaways

- PAIO, a **user-level** framework to build **custom-made** storage **data plane stages**
- Combines ideas from **Software-Defined Storage** and **context propagation**
- **Decouples** system-specific optimizations to **dedicated I/O layers**
- **User-level data planes** enable similar **control** and **I/O performance** as system-specific optimizations
  - Can be applied over (a lot of) different storage scenarios ...

## Q & A

**Ricardo Macedo**

✉ [ricardo.g.macedo@inesctec.pt](mailto:ricardo.g.macedo@inesctec.pt)

/github.com/dsrhaslab

[dsr-haslab.github.io](http://dsr-haslab.github.io)

